

DEEPMLEARNING ARCHITECTURES

FOCUS: CNN & RNN

SYLLA N'FALY | ENG. STUDENT, COMPUTER SCIENCE

I'M,

SYLLA N'FALY

Engineering student in
Computer Science at
ENSA Tangier, Morocco

A passionate Data
Science enthusiast



“Seeing” the progress of deep learning throughout the years



2015

Goodfellow et al.



2018

Karras, Laine, Aila.



2020

MIT Intro to Deep Learning



3 months ago

That is easily the cleanest visual deepfake I've ever seen. It must have taken ages to render, because it just looks flawless.



5 months ago

THAT INTRO TO THE LECTURE IS SAVAGE!!!



2 months ago

WOW WOW WOW i am amazed.



3 months ago

This is the best example of a Course that sells itself. 😊

2020

...creating this 2 minute video required...

2 hours of professional audio

50 hours of HD video

Static, pre-defined script

Over \$15K USD of compute

2020

...creating this 2 minute video required...

2 hours of professional audio

50 hours of HD video

Static, pre-defined script

Over \$15K USD of compute

2025

...fast forward a few years...

?

OUR PLAN

01

The Foundations

02

Computer Vision

FCNN & CNN

03

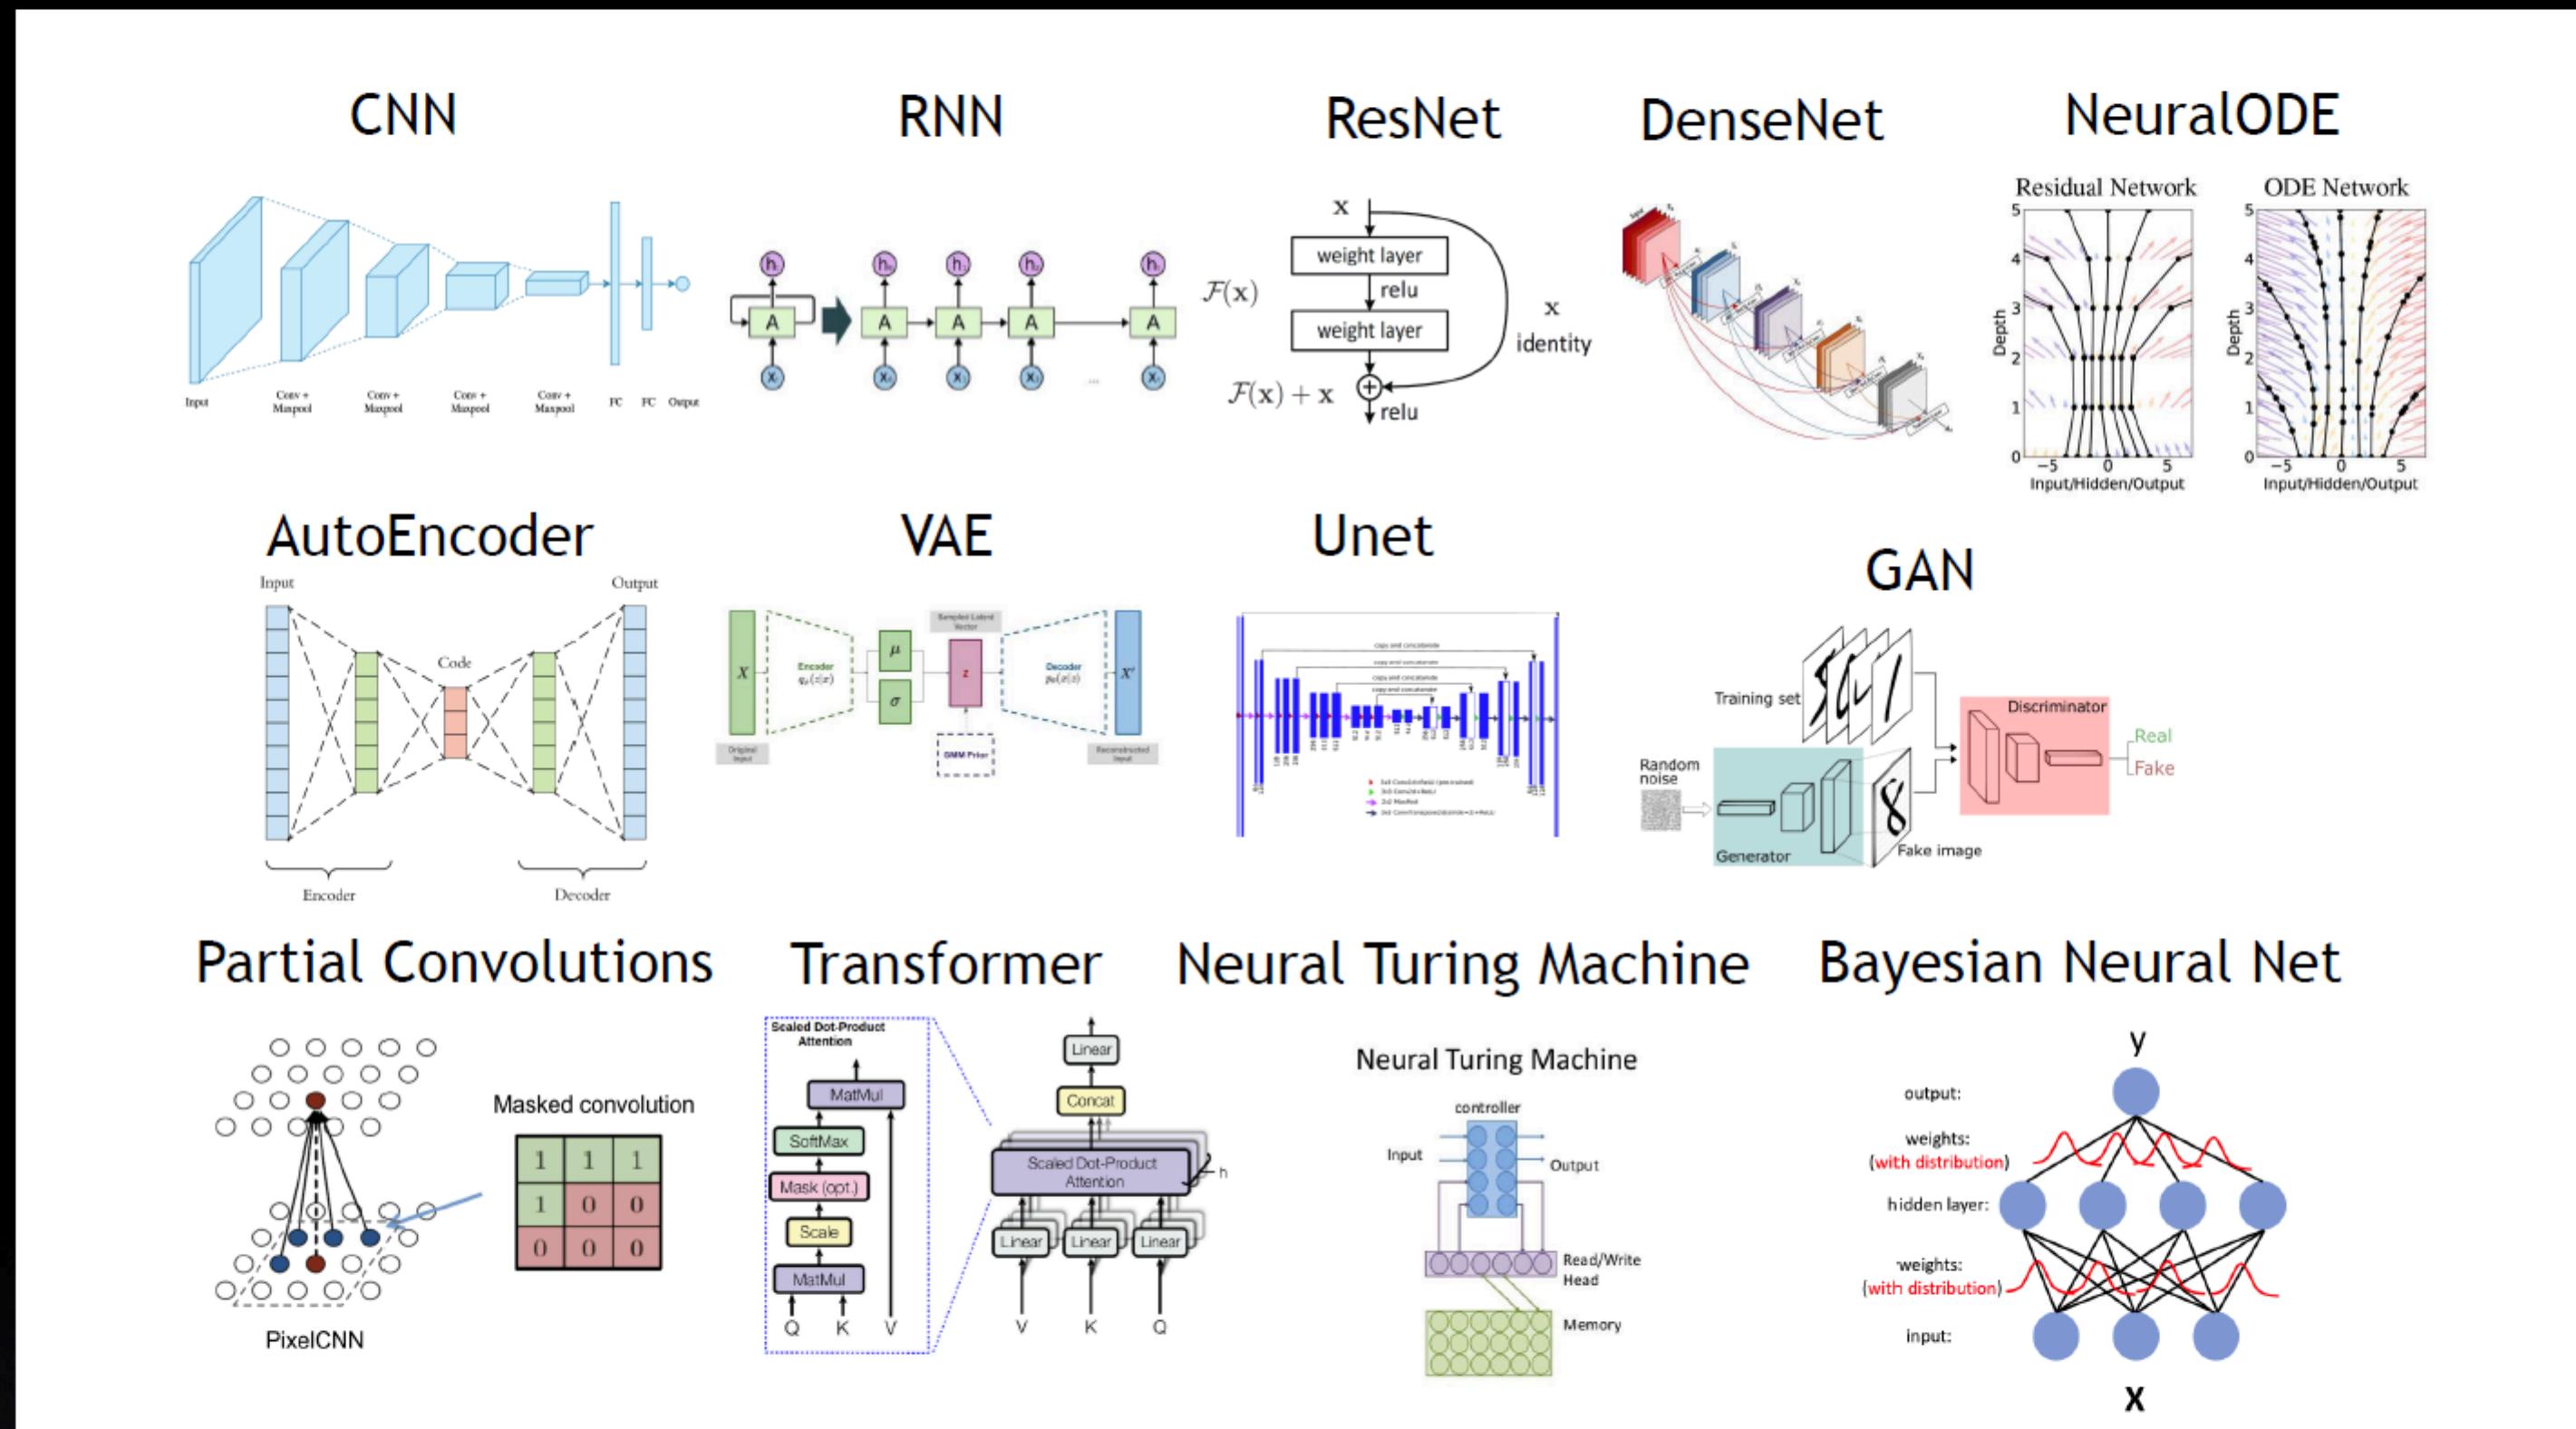
Deep Sequence Modeling

RNN, LSTM & GRU

04

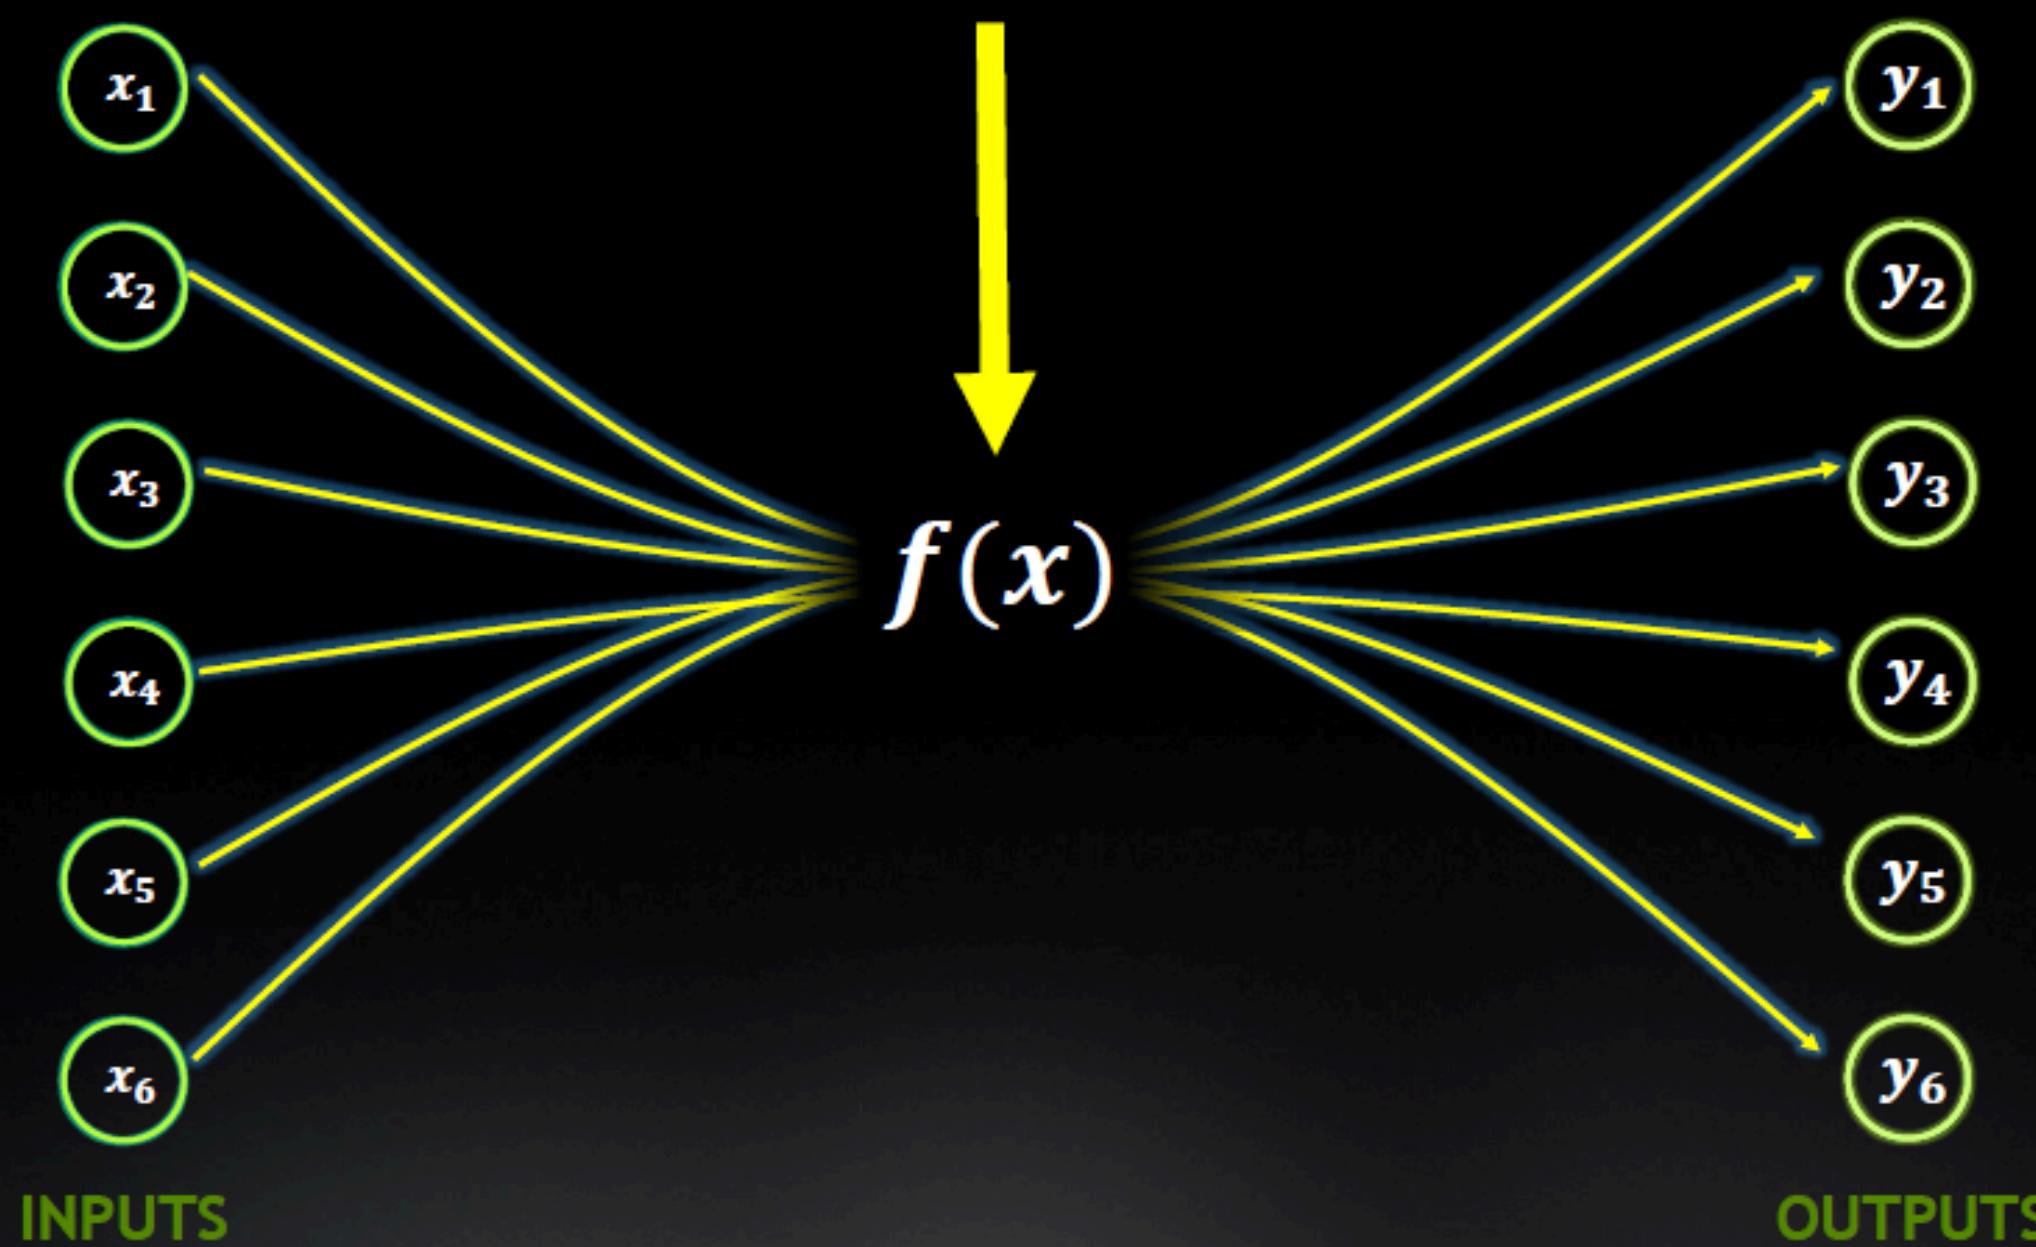
Q/R

DEEP LEARNING MODEL ZOO



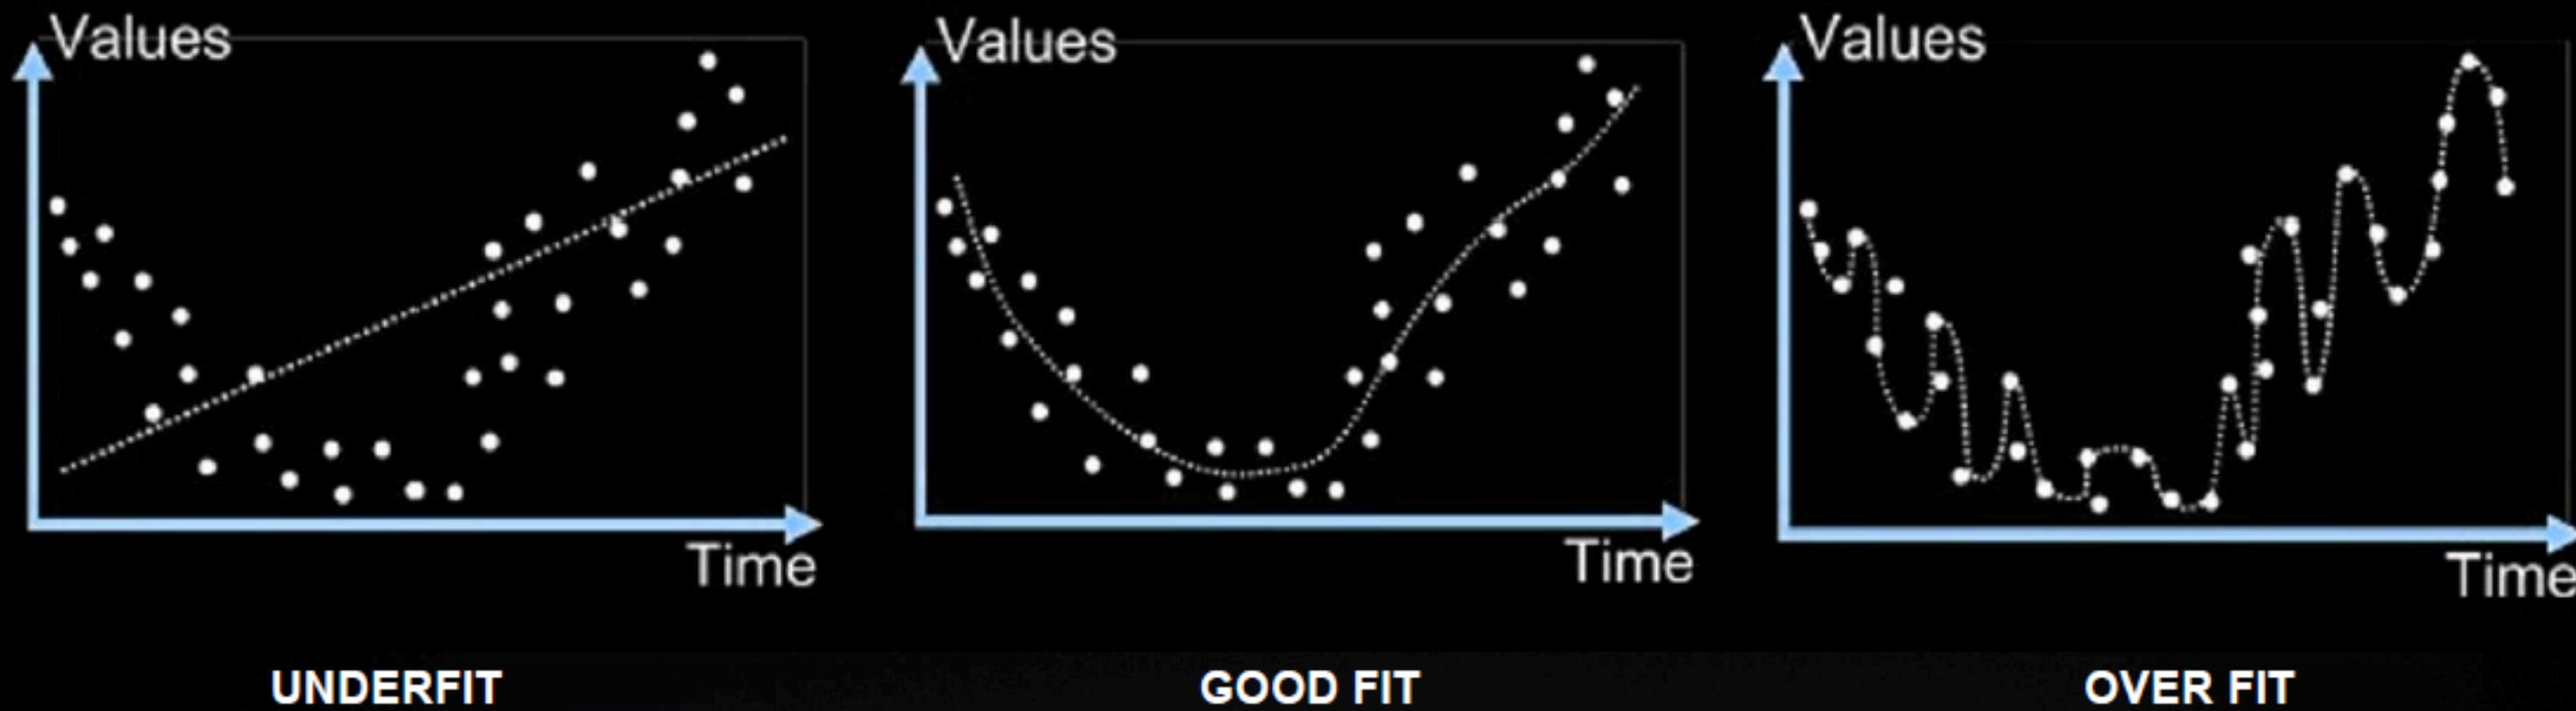
REVERSE-ENGINEER FUNCTIONS FROM EXAMPLES

Find this, automatically



ADJUST MODEL CAPACITY TO FIT YOUR DATA

A good model is one that generalizes to new data



UNDERFIT

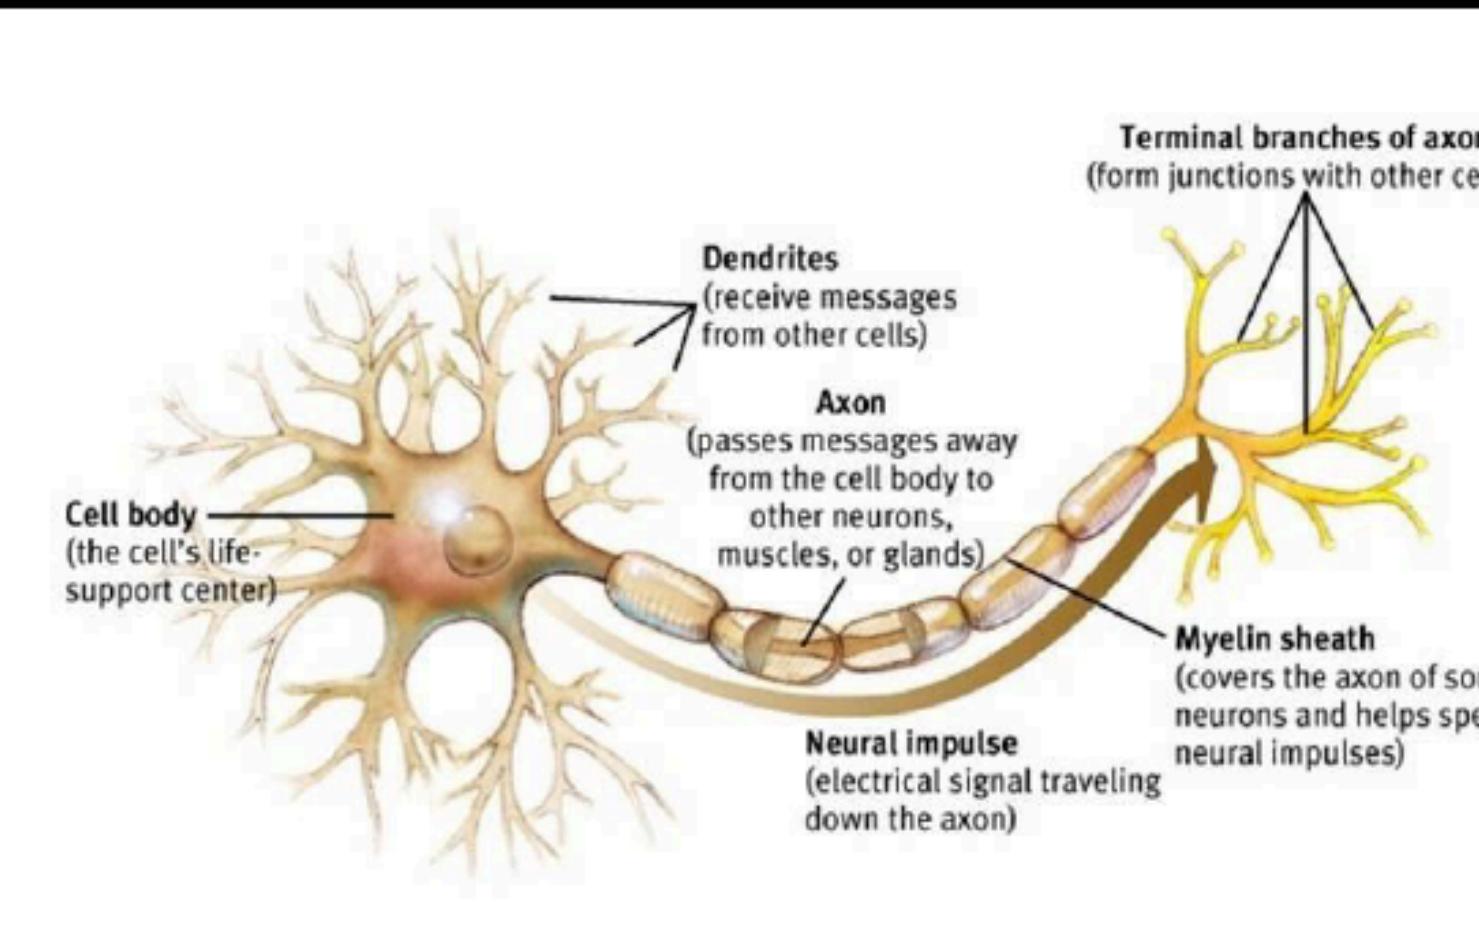
GOOD FIT

OVER FIT

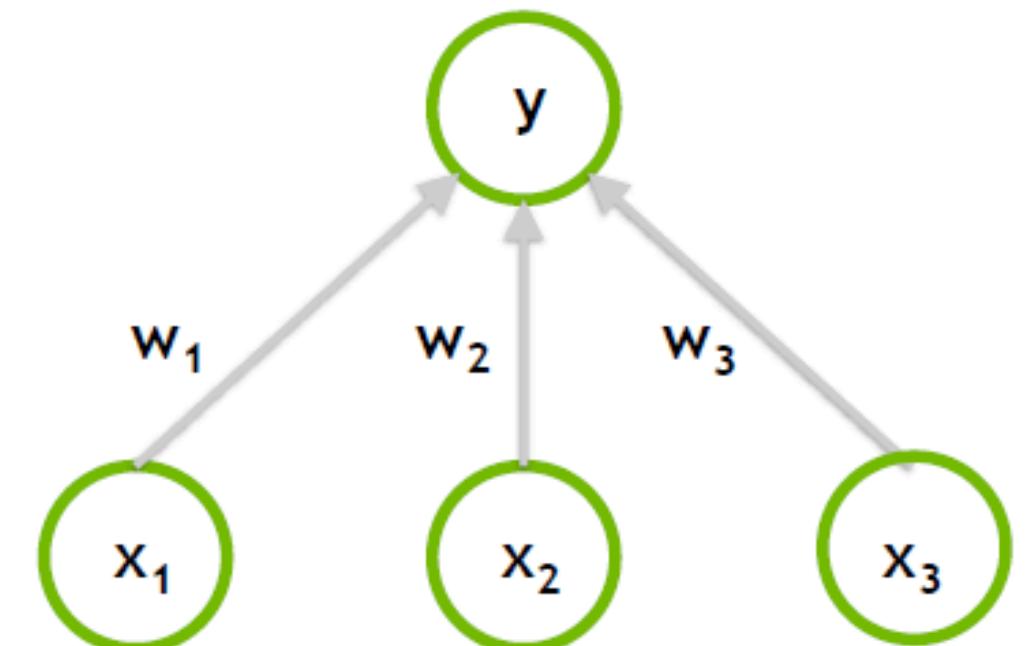
ARTIFICIAL NEURONS

Are simple equations with a set of adjustable parameters

Biological neuron



Artificial neuron

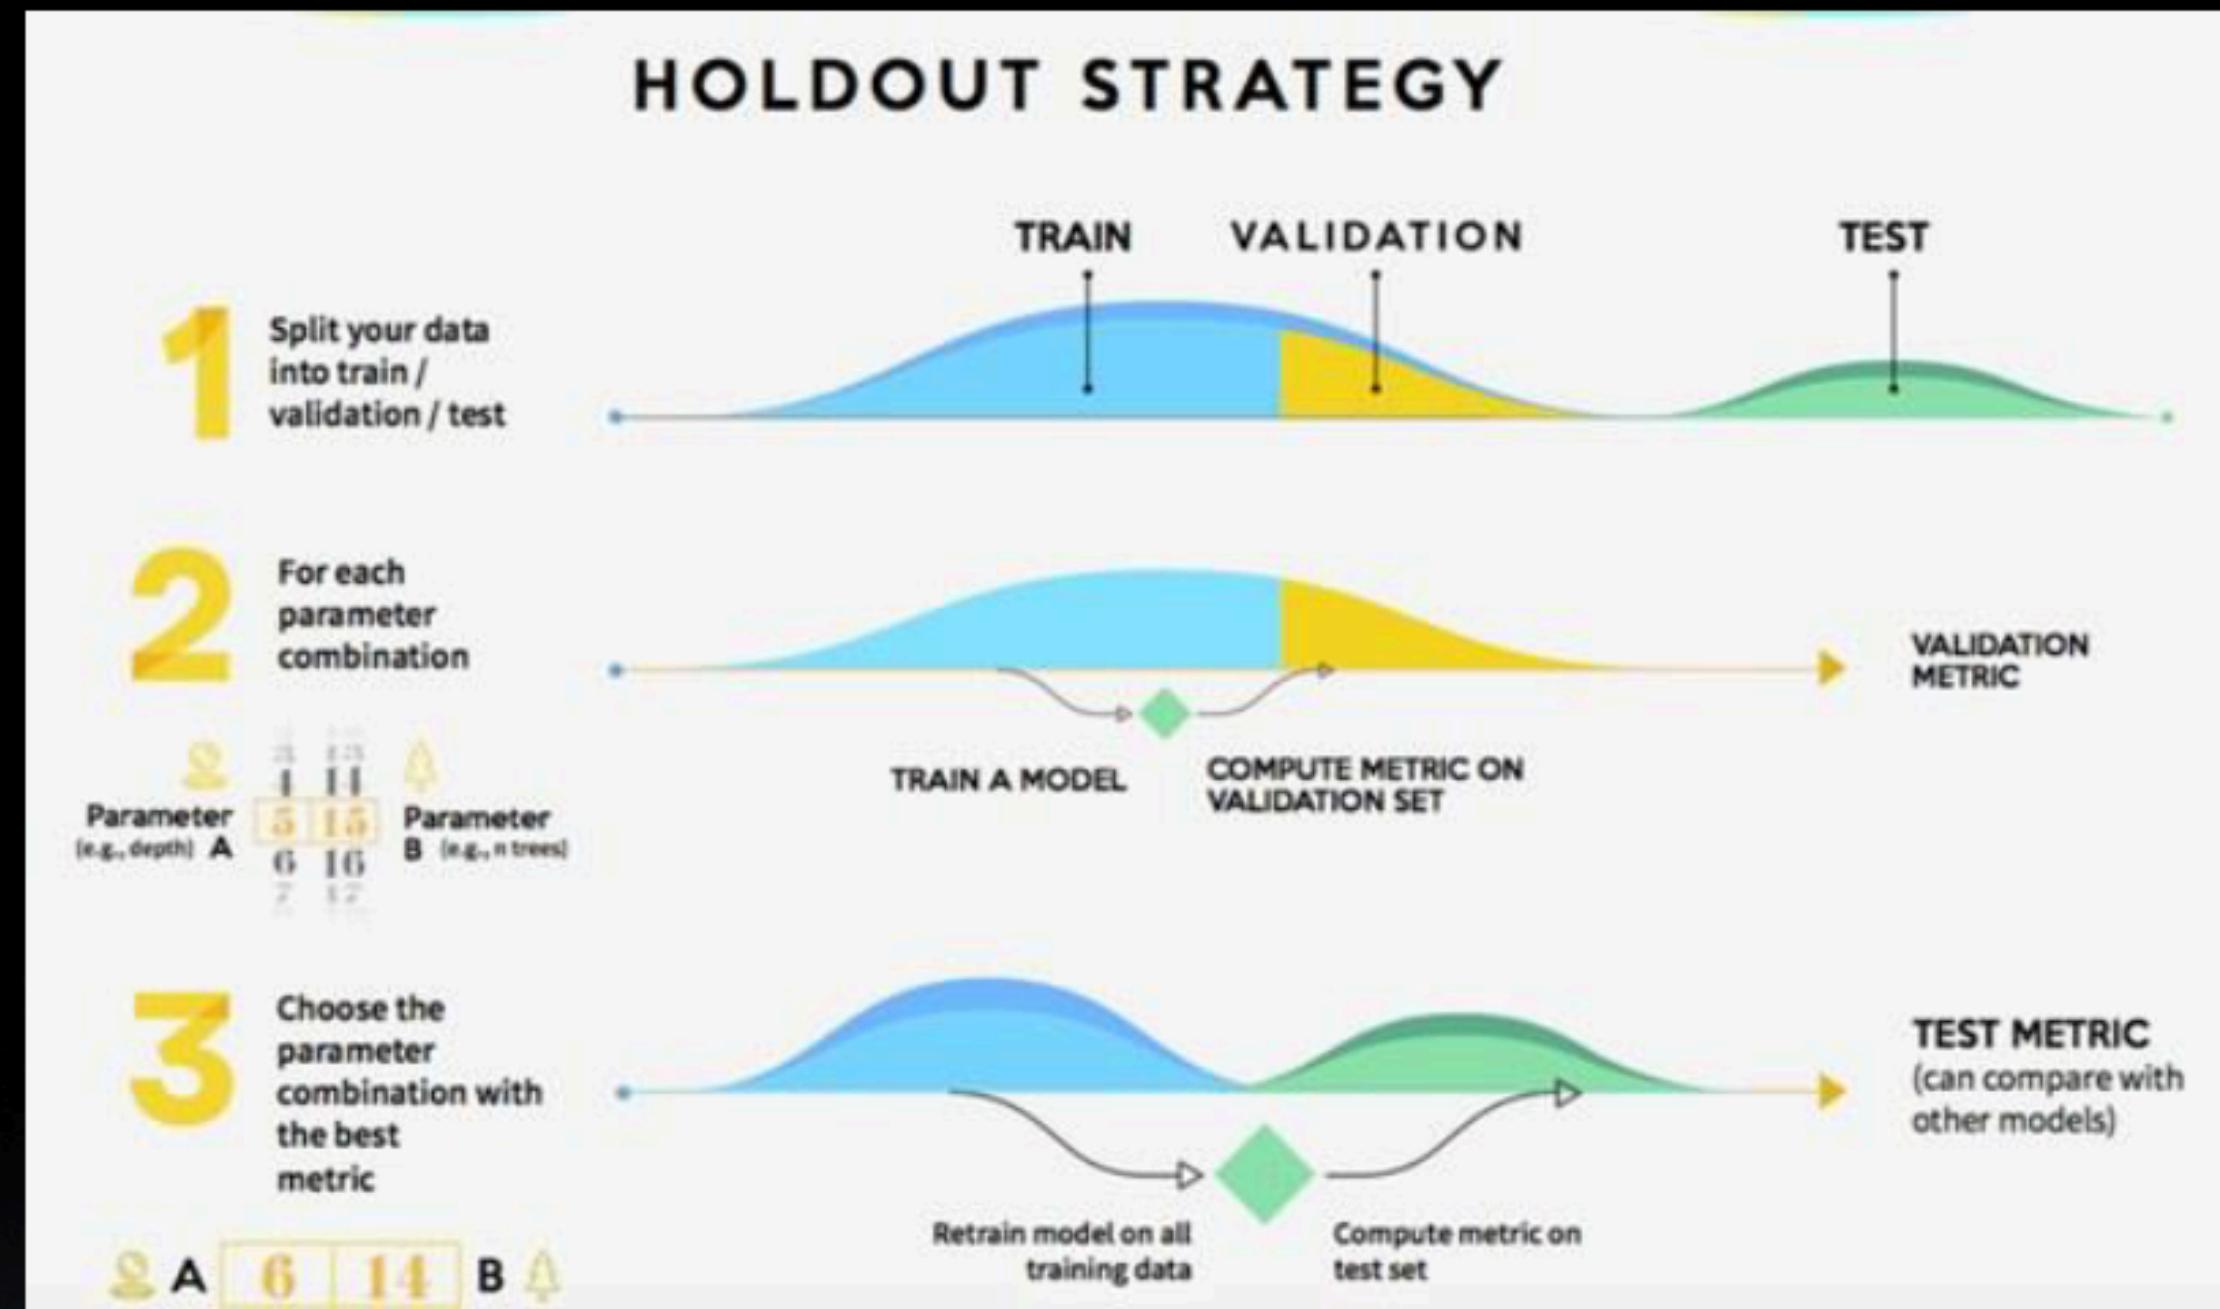


$$y = f(w_1x_1 + w_2x_2 + w_3x_3)$$

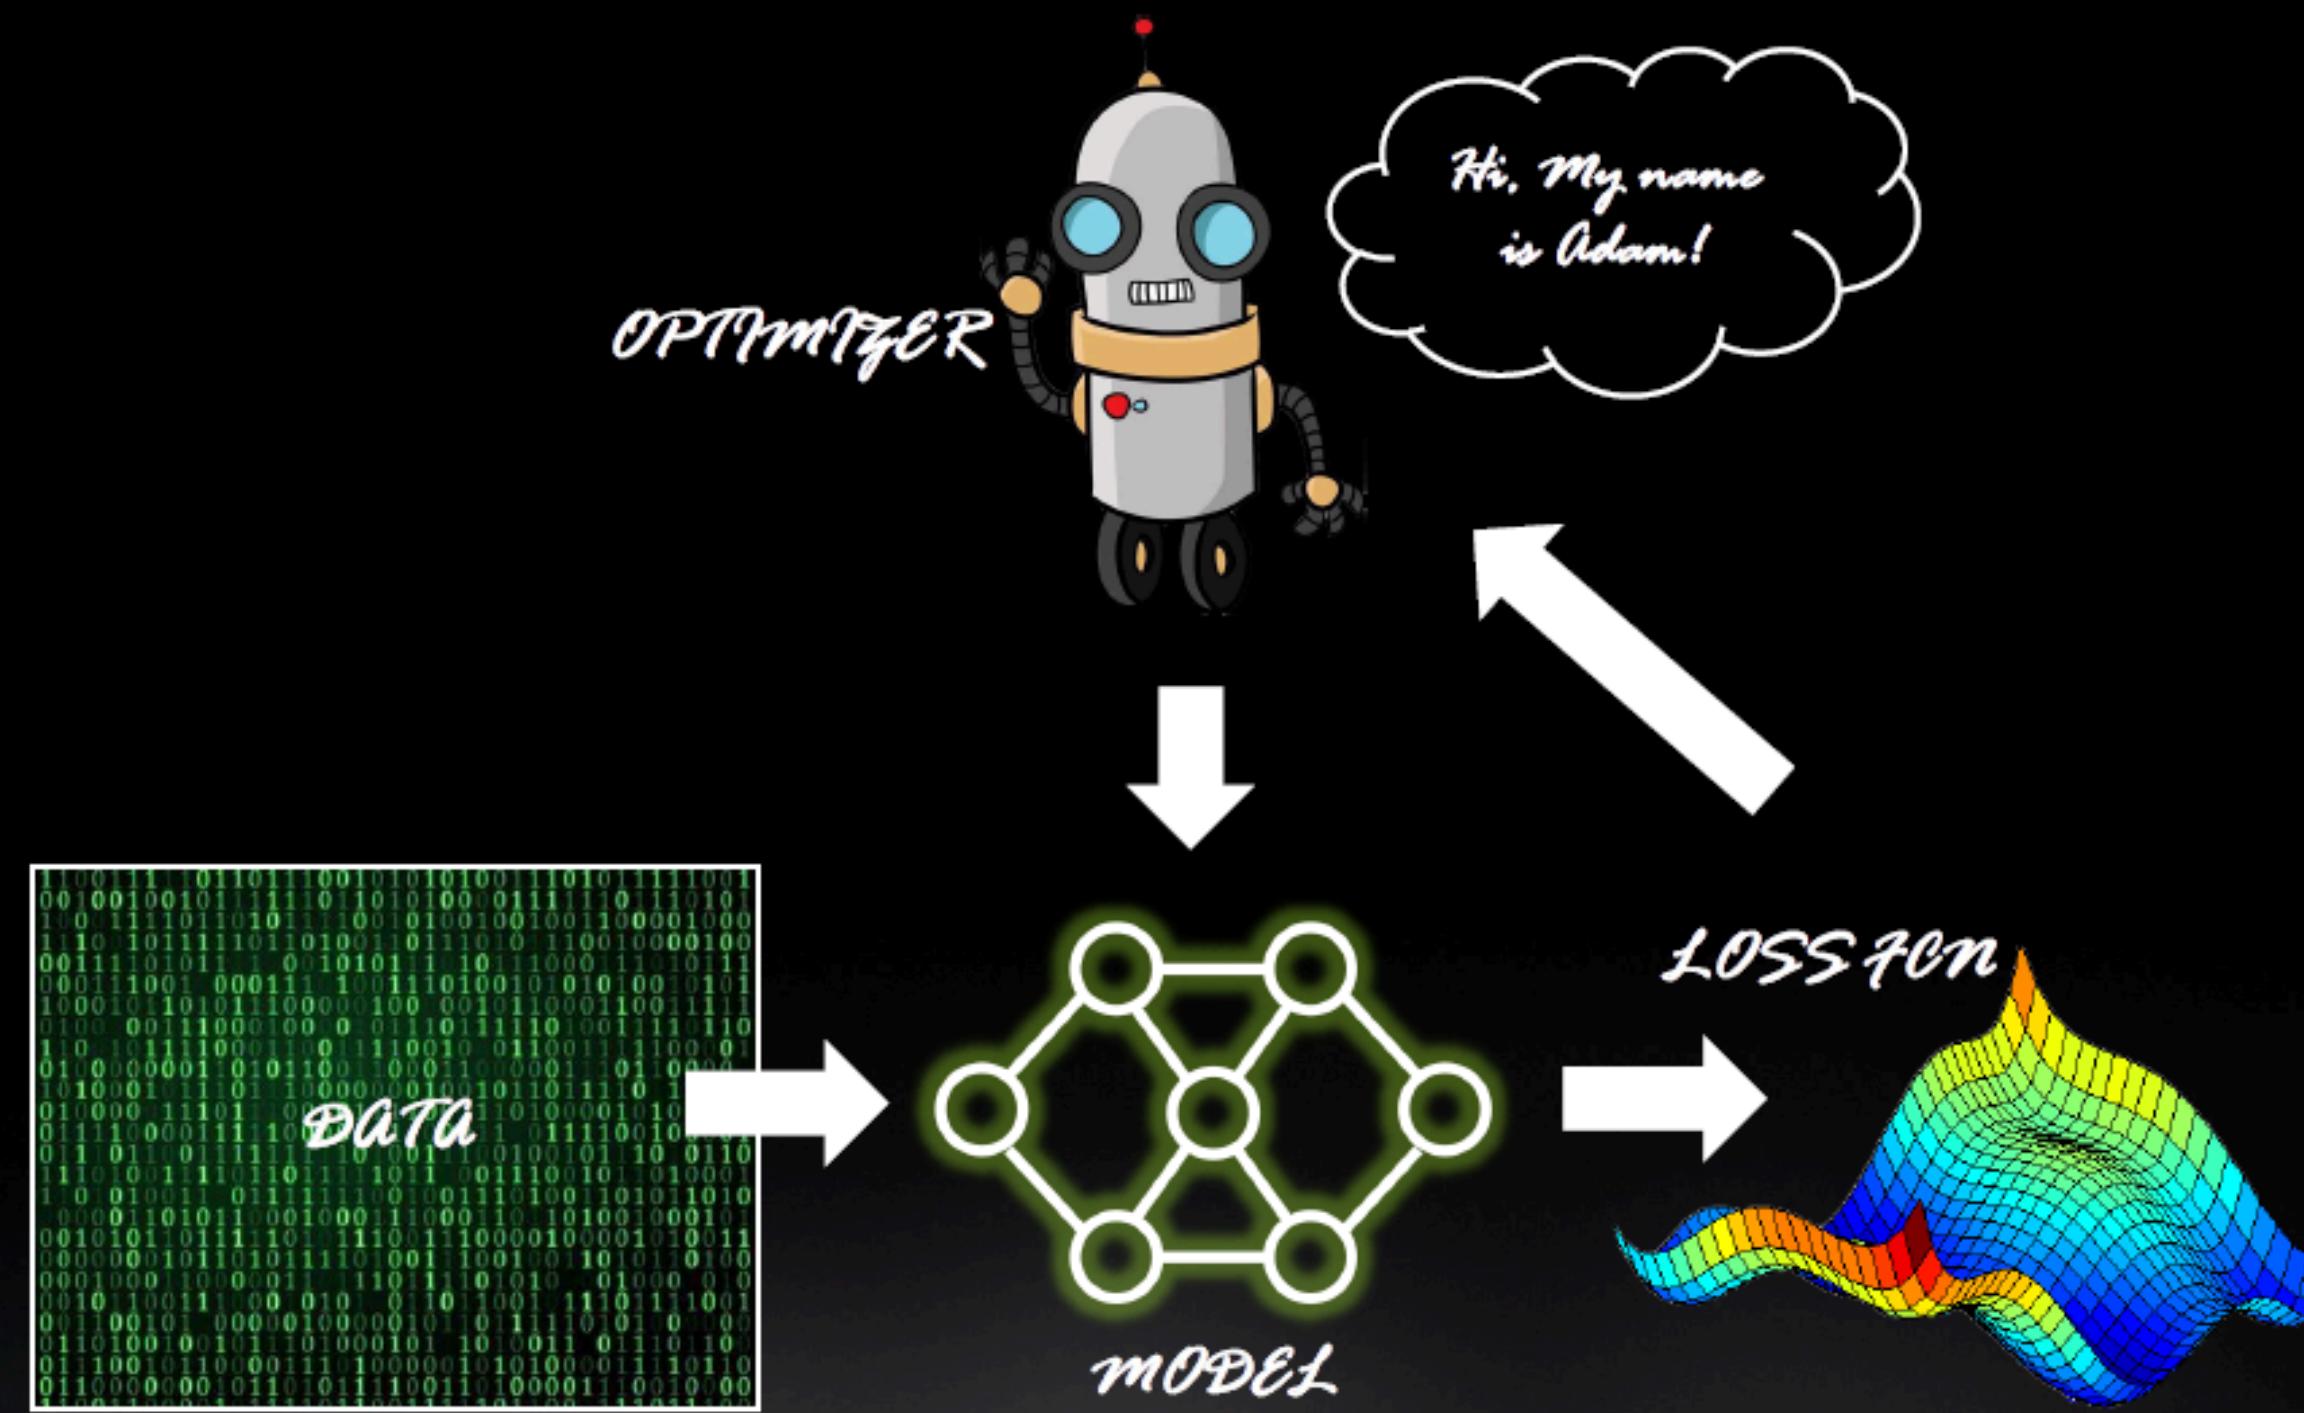
<https://towardsdatascience.com/the-differences-between-artificial-and-biological-neural-networks-a8b46db828b7>

KEEP TRAINING AND TEST DATA SEPERATE

KEEP TEST, TRAINING, AND VALIDATION DATA SEPERATE



TRAINING DATA, MODEL, LOSS, AND OPTIMIZER



WHAT YOU NEED TO MAKE IT WORK

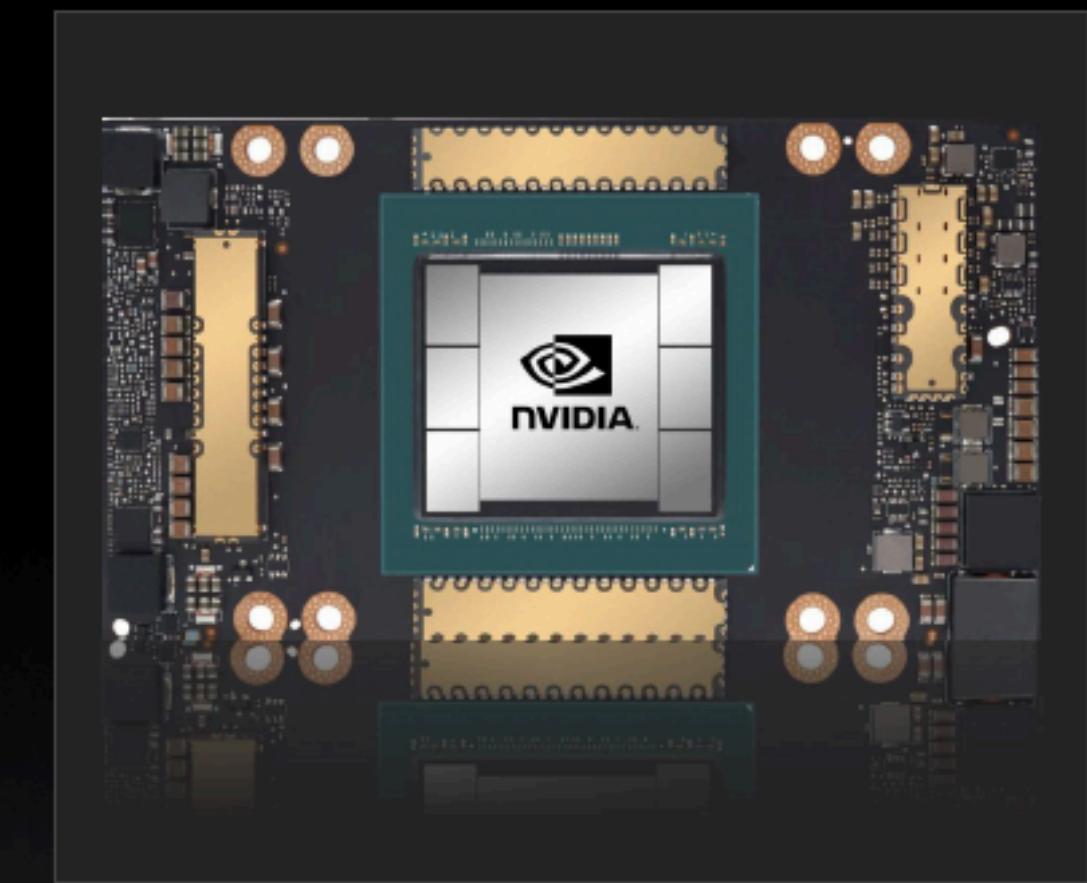
You need three main ingredients (and some skill)



LARGE QUANTITIES OF DATA

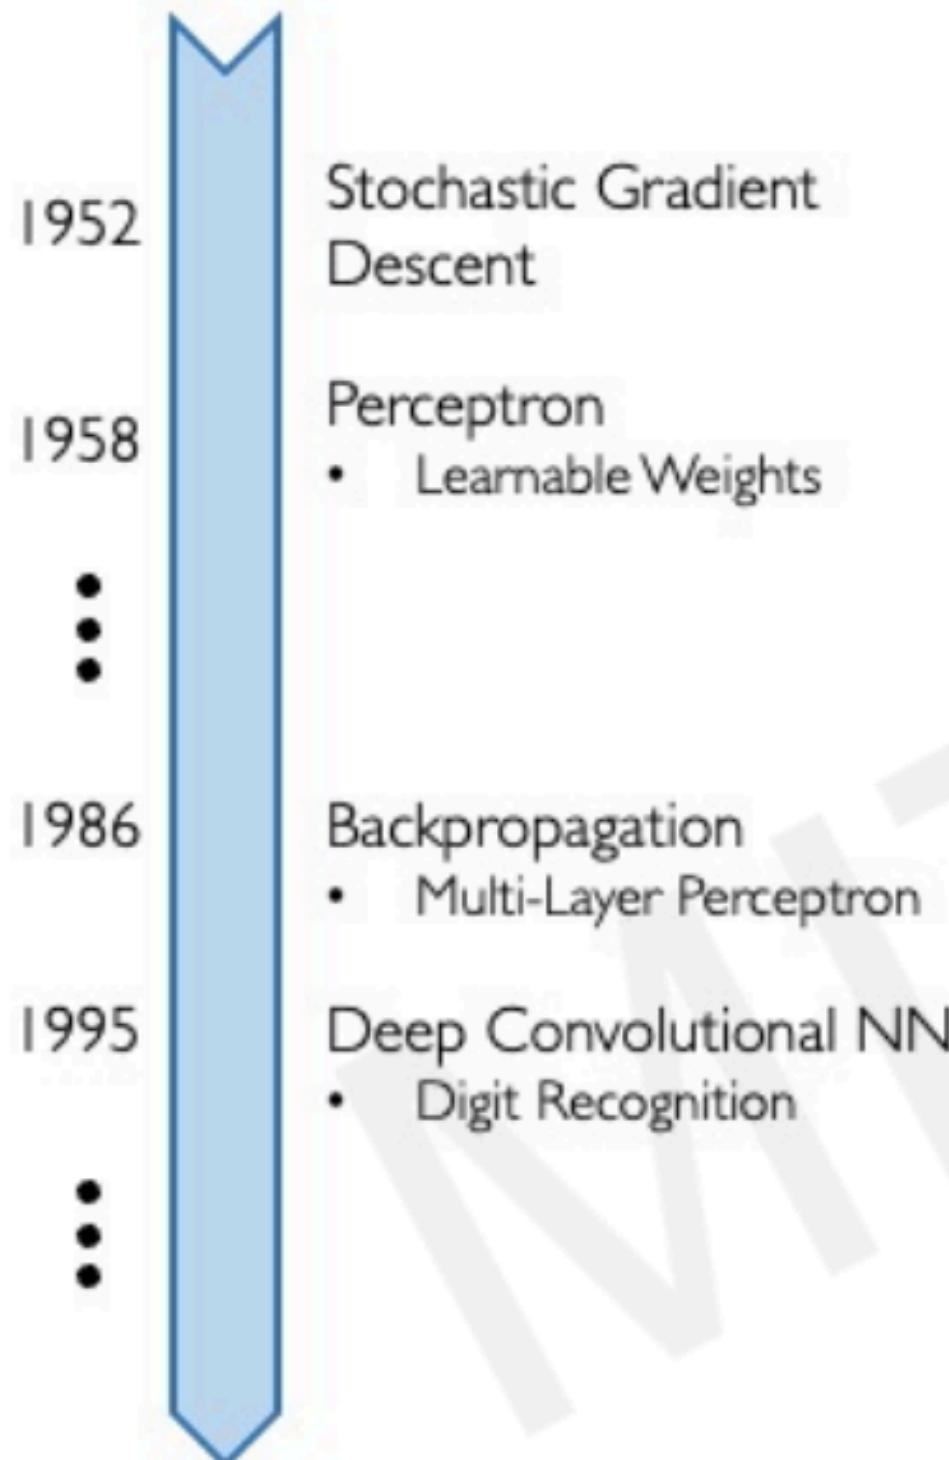


ML FRAMEWORK



GPU ACCELERATOR

Why Now?

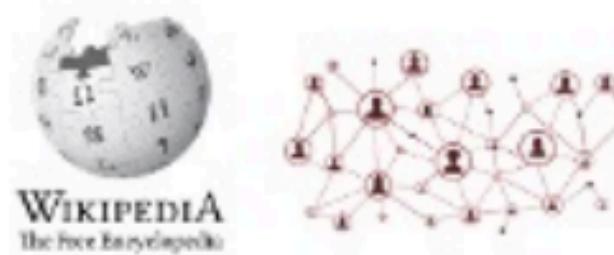


Neural Networks date back decades, so why the dominance?

I. Big Data

- Larger Datasets
- Easier Collection & Storage

IMAGENET



2. Hardware

- Graphics Processing Units (GPUs)
- Massively Parallelizable

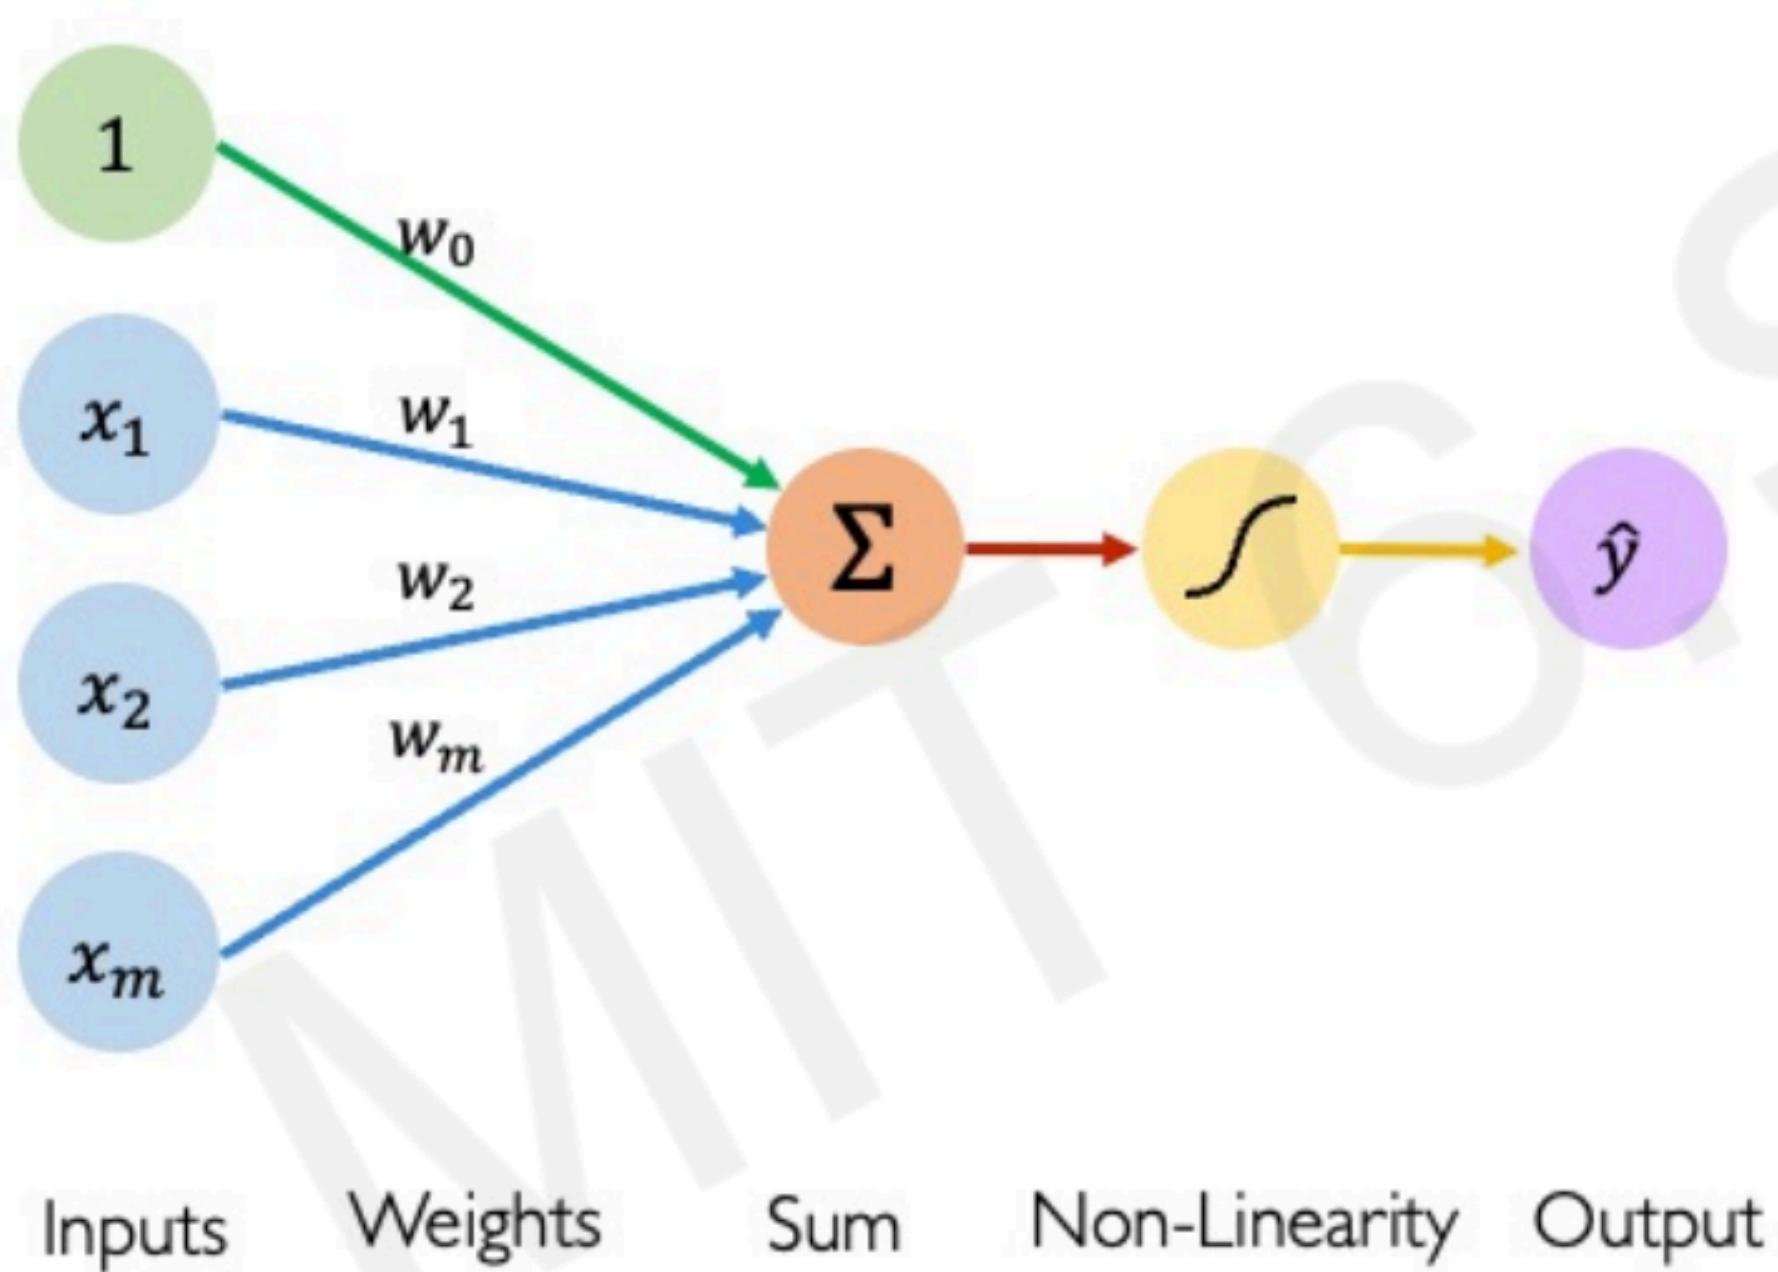


3. Software

- Improved Techniques
- New Models
- Toolboxes



The Perceptron: Forward Propagation

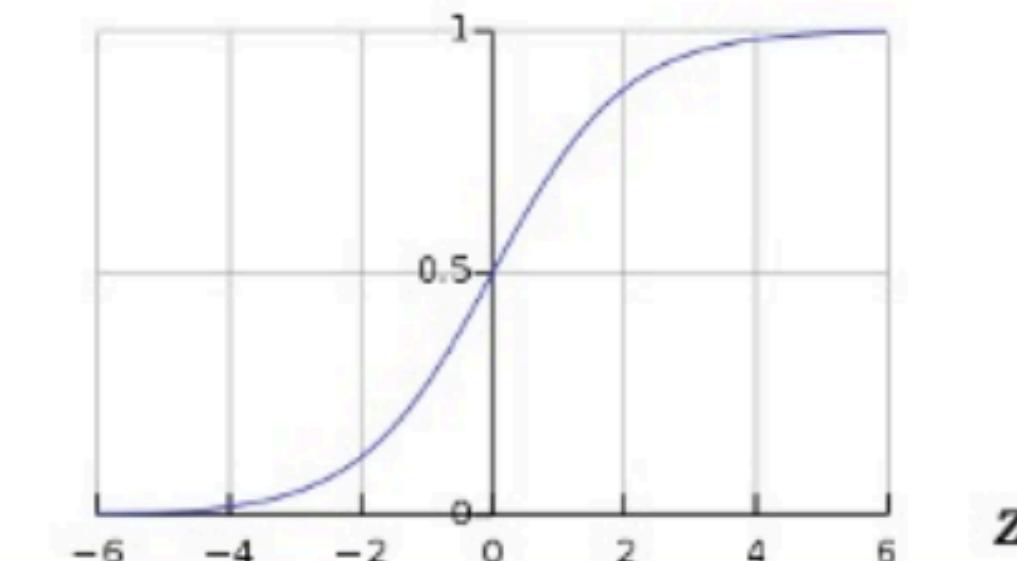


Activation Functions

$$\hat{y} = g(w_0 + \mathbf{X}^T \mathbf{W})$$

- Example: sigmoid function

$$g(z) = \sigma(z) = \frac{1}{1 + e^{-z}}$$



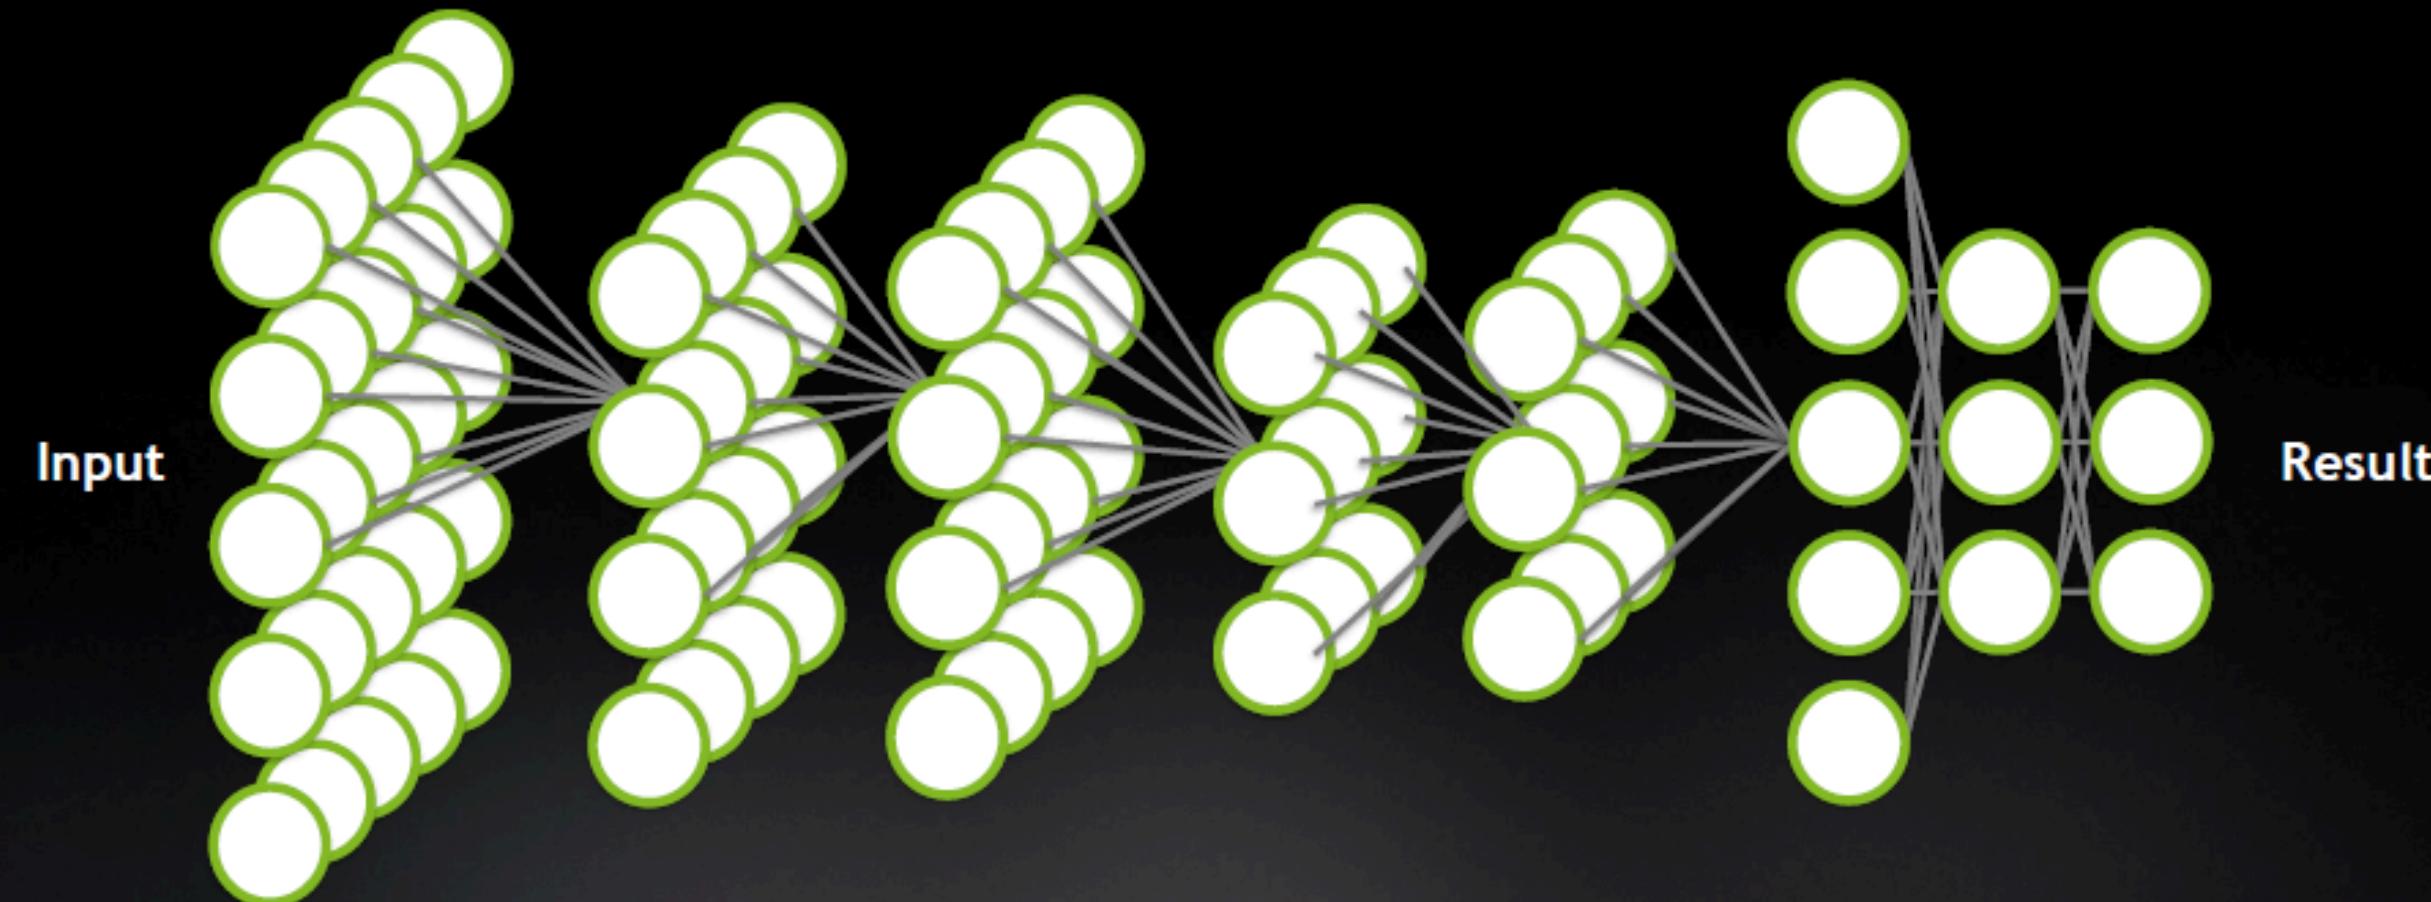
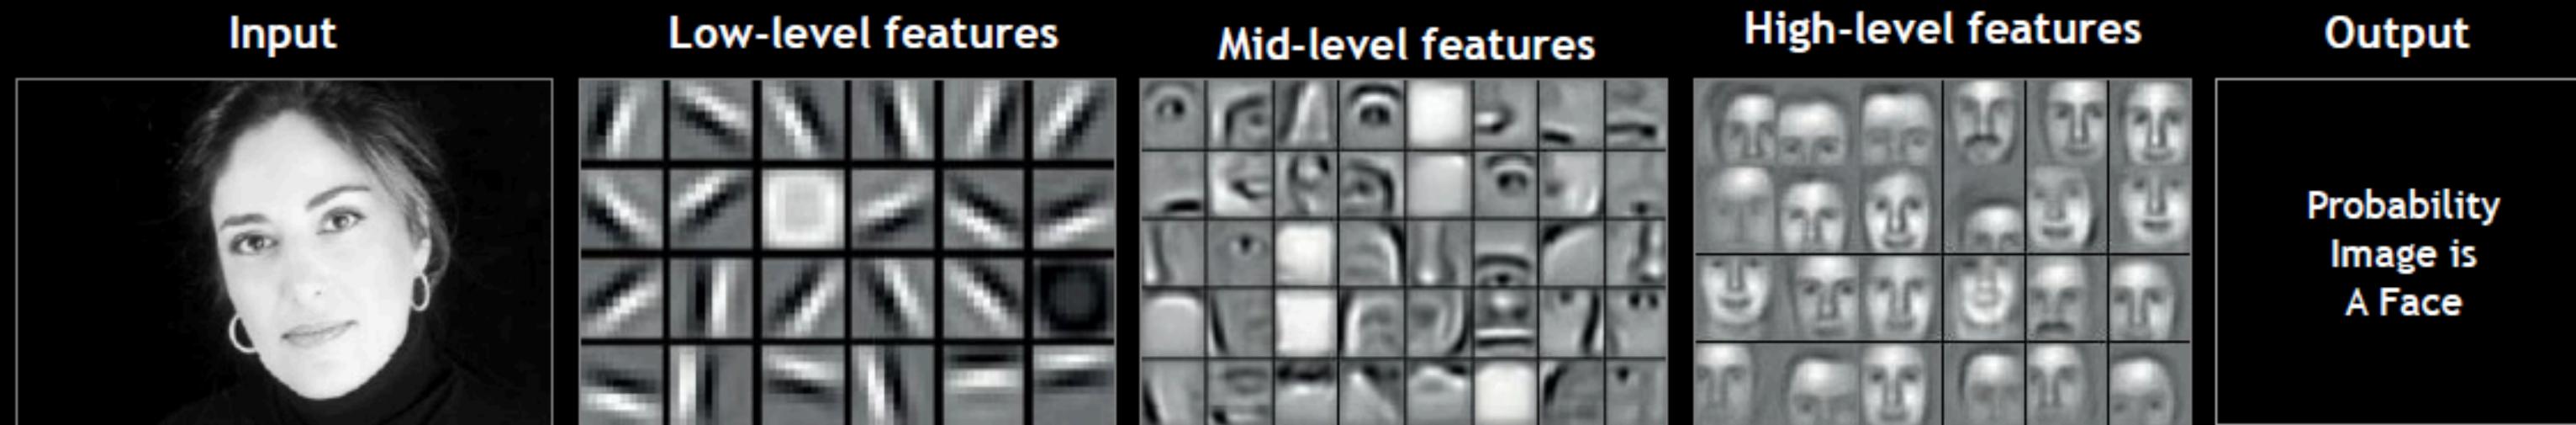
COMPUTER VISION

FCNN & CNN

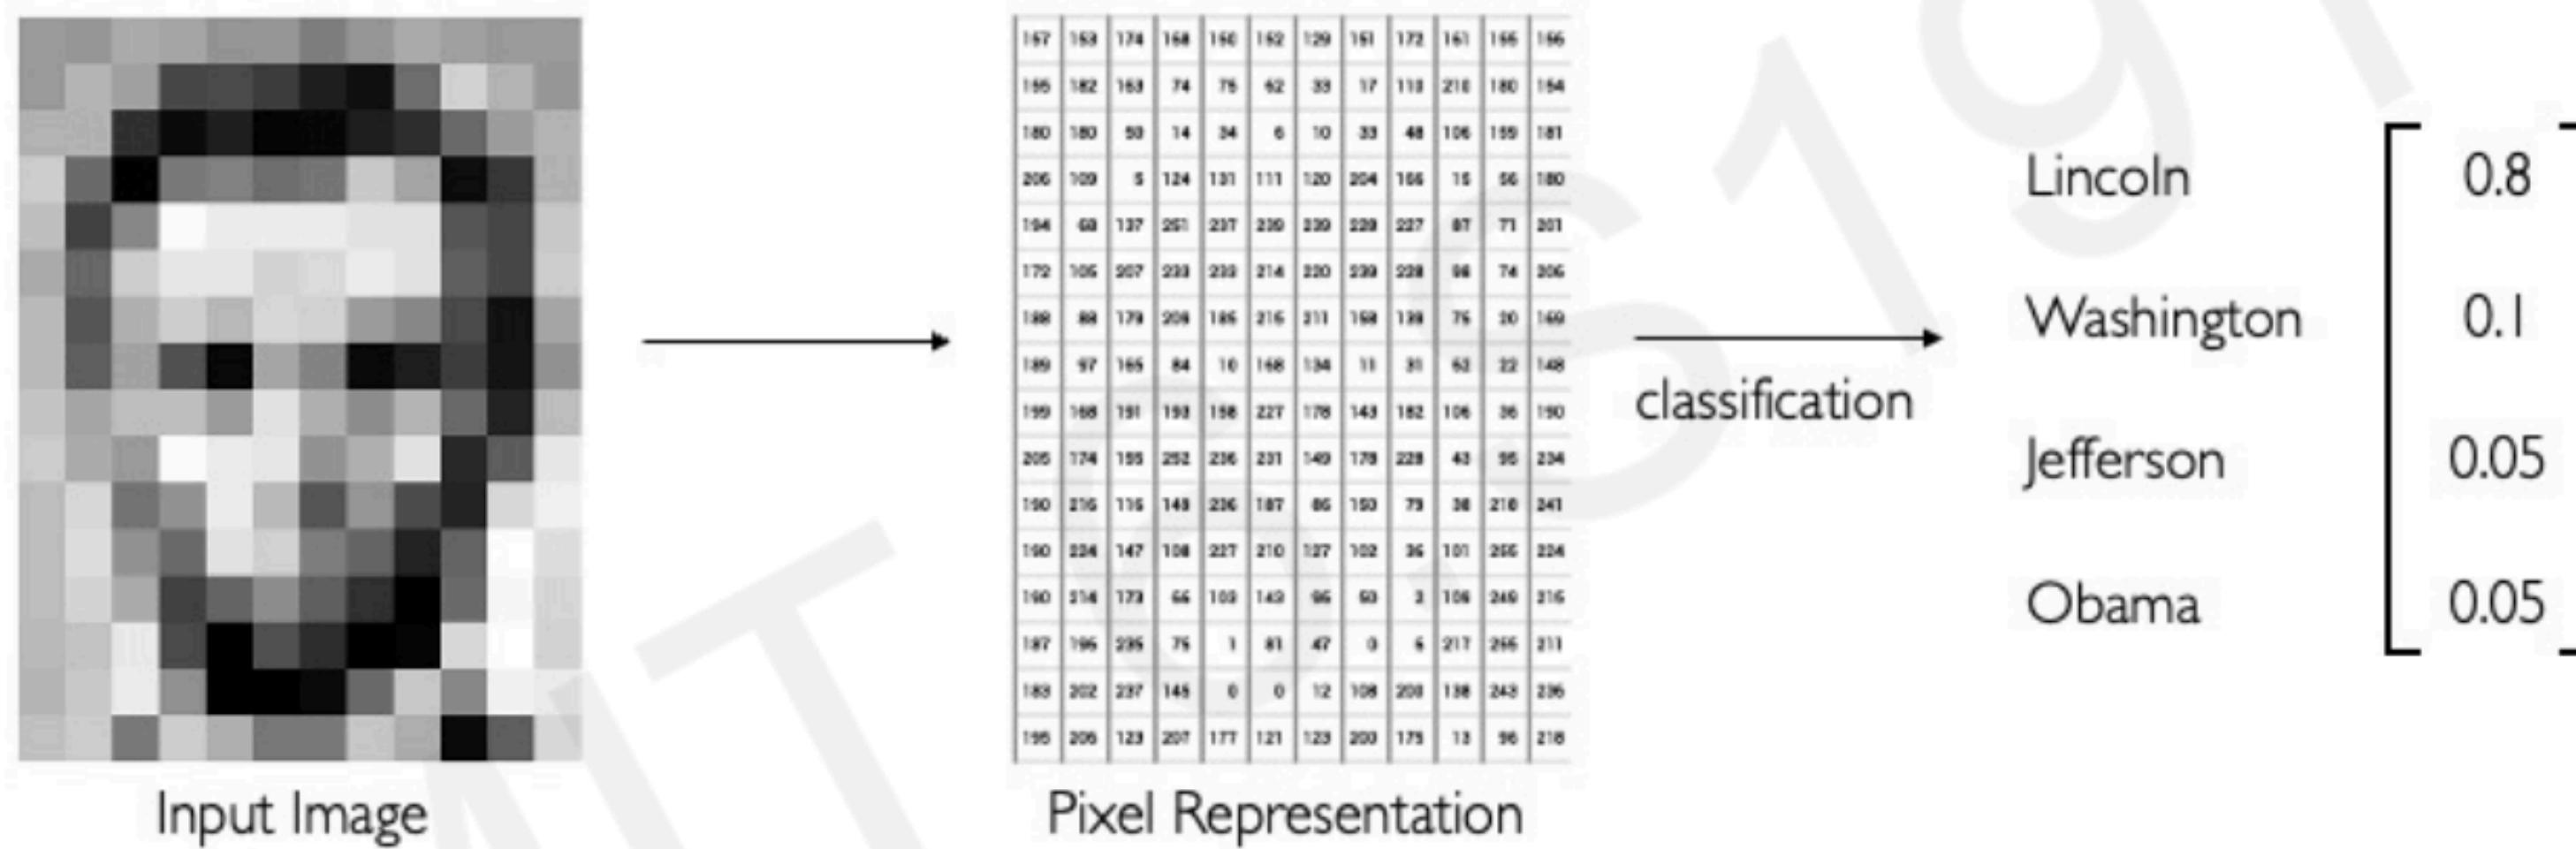


DEEPER NEURAL NETWORKS

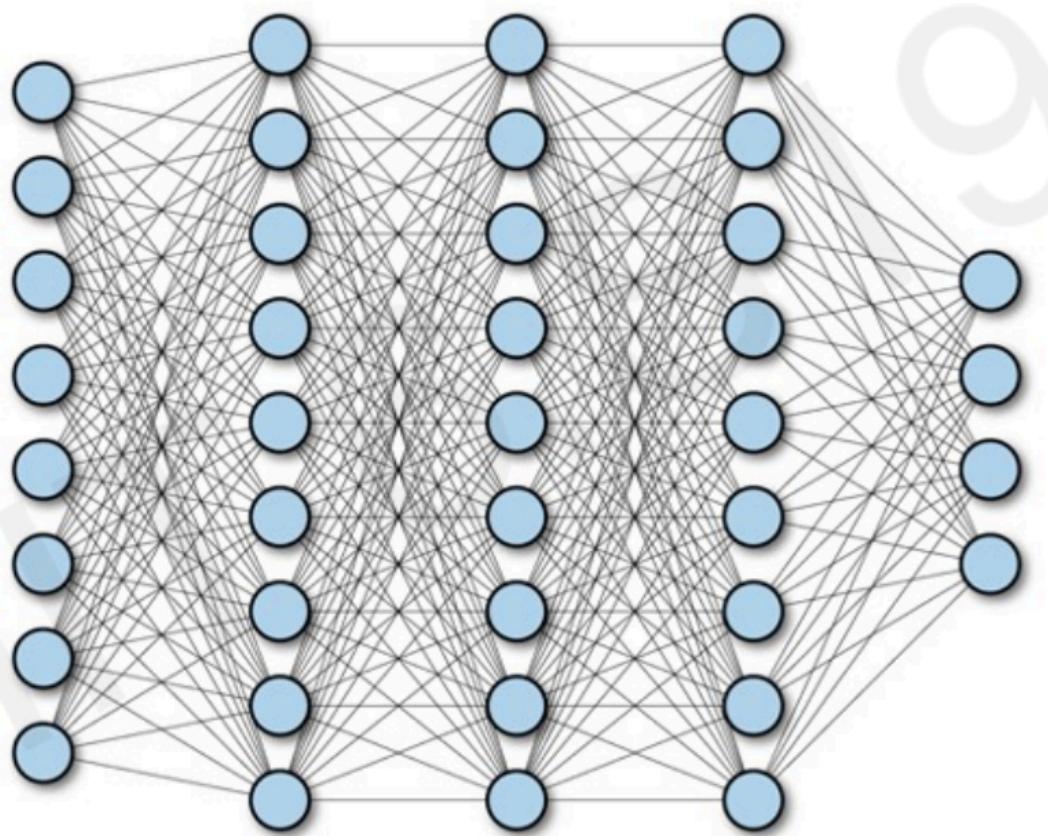
More layers allows for more levels of abstraction



Tasks in Computer Vision

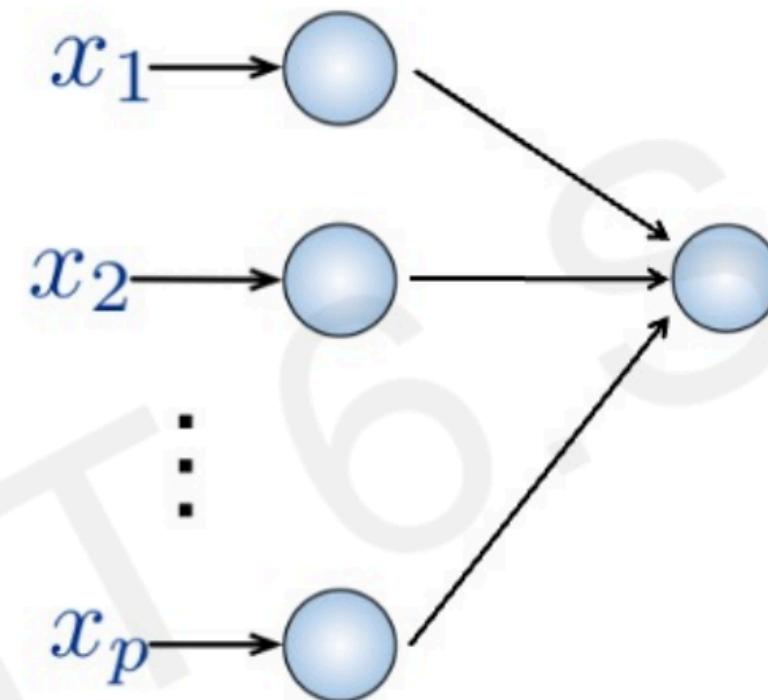


Fully Connected Neural Network



Input:

- 2D image
- Vector of pixel values



Fully Connected:

- Connect neuron in hidden layer to all neurons in input layer
- No spatial information!
- And many, many parameters!

How can we use **spatial structure** in the input to inform the architecture of the network?

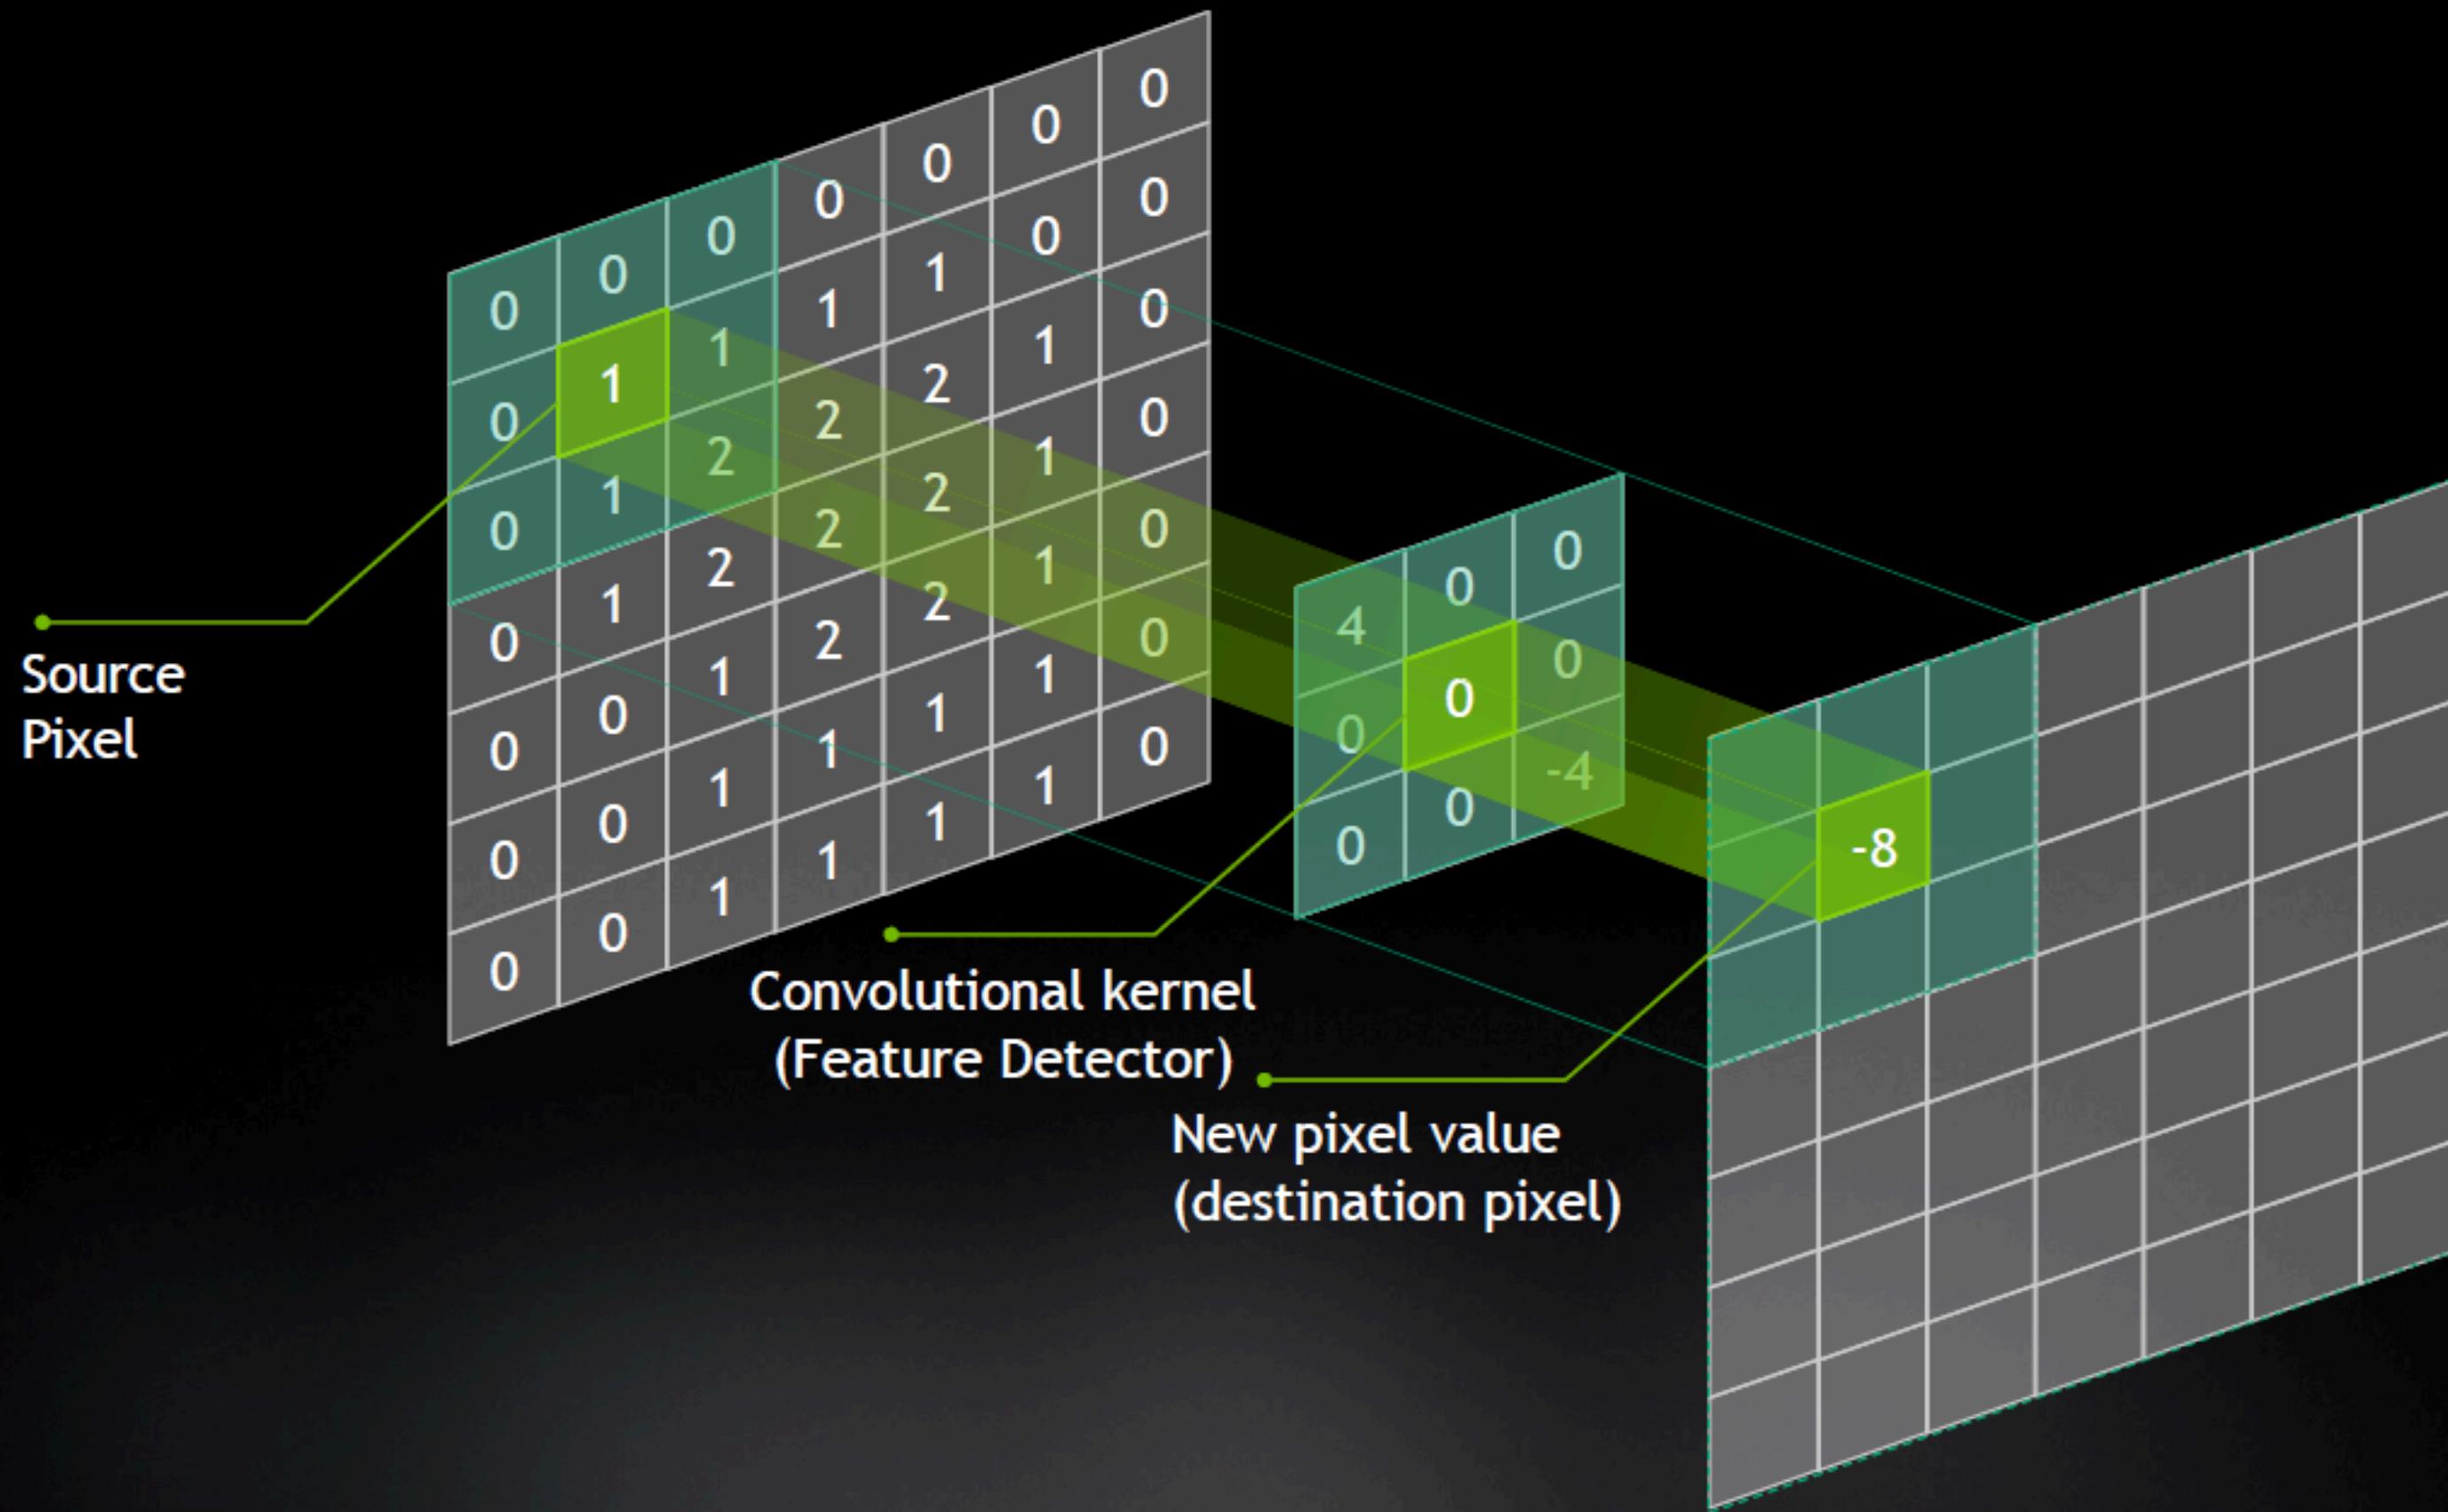
THE
CONVOLUTION



A dark, circular background featuring concentric circles of varying shades of gray. A horizontal dotted line is positioned at the bottom center of the image.

WHAT IS A CONVOLUTION?

A small matrix transformation, applied at each point of the image

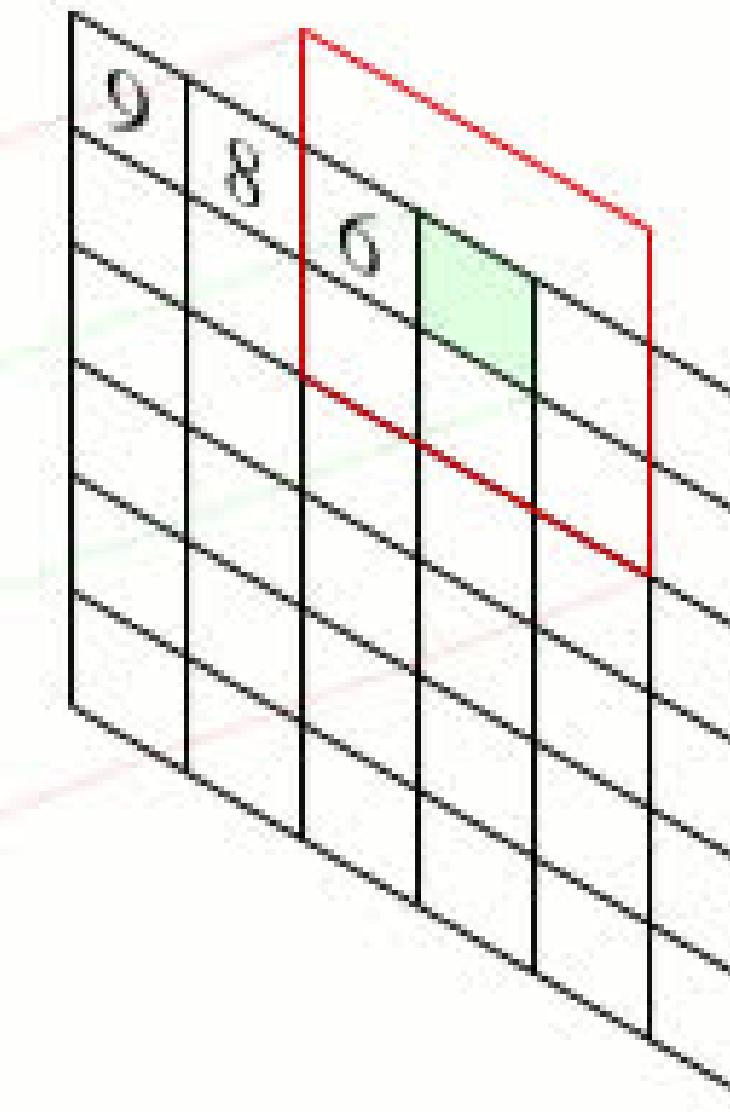


output

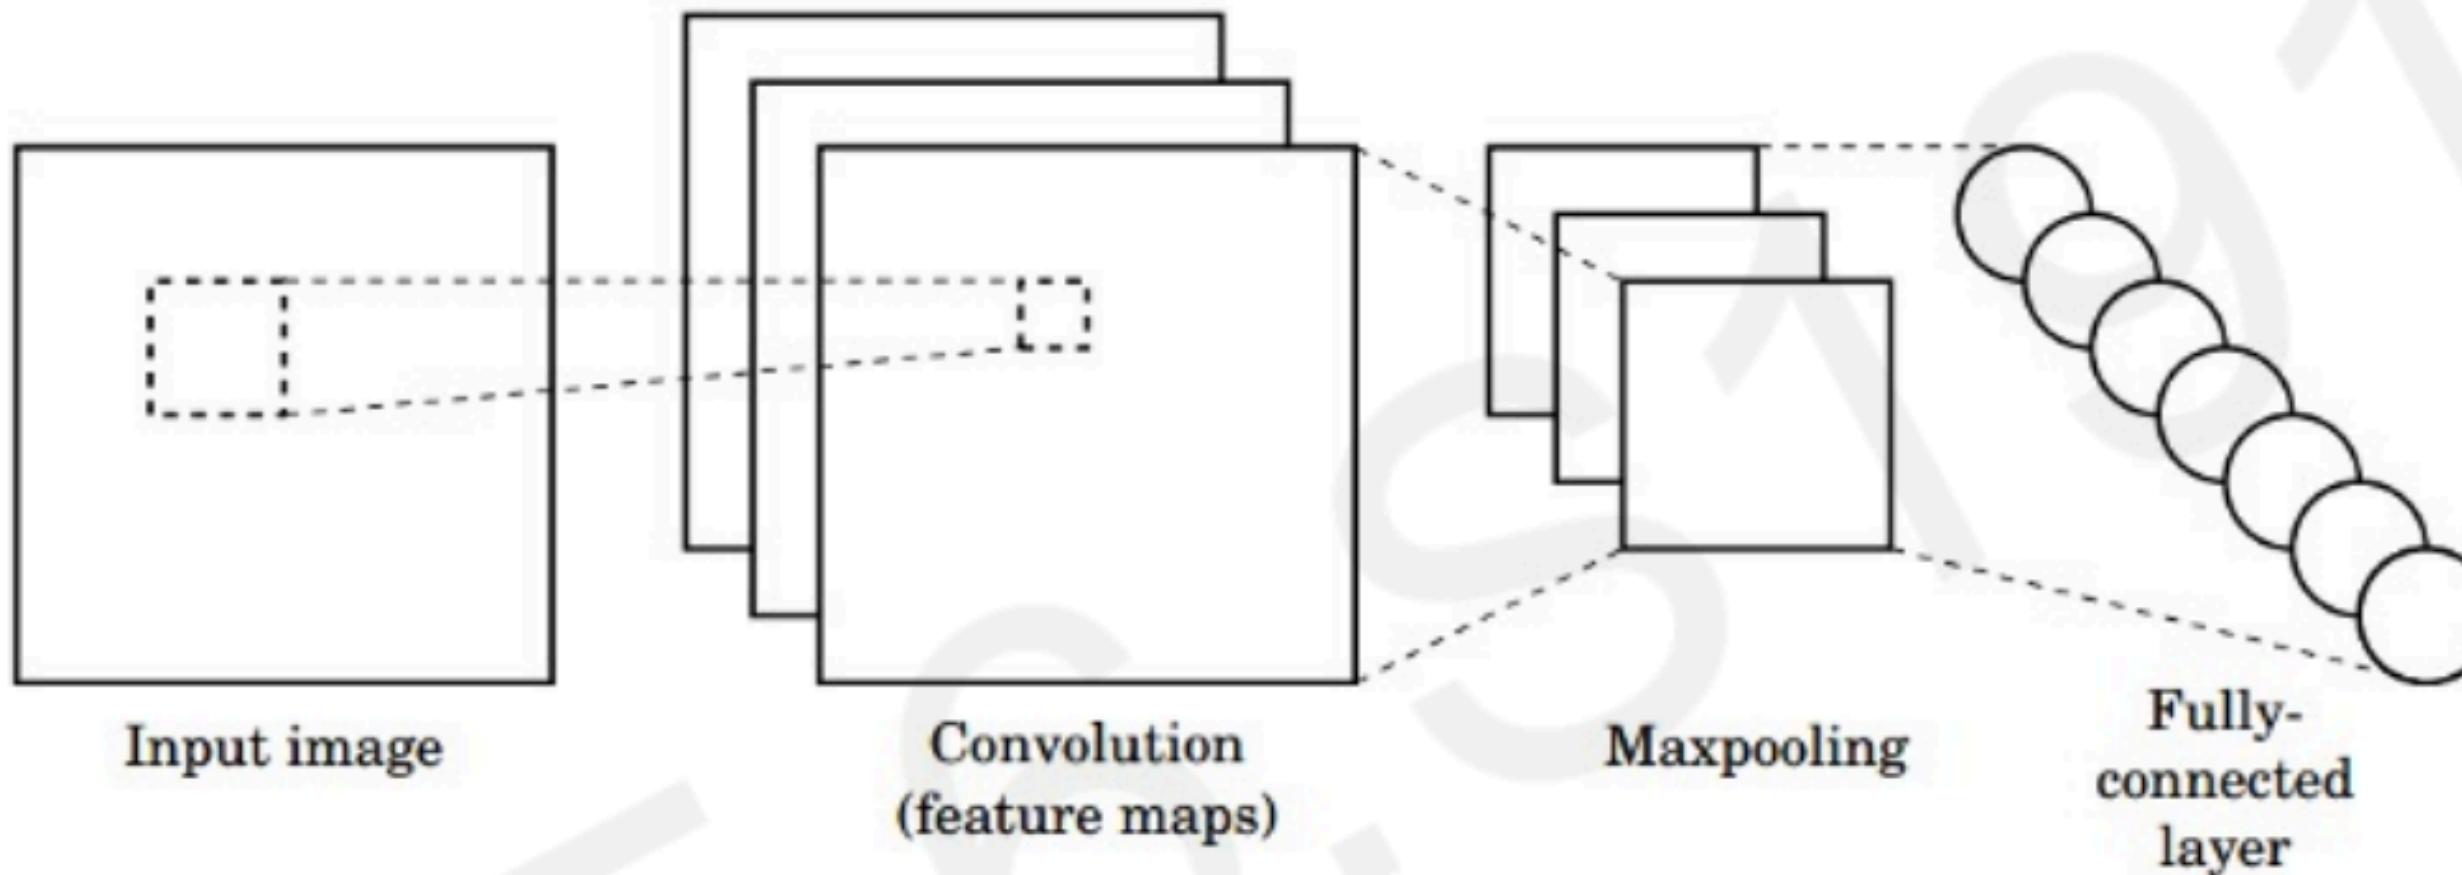
| | | | | | | | | | |
|---|---|---|---|---|---|---|---|---|---|
| 7 | 0 | 1 | 0 | 0 | 1 | 0 | 0 | 0 | 0 |
| 0 | 4 | 3 | 3 | 4 | 0 | 0 | 7 | 0 | 0 |
| 6 | 0 | 4 | 3 | 3 | 2 | 2 | 0 | 0 | 0 |
| 3 | 1 | 1 | 1 | 1 | 1 | 1 | 0 | 0 | 0 |
| 2 | 0 | 1 | 1 | 1 | 1 | 1 | 0 | 0 | 0 |
| 5 | 0 | 1 | 1 | 1 | 1 | 1 | 0 | 0 | 0 |
| 4 | 0 | 1 | 1 | 1 | 1 | 1 | 0 | 0 | 0 |
| 3 | 0 | 1 | 1 | 1 | 1 | 1 | 0 | 0 | 0 |
| 2 | 0 | 1 | 1 | 1 | 1 | 1 | 0 | 0 | 0 |
| 1 | 0 | 1 | 1 | 1 | 1 | 1 | 0 | 0 | 0 |

input

| | | | | |
|----|----|---|----|---|
| 0 | -1 | 0 | 1 | 0 |
| -1 | 0 | 5 | -1 | 0 |
| 0 | -1 | 0 | 1 | 0 |
| -1 | 0 | 1 | 0 | 0 |



CNNs for Classification



```
tf.keras.layers.Conv2D
```

```
tf.keras.activations.*
```

```
tf.keras.layers.MaxPool2D
```

1. Convolution: Apply filters to generate feature maps.

2. Non-linearity: Often ReLU.

3. Pooling: Downsampling operation on each feature map.



```
torch.nn.Conv2d
```

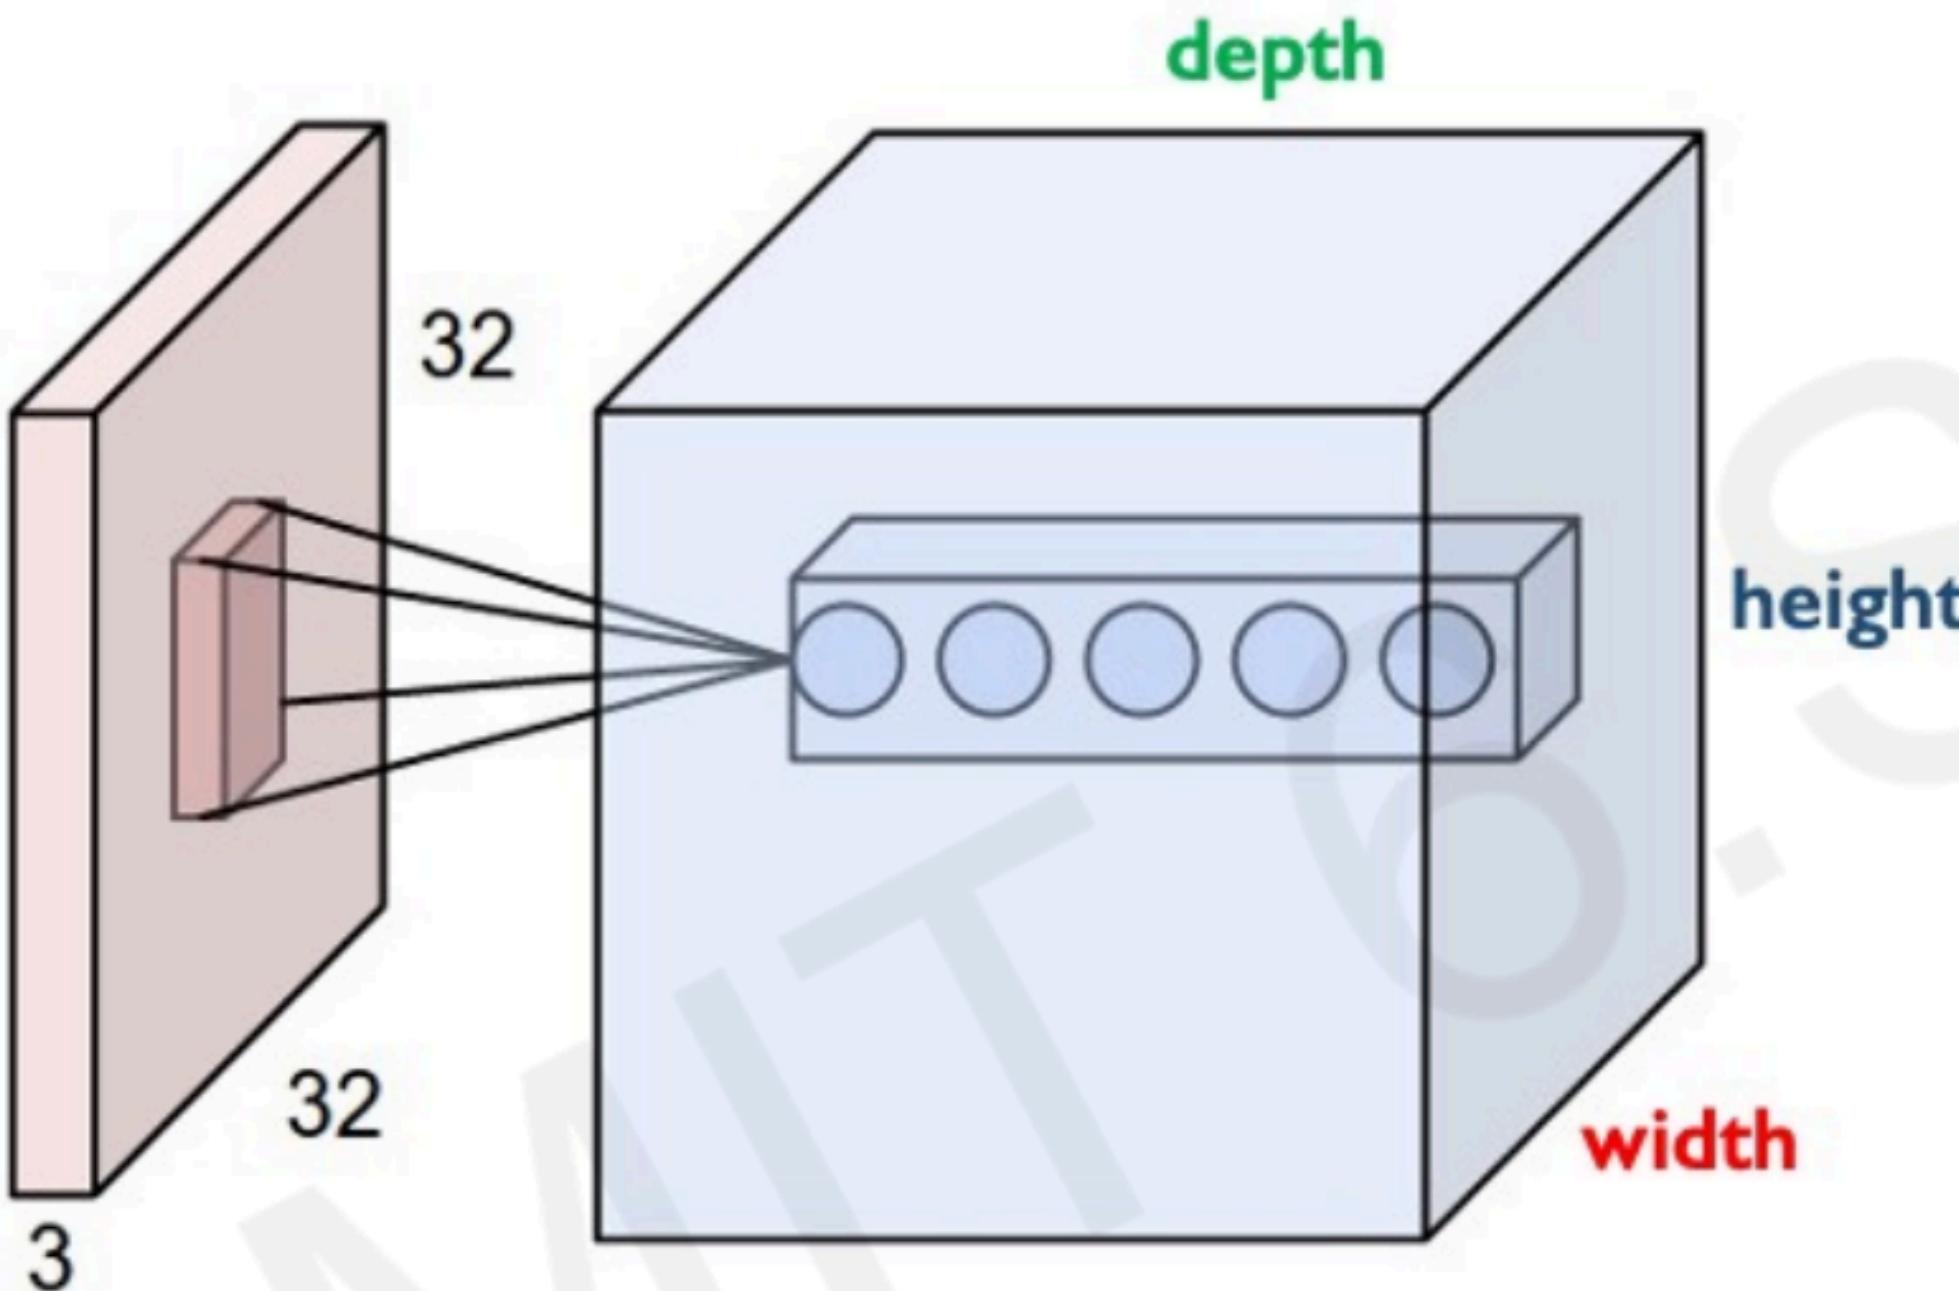
```
torch.nn.ReLU,...
```

```
torch.nn.MaxPool2d
```

Train model with image data.

Learn weights of filters in convolutional layers.

CNNs: Spatial Arrangement of Output Volume



Layer Dimensions:

$h \times w \times d$

where h and w are spatial dimensions
d (depth) = number of filters

Stride:

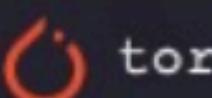
Filter step size

Receptive Field:

Locations in input image that
a node is path connected to



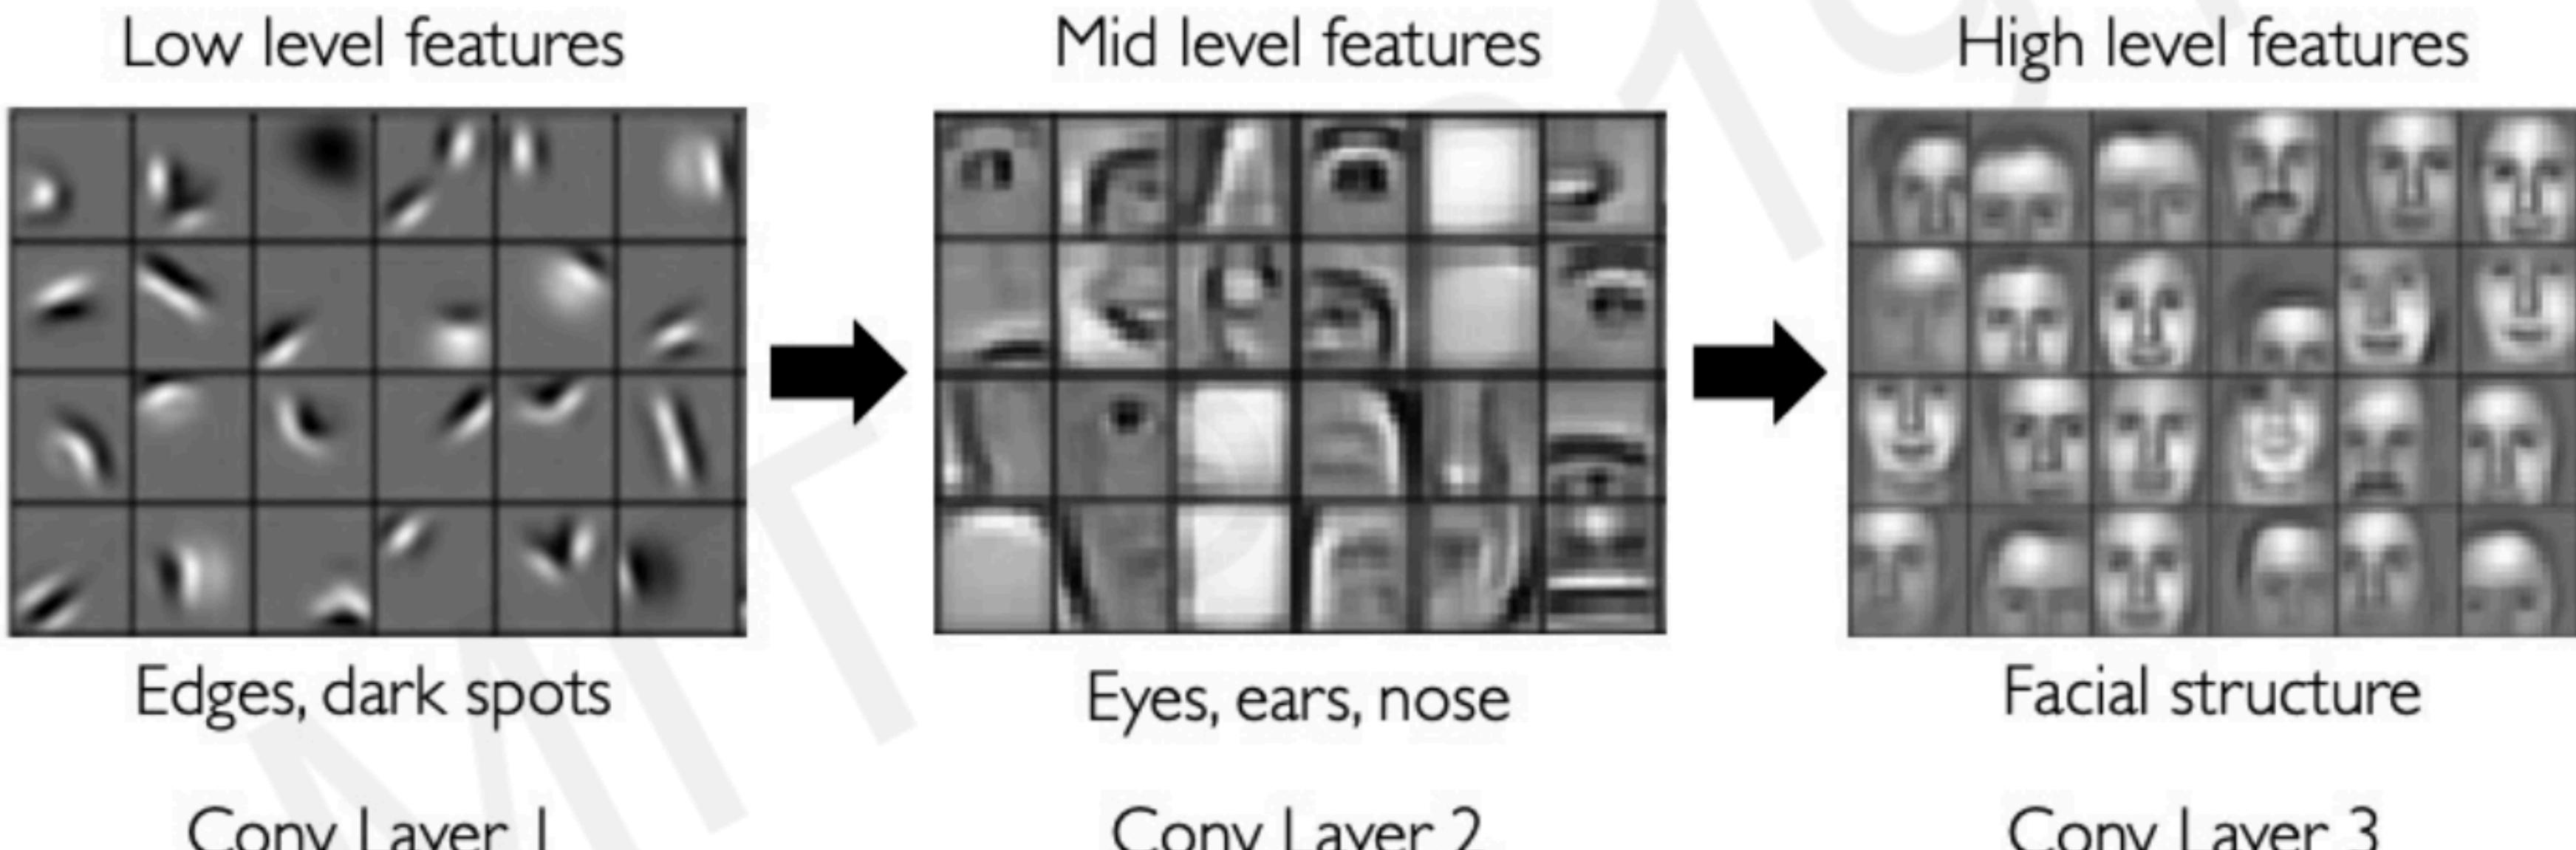
```
tf.keras.layers.Conv2D( filters=d, kernel_size=(h,w), strides=s )
```



```
torch.nn.Conv2d( in_channels=3, out_channels=d, kernel_size=(h,w), stride=s )
```

Need to specify
input dimensionality!

Representation Learning in Deep CNNs





```
model_cnn = Sequential([
    # 32 filtres pour apprendre 32 motifs différents.
    Conv2D(32, kernel_size=(3, 3), activation='relu', input_shape=input_shape),

    # On résume l'information
    MaxPooling2D(pool_size=(2, 2)),

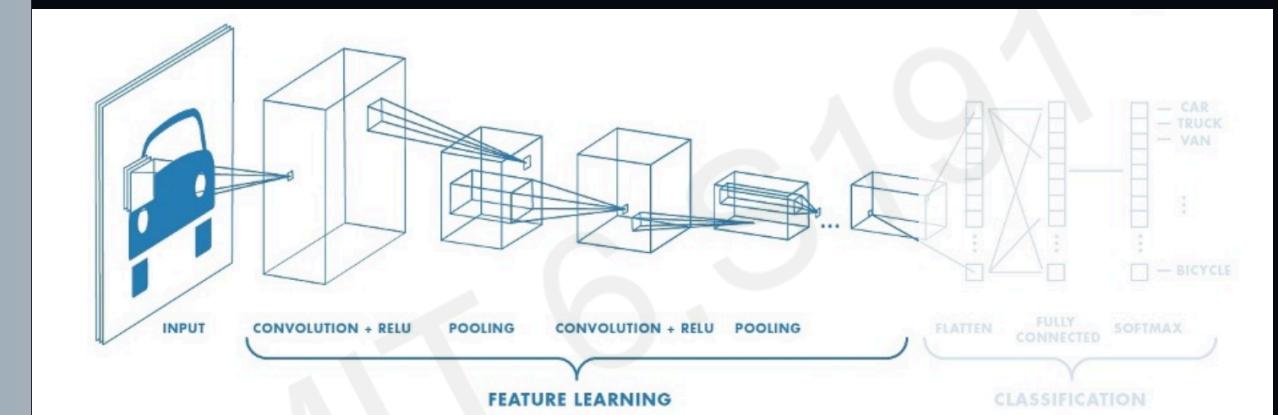
    # On utilise 64 filtres pour capturer plus de complexité.
    Conv2D(64, kernel_size=(3, 3), activation='relu'),
    MaxPooling2D(pool_size=(2, 2)),

    # On transforme notre carte 2D de features en un vecteur 1D.
    Flatten(),

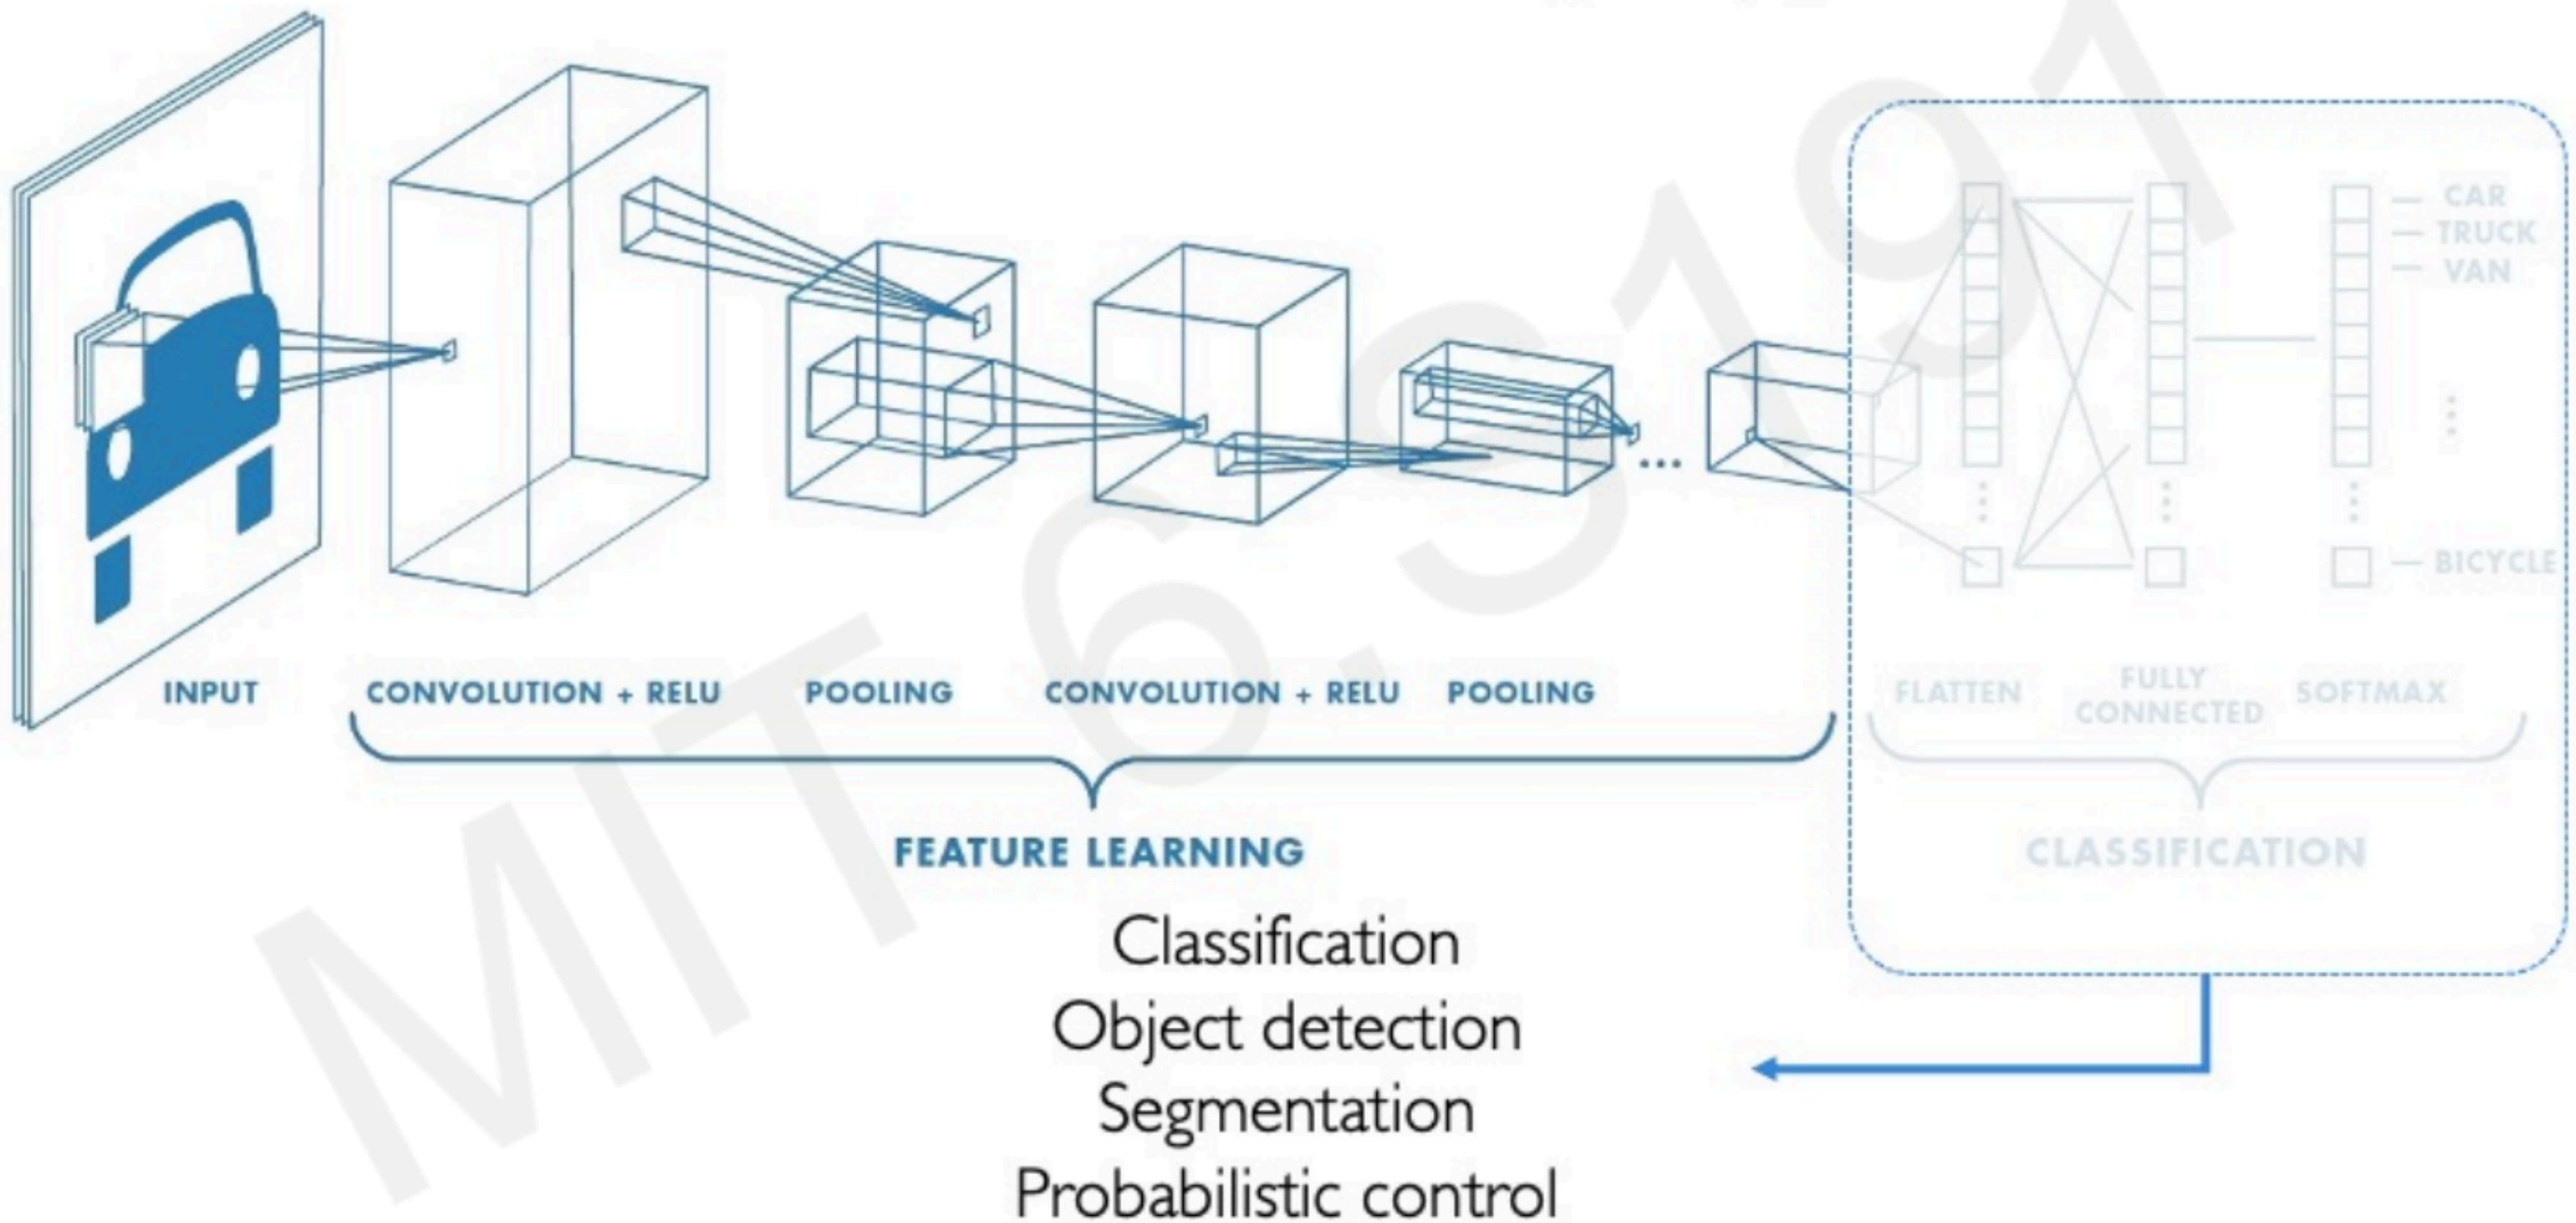
    # Il "éteint" aléatoirement 50% des neurones pendant l'entraînement.
    Dropout(0.5),

    # 'softmax' transforme les sorties en une distribution de probabilité.
    Dense(10, activation='softmax')
])

# Affichons le plan de notre architecture
model_cnn.summary()
```



An Architecture for Many Applications



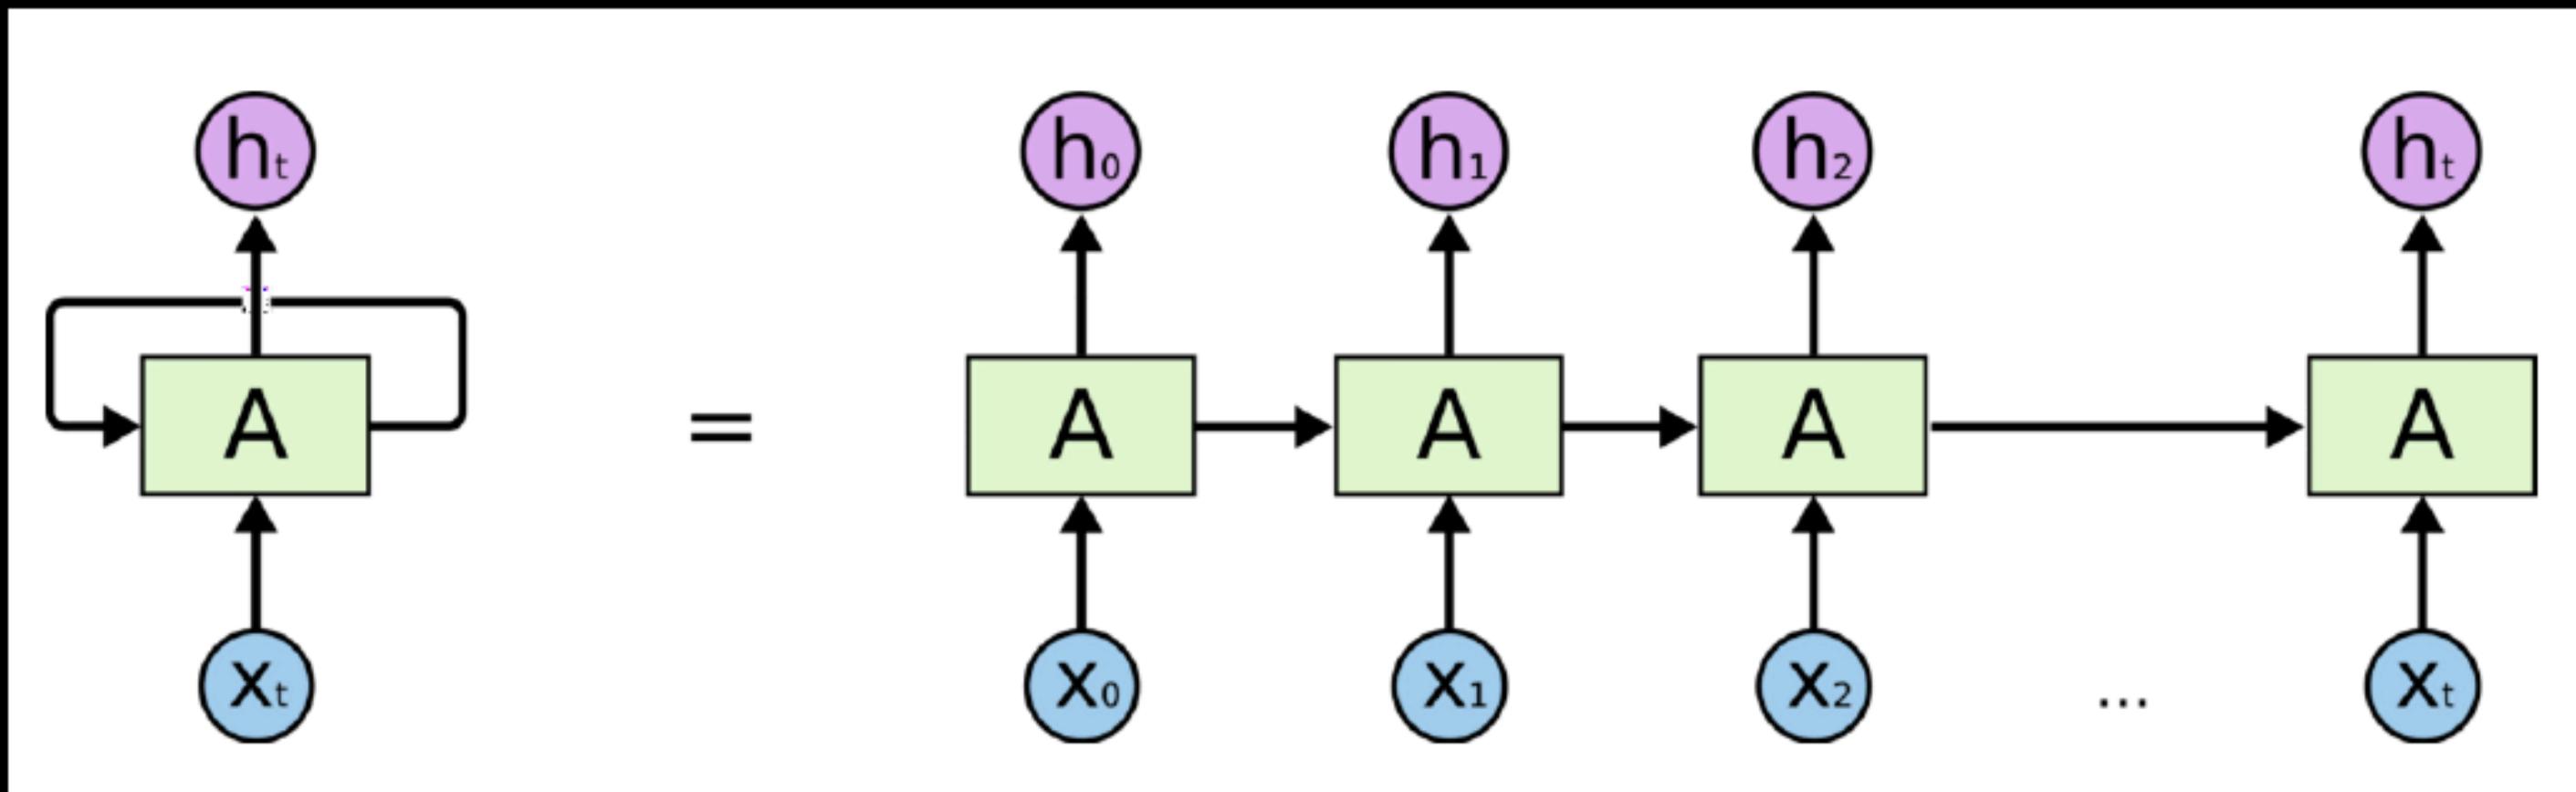
DEEP SEQUENCE MODELING

RNN, LSTM & GRU



RECURRENT NEURAL NETS

(1986) Neural networks for sequences

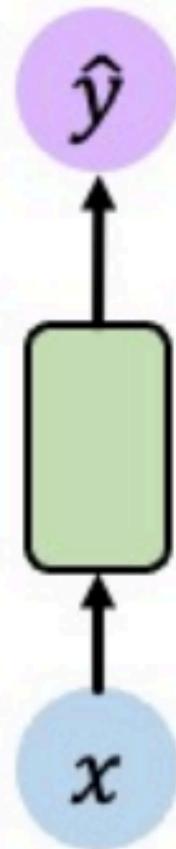


RNN, Unrolled over time

<https://stanford.edu/~shervine/teaching/cs-230/cheatsheet-recurrent-neural-networks>

<http://www.wildml.com/2015/09/recurrent-neural-networks-tutorial-part-1-introduction-to-rnns/>

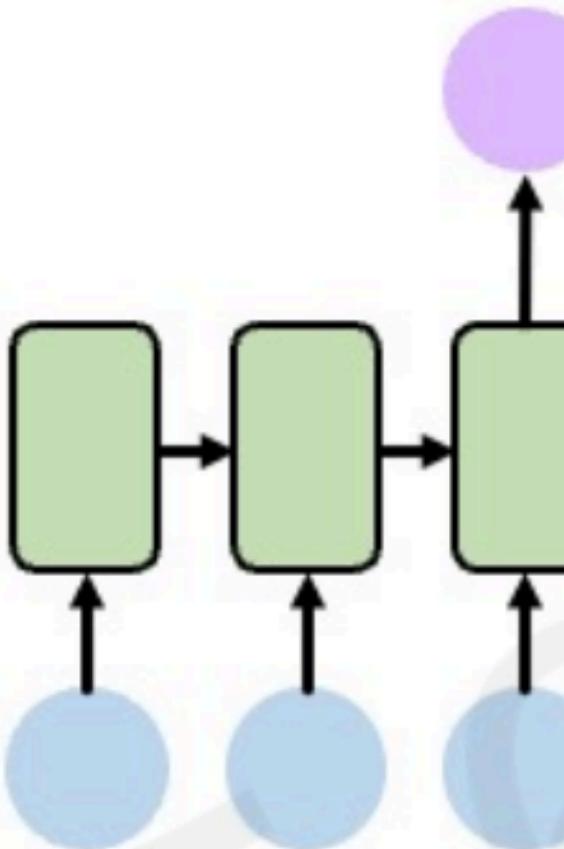
Sequence Modeling Applications



One to One
Binary Classification

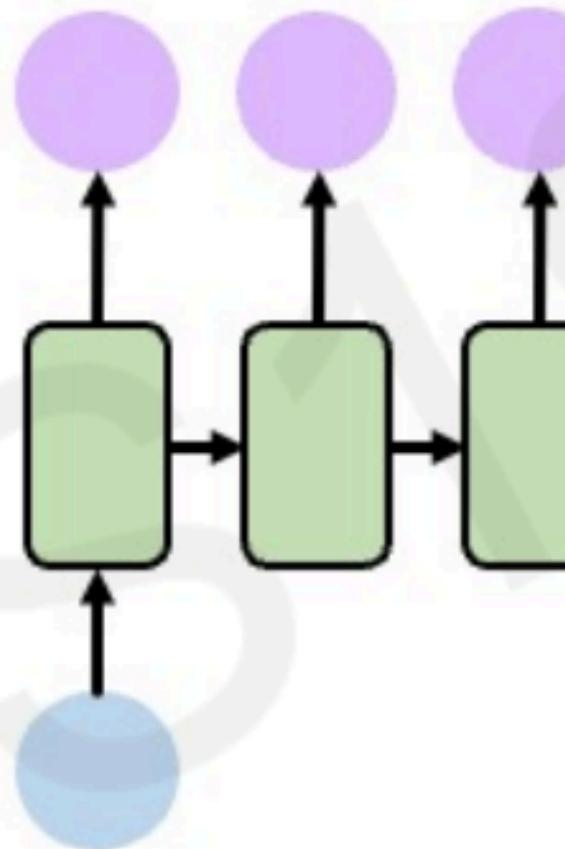


"Will I pass this class?"
Student → Pass?



Many to One
Sentiment Classification

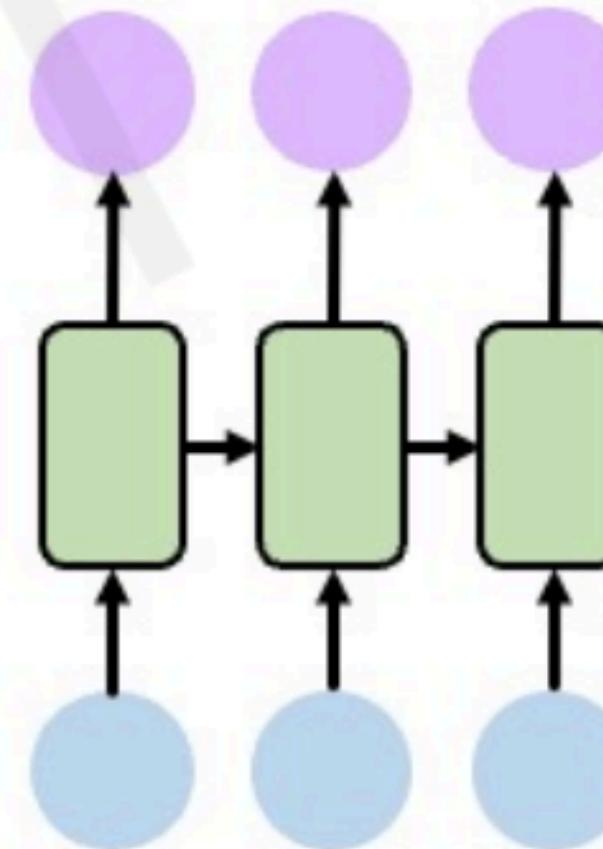
Ivar Hagendoorn
@IvarHagendoorn
Follow
The @MIT Introduction to #DeepLearning is definitely one of the best courses of its kind currently available online introtodeeplearning.com
12:45 PM - 12 Feb 2018



One to Many
Image Captioning



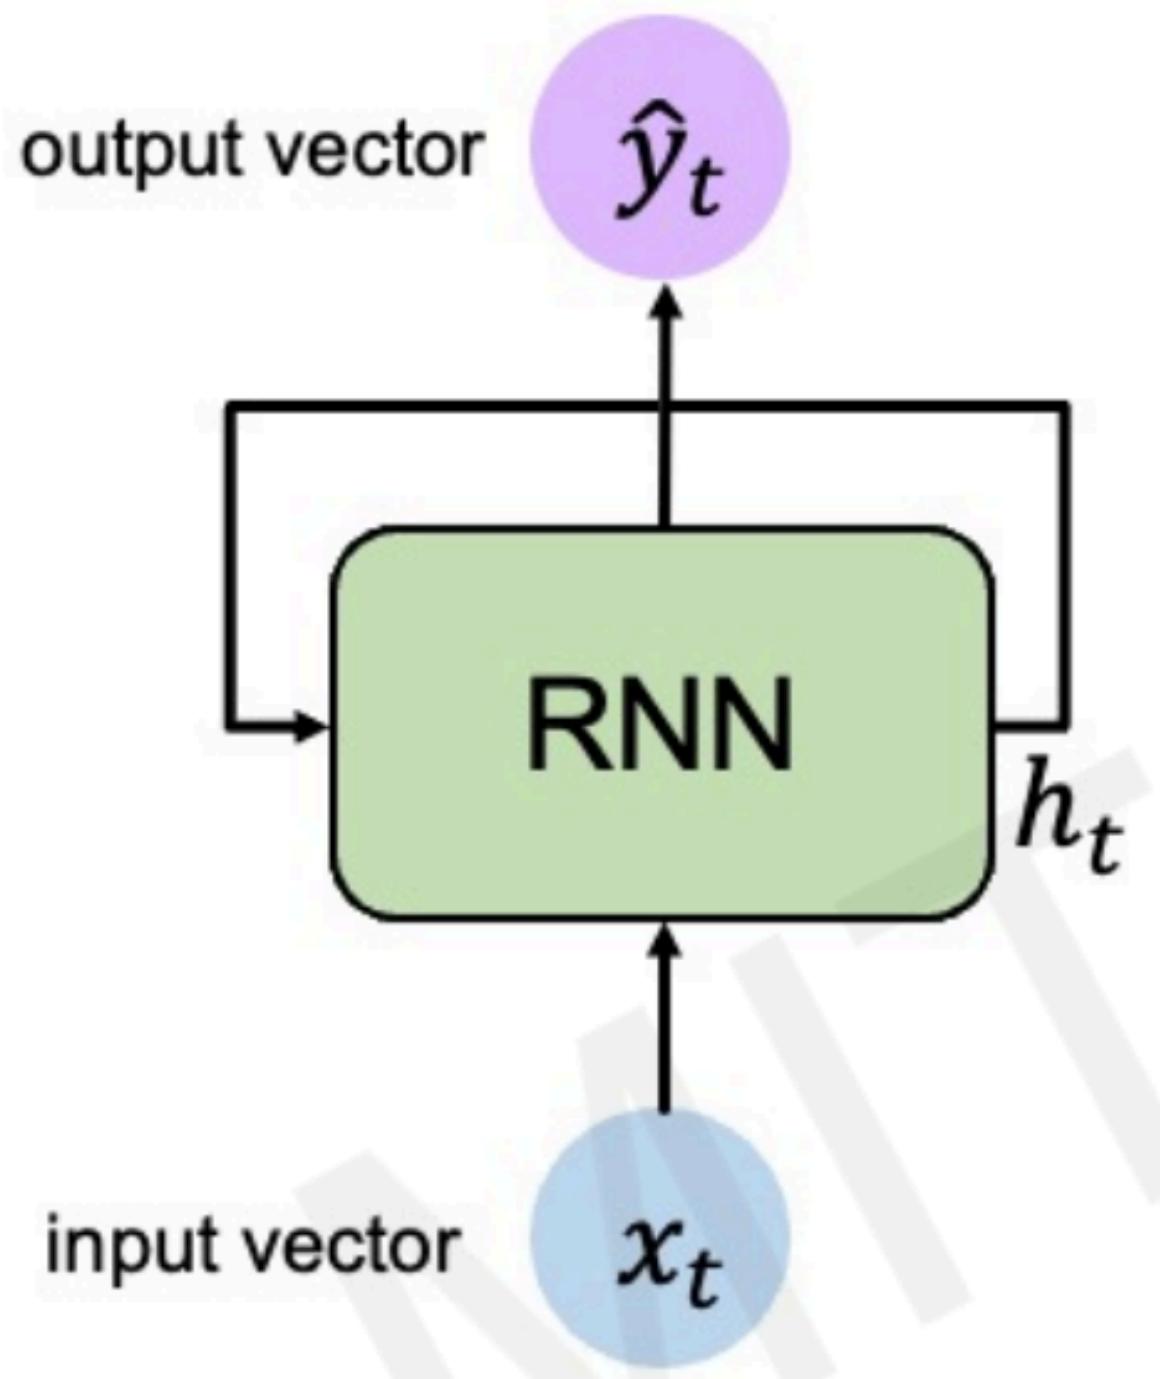
"A baseball player throws a ball."



Many to Many
Machine Translation



Recurrent Neural Networks (RNNs)



Apply a **recurrence relation** at every time step to process a sequence:

$$h_t = f_W(x_t, h_{t-1})$$

cell state function
with weights input
 w
 old state

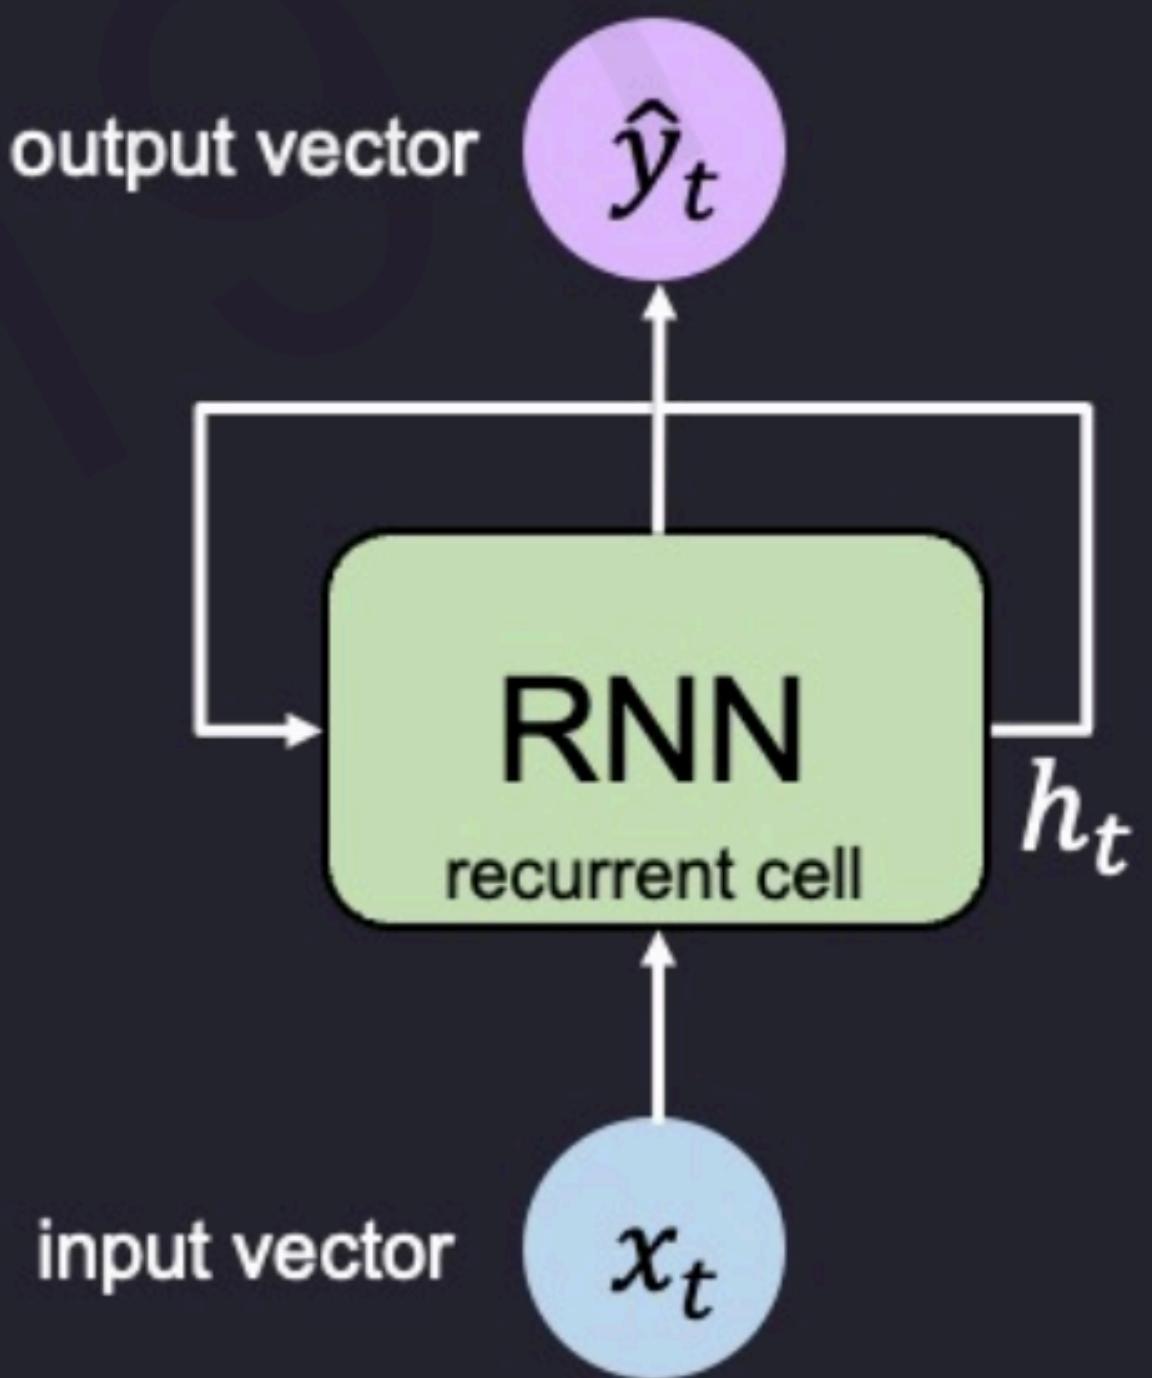
Note: the same function and set of parameters are used at every time step

RNNs have a **state**, h_t , that is updated **at each time step** as a sequence is processed

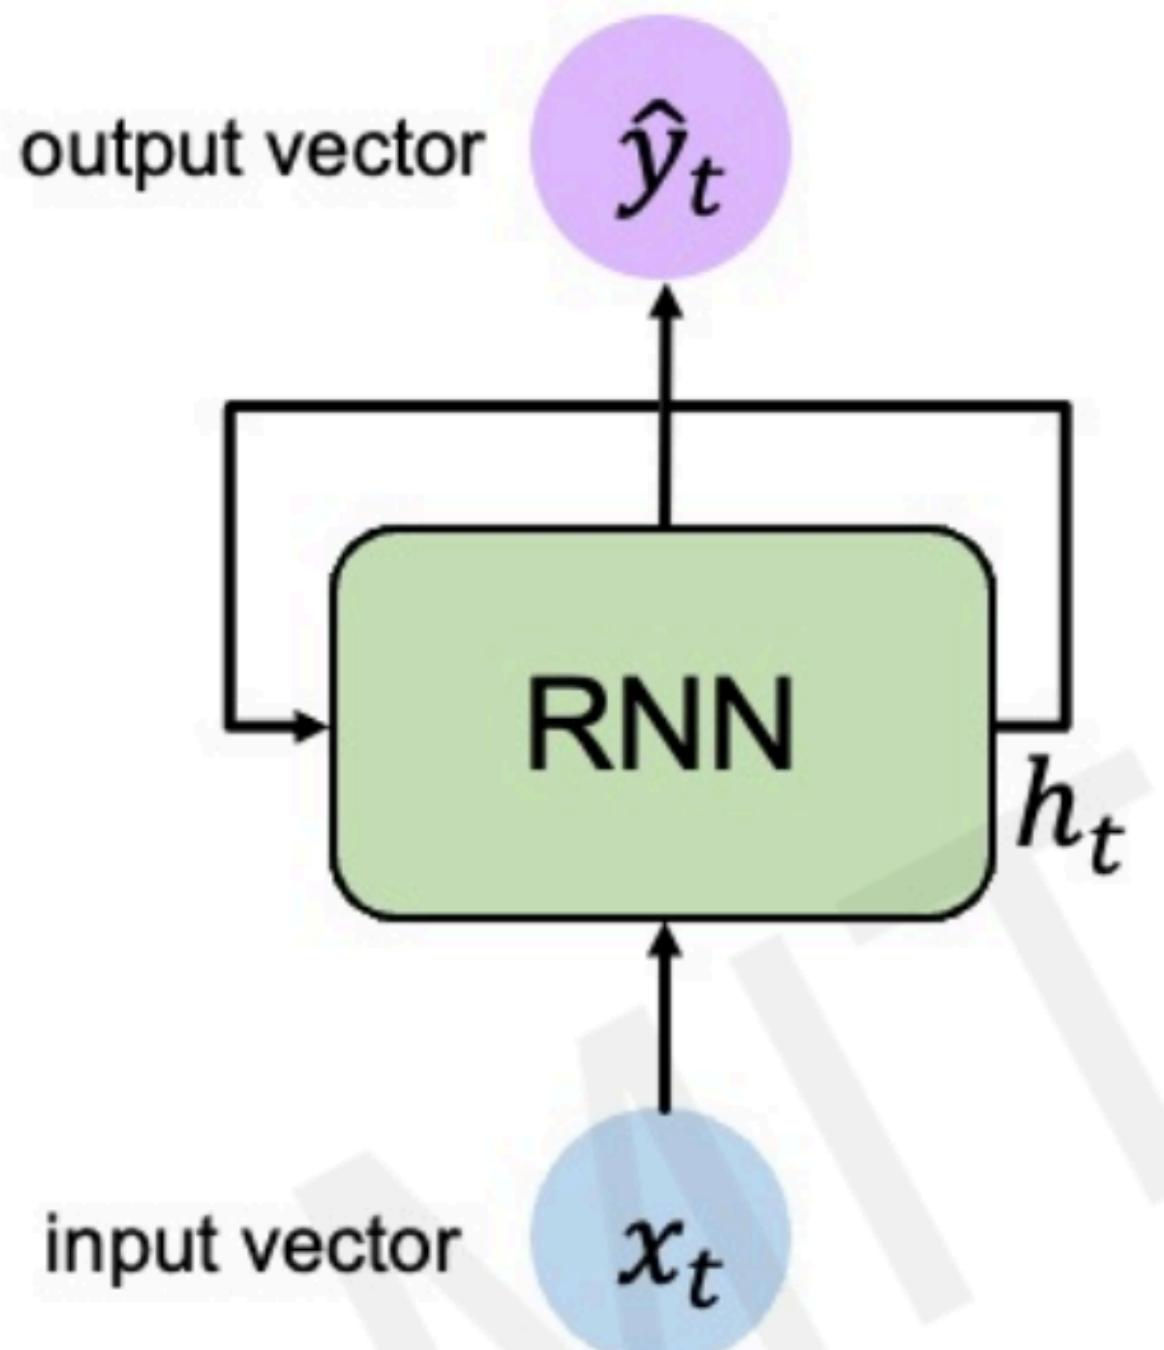
RNN Intuition

```
my_rnn = RNN()  
hidden_state = [0, 0, 0, 0]  
  
sentence = ["I", "love", "recurrent", "neural"]  
  
for word in sentence:  
    prediction, hidden_state = my_rnn(word, hidden_state)
```

```
next_word_prediction = prediction  
# >>> "networks!"
```



RNN State Update and Output



Output Vector

$$\hat{y}_t = W_{hy}^T h_t$$

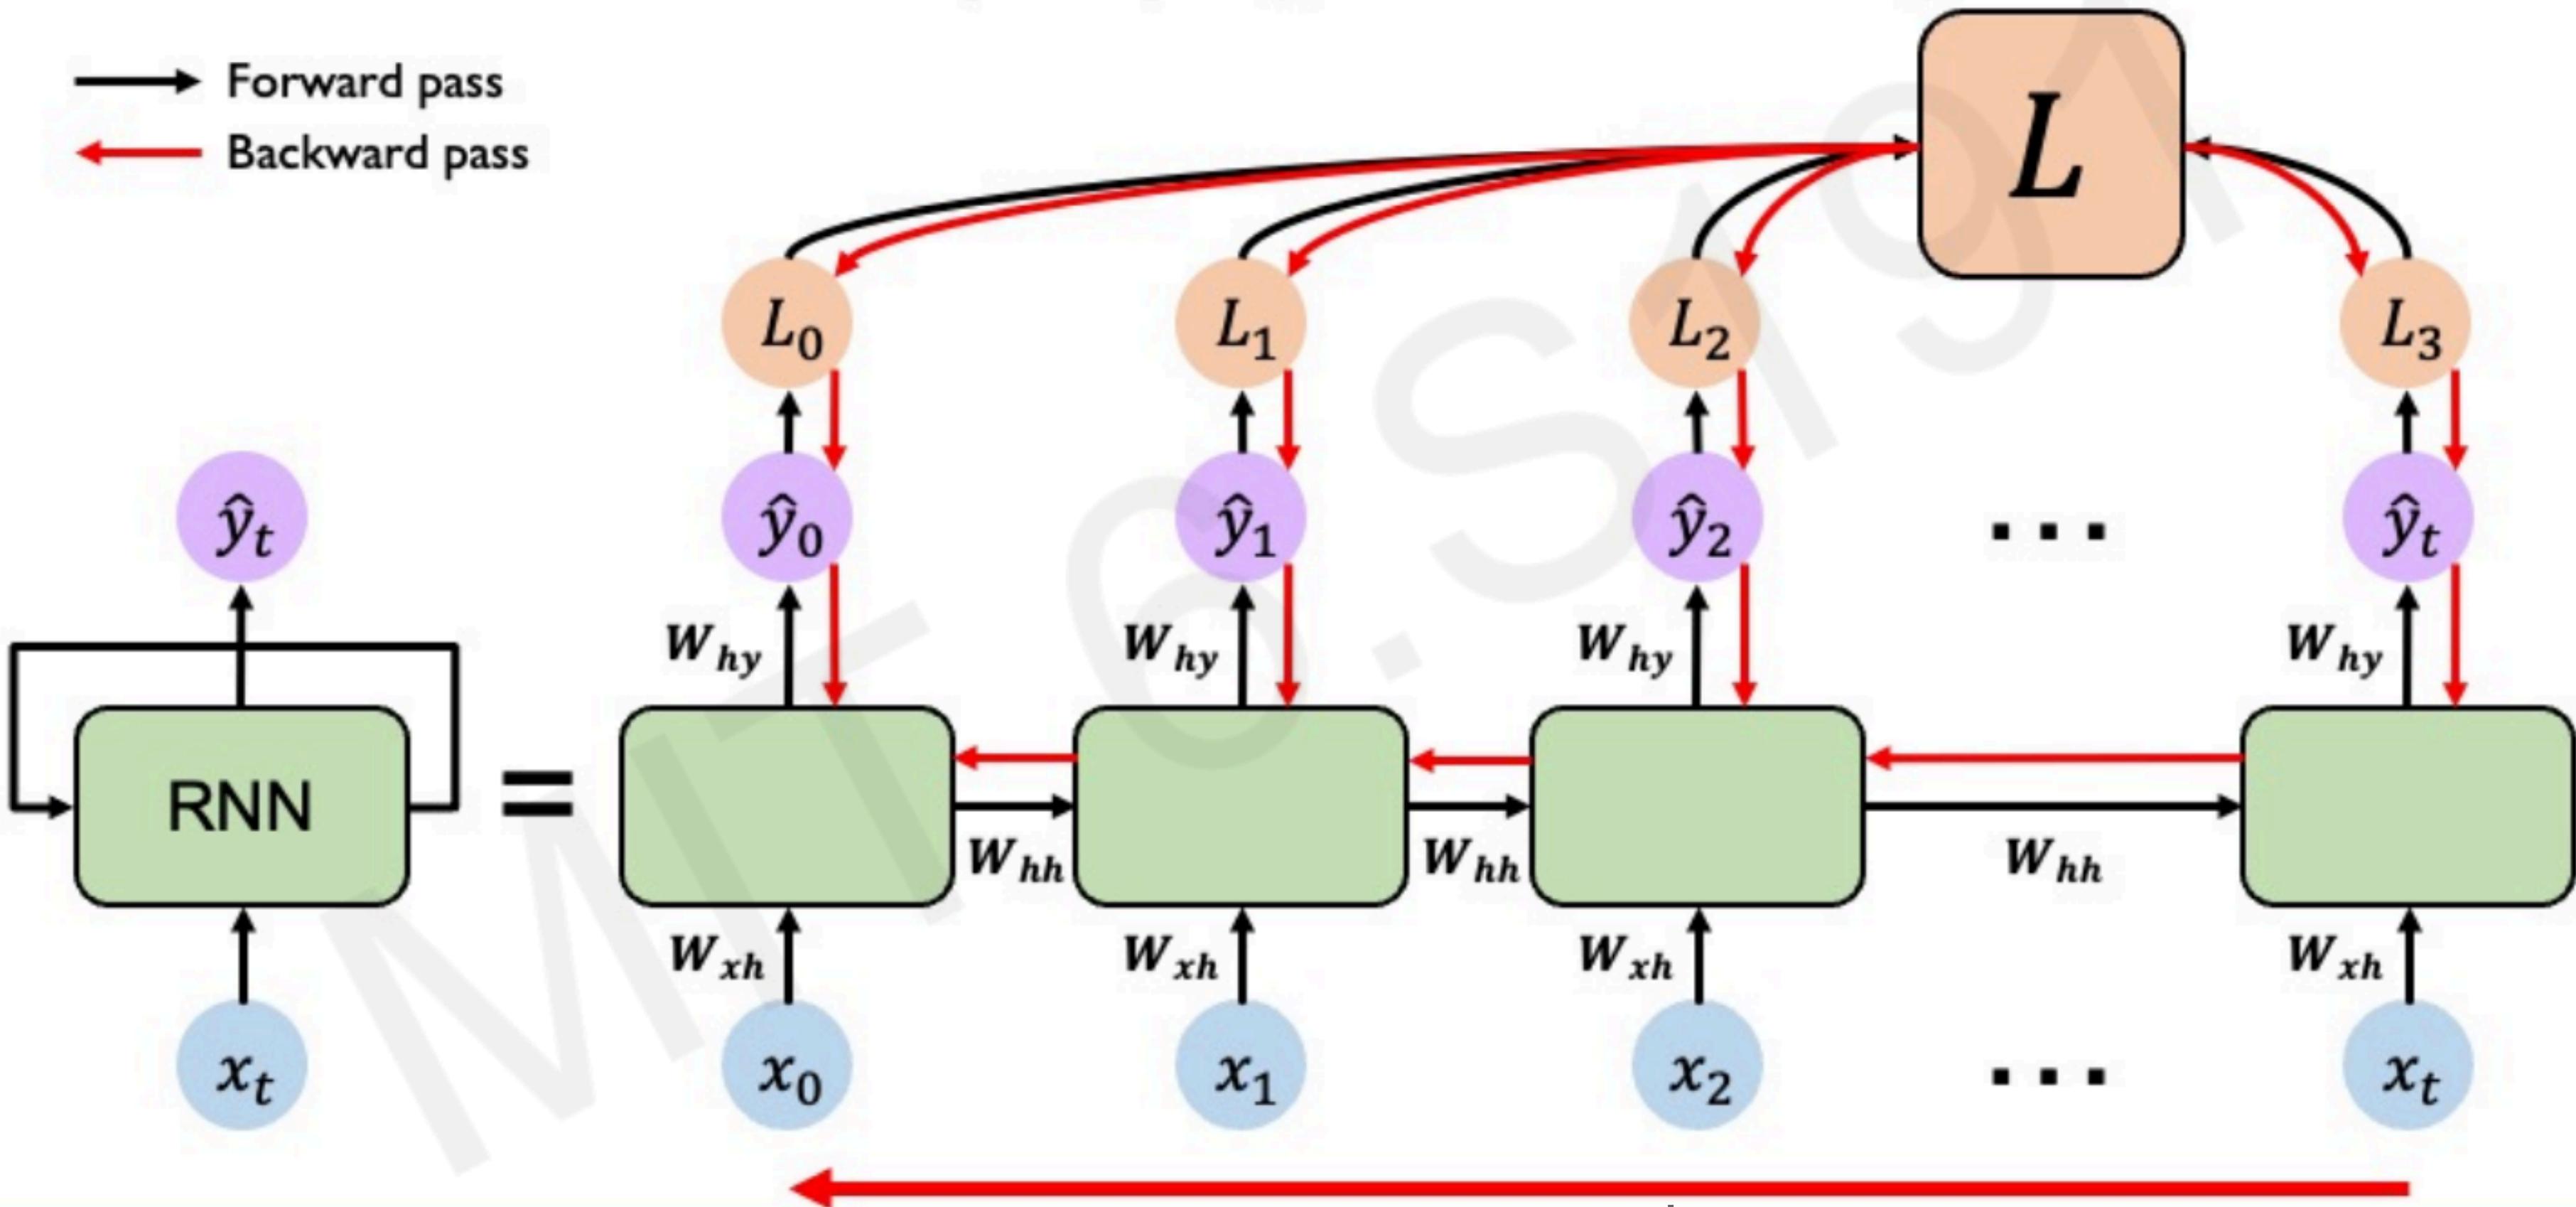
Update Hidden State

$$h_t = \tanh(W_{hh}^T h_{t-1} + W_{xh}^T x_t)$$

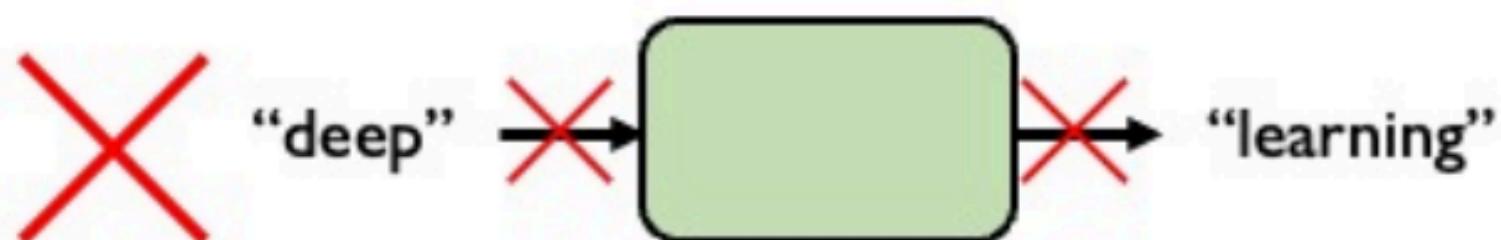
Input Vector

$$x_t$$

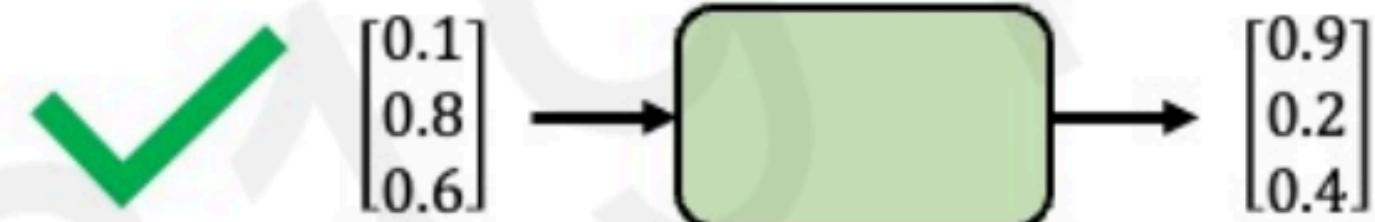
RNNs: Backpropagation Through Time



Encoding Language for a Neural Network

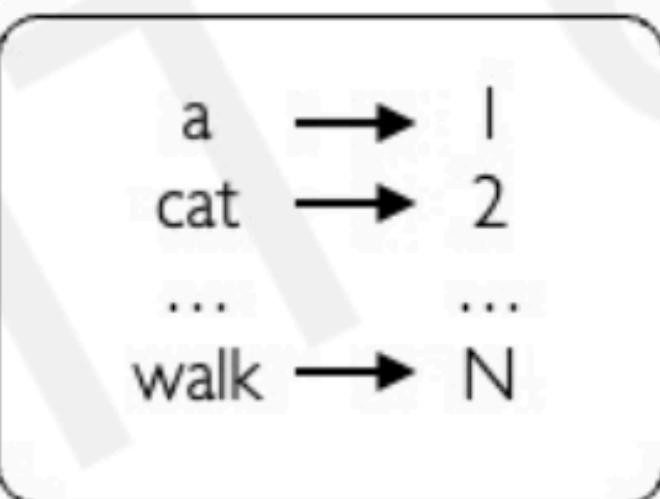
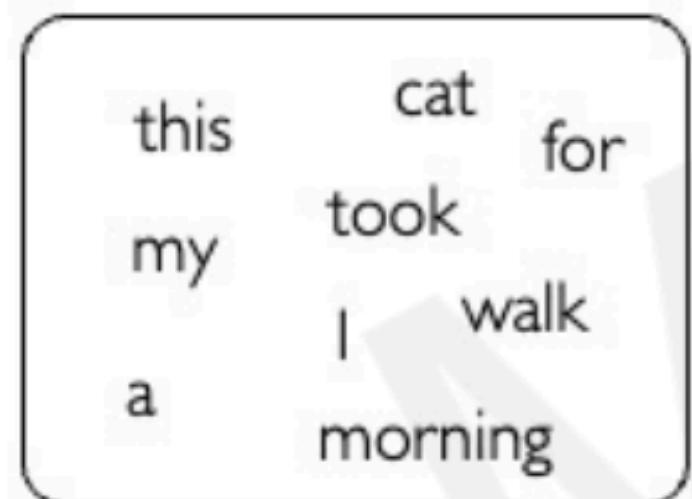


Neural networks cannot interpret words



Neural networks require numerical inputs

Embedding: transform indexes into a vector of fixed size.



1. Vocabulary:

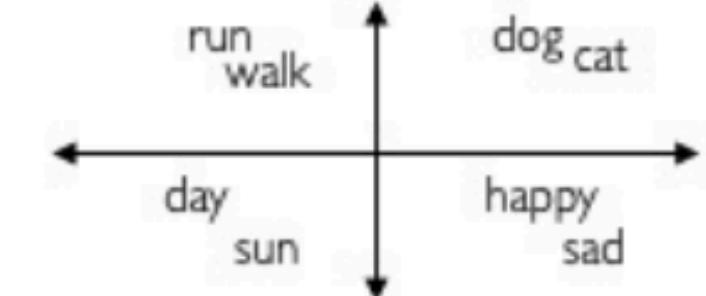
Corpus of words

One-hot embedding

"cat" = $[0, 1, 0, 0, 0, 0]$

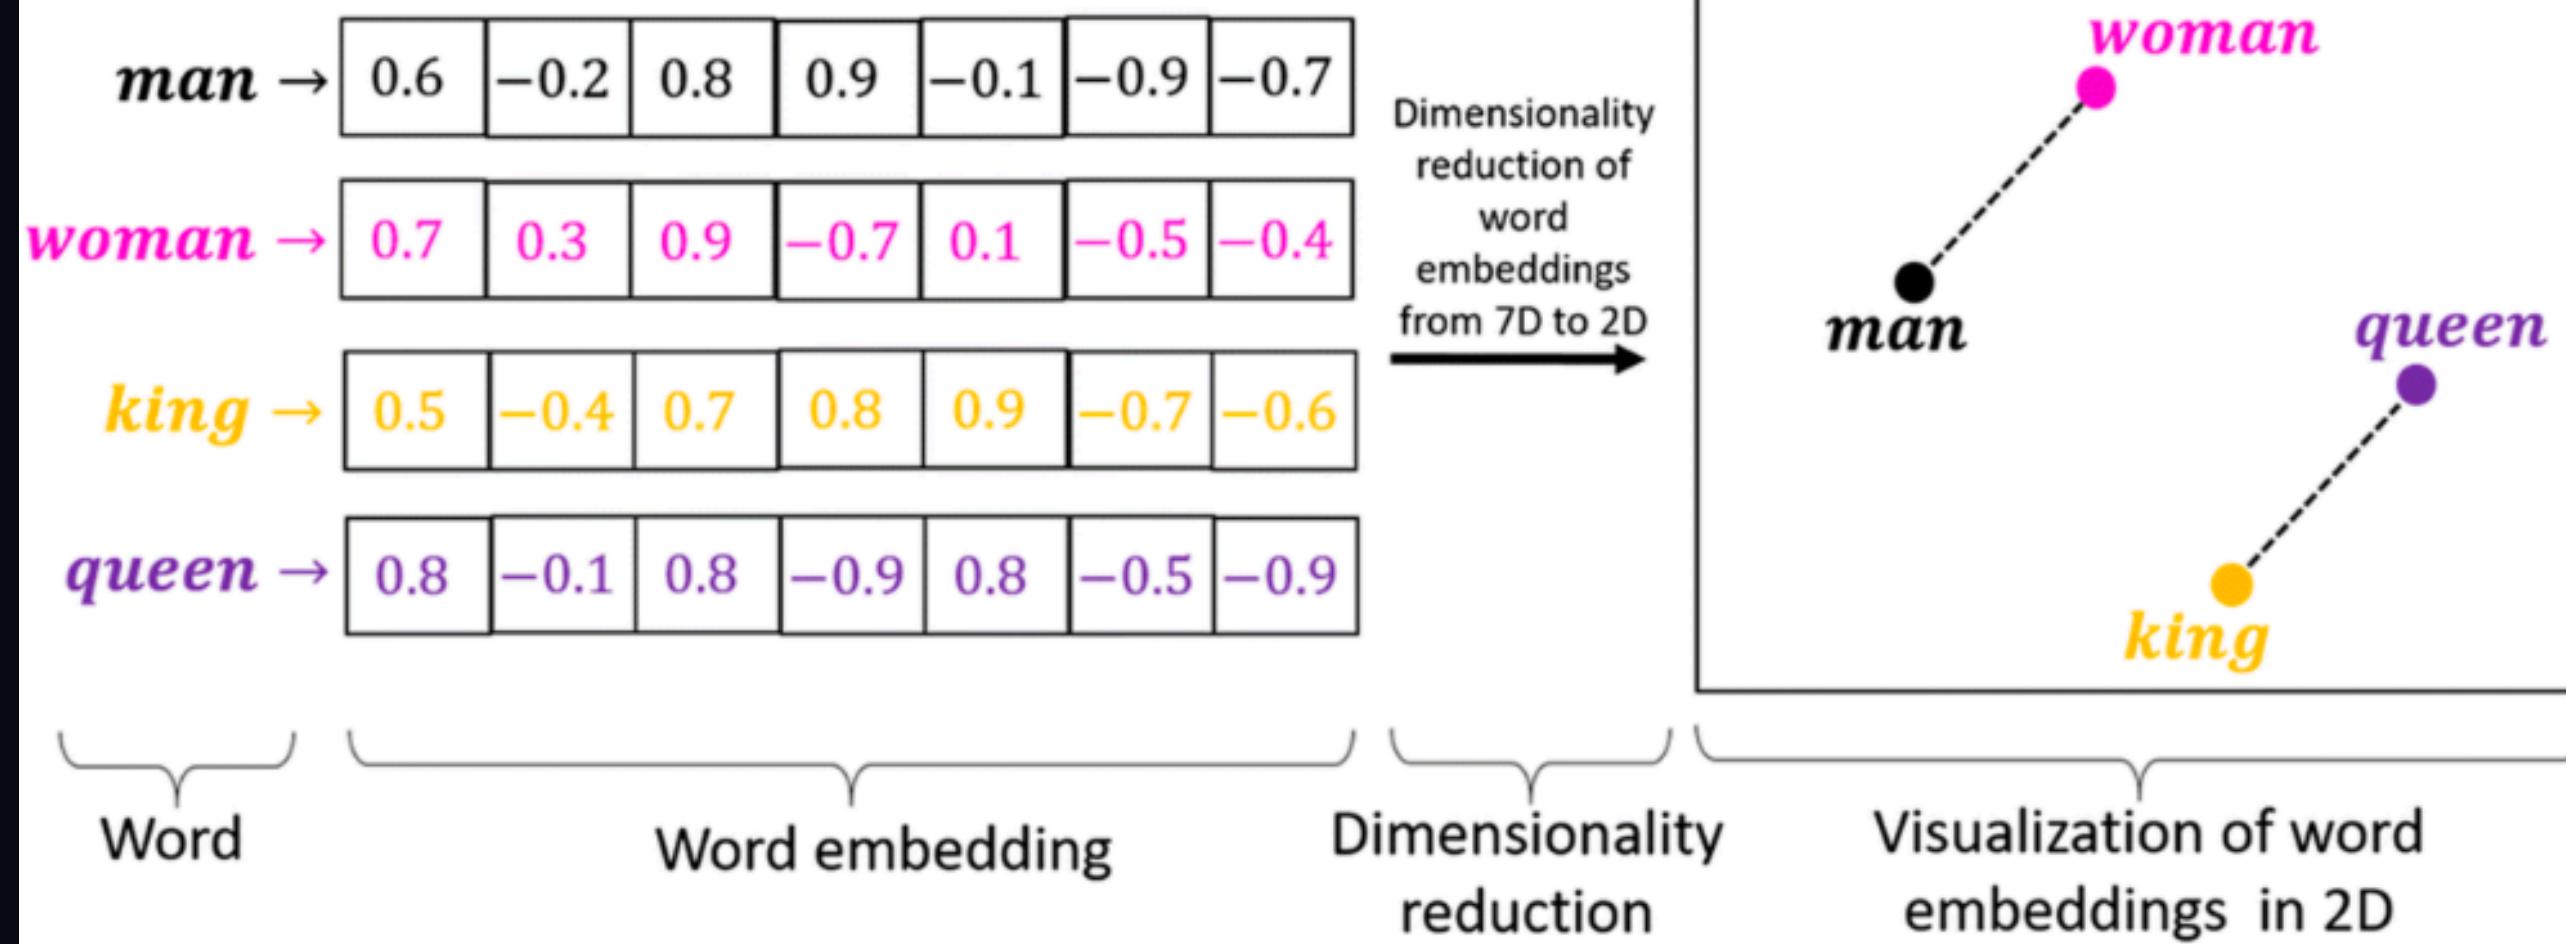
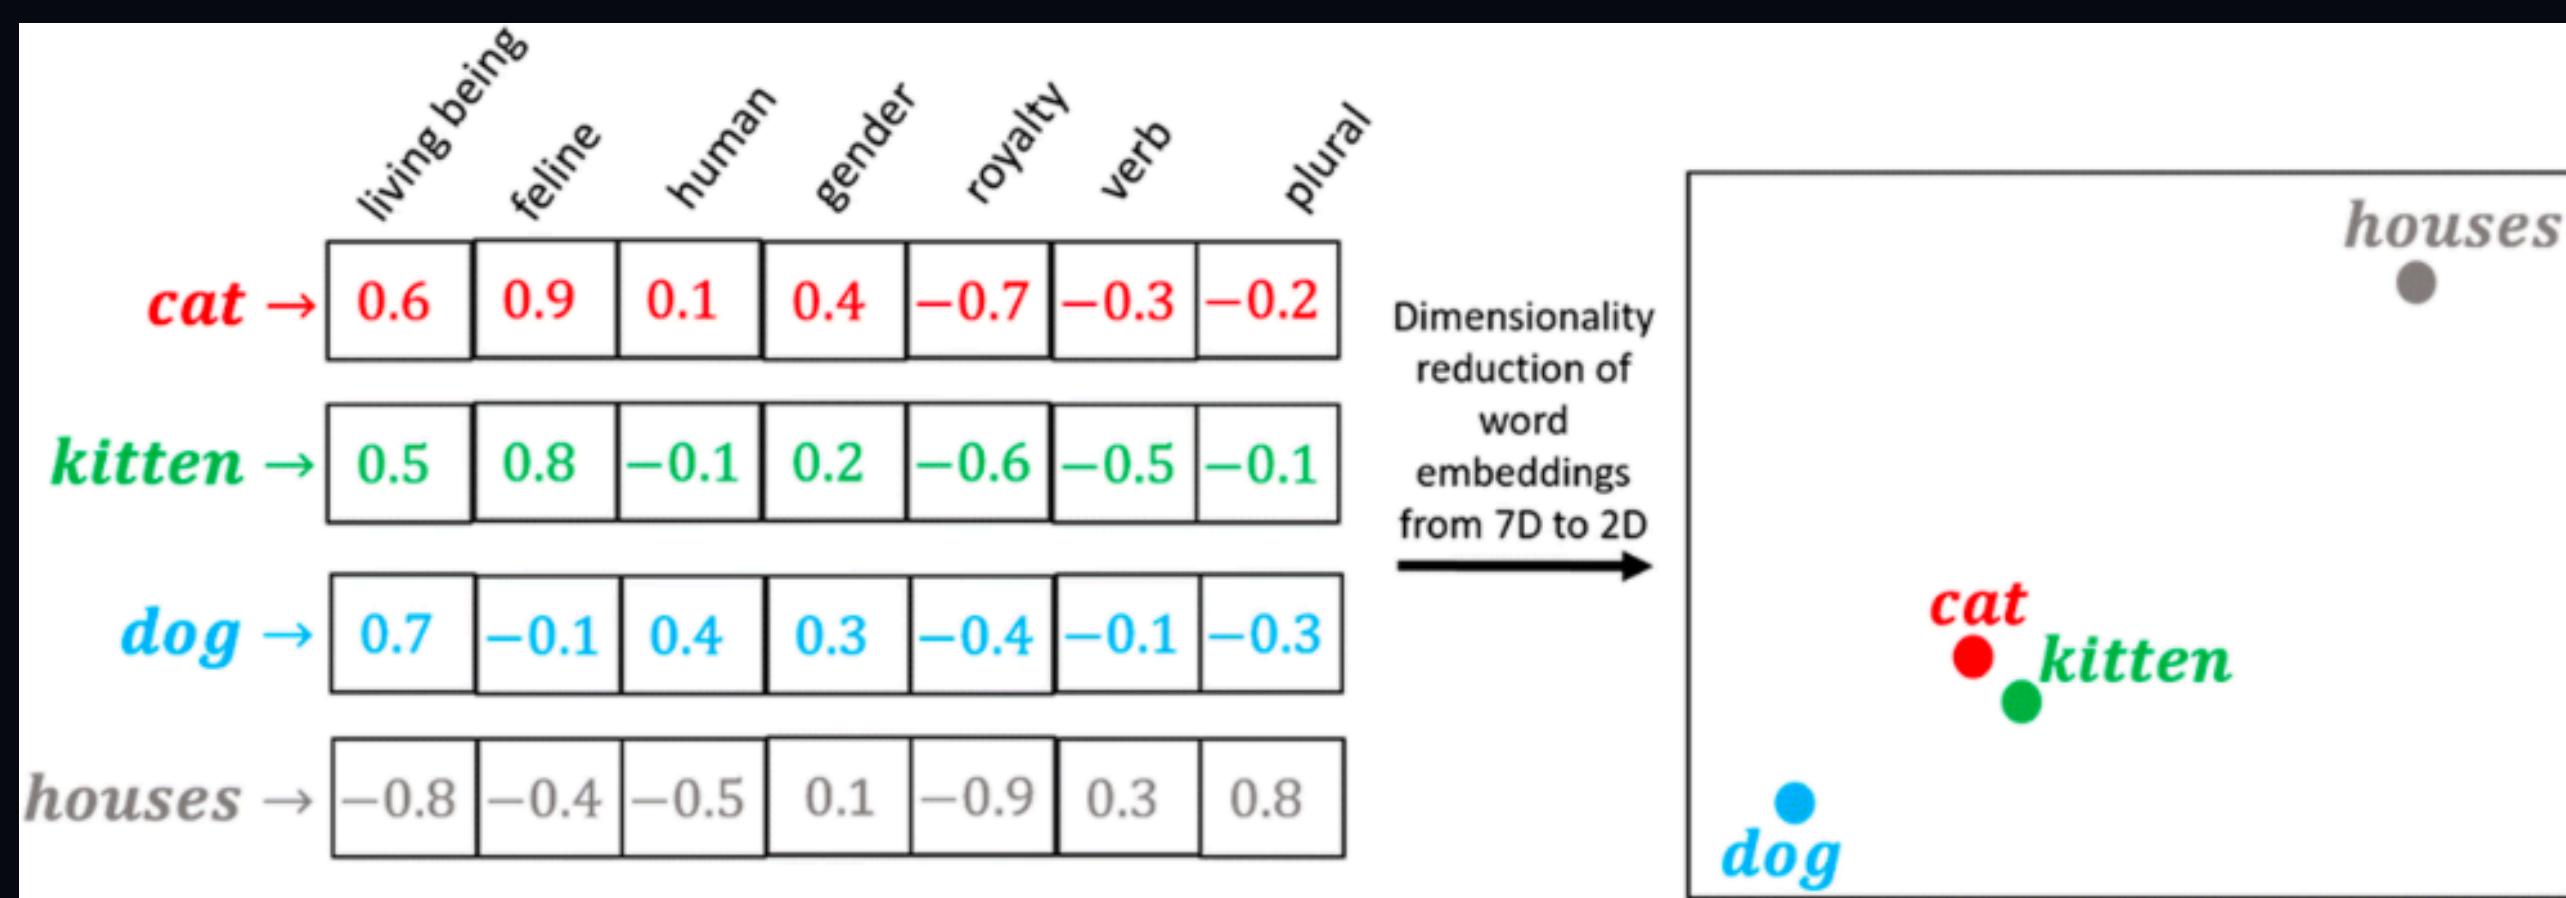
i-th index

Learned embedding

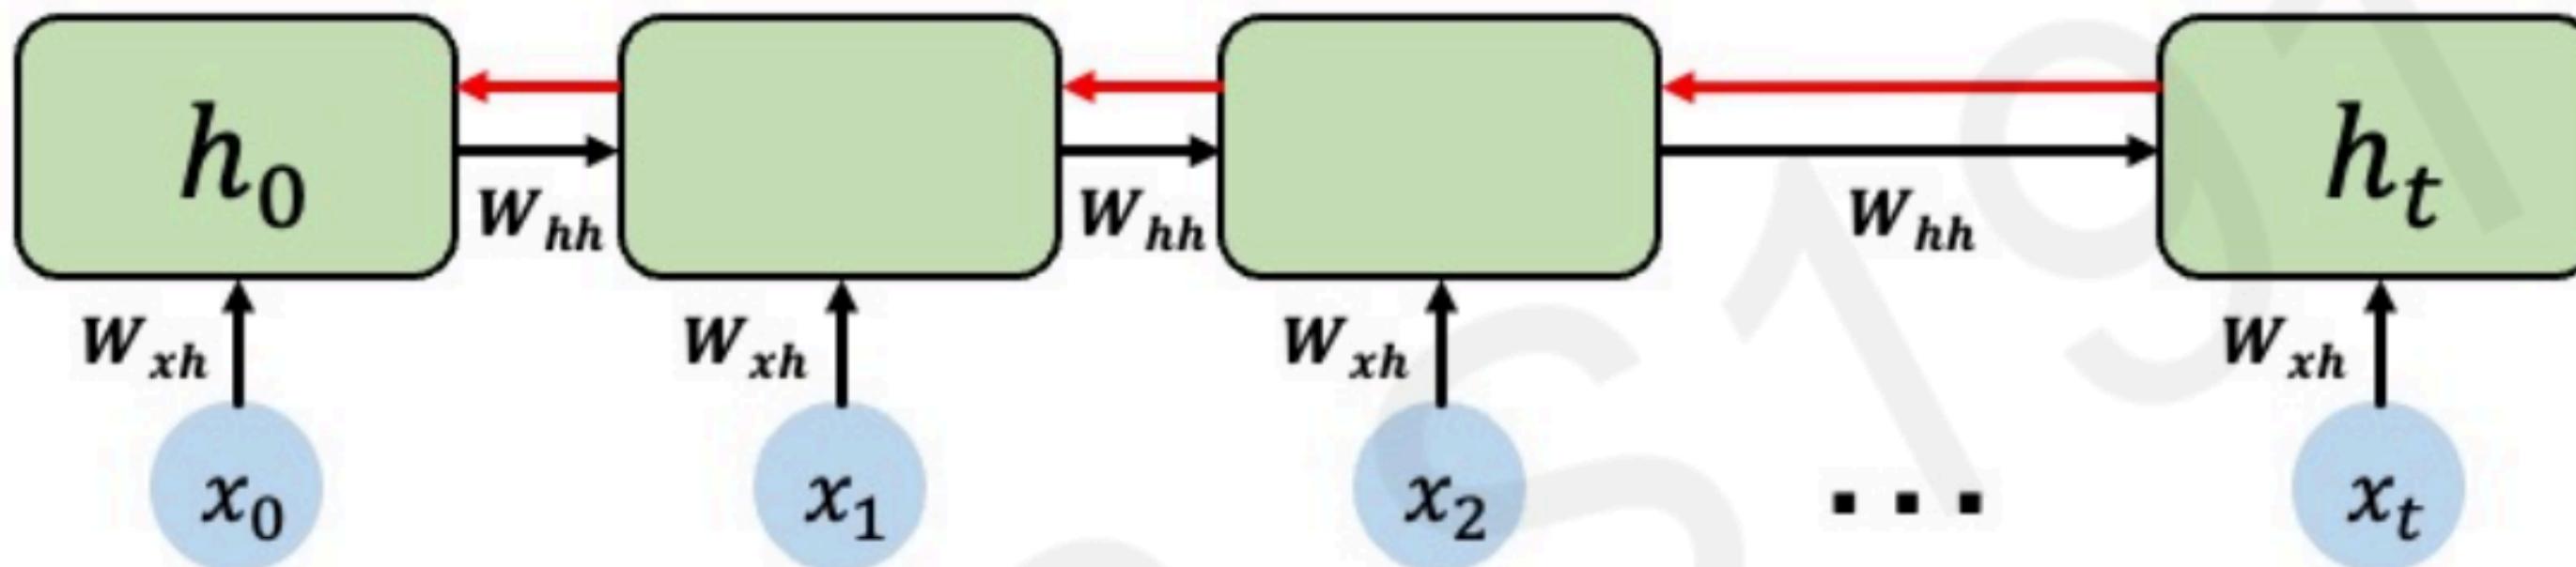


3. Embedding:

Index to fixed-sized vector



Standard RNN Gradient Flow: Vanishing Gradients



Computing the gradient wrt h_0 involves **many factors of W_{hh}** + **repeated gradient computation!**

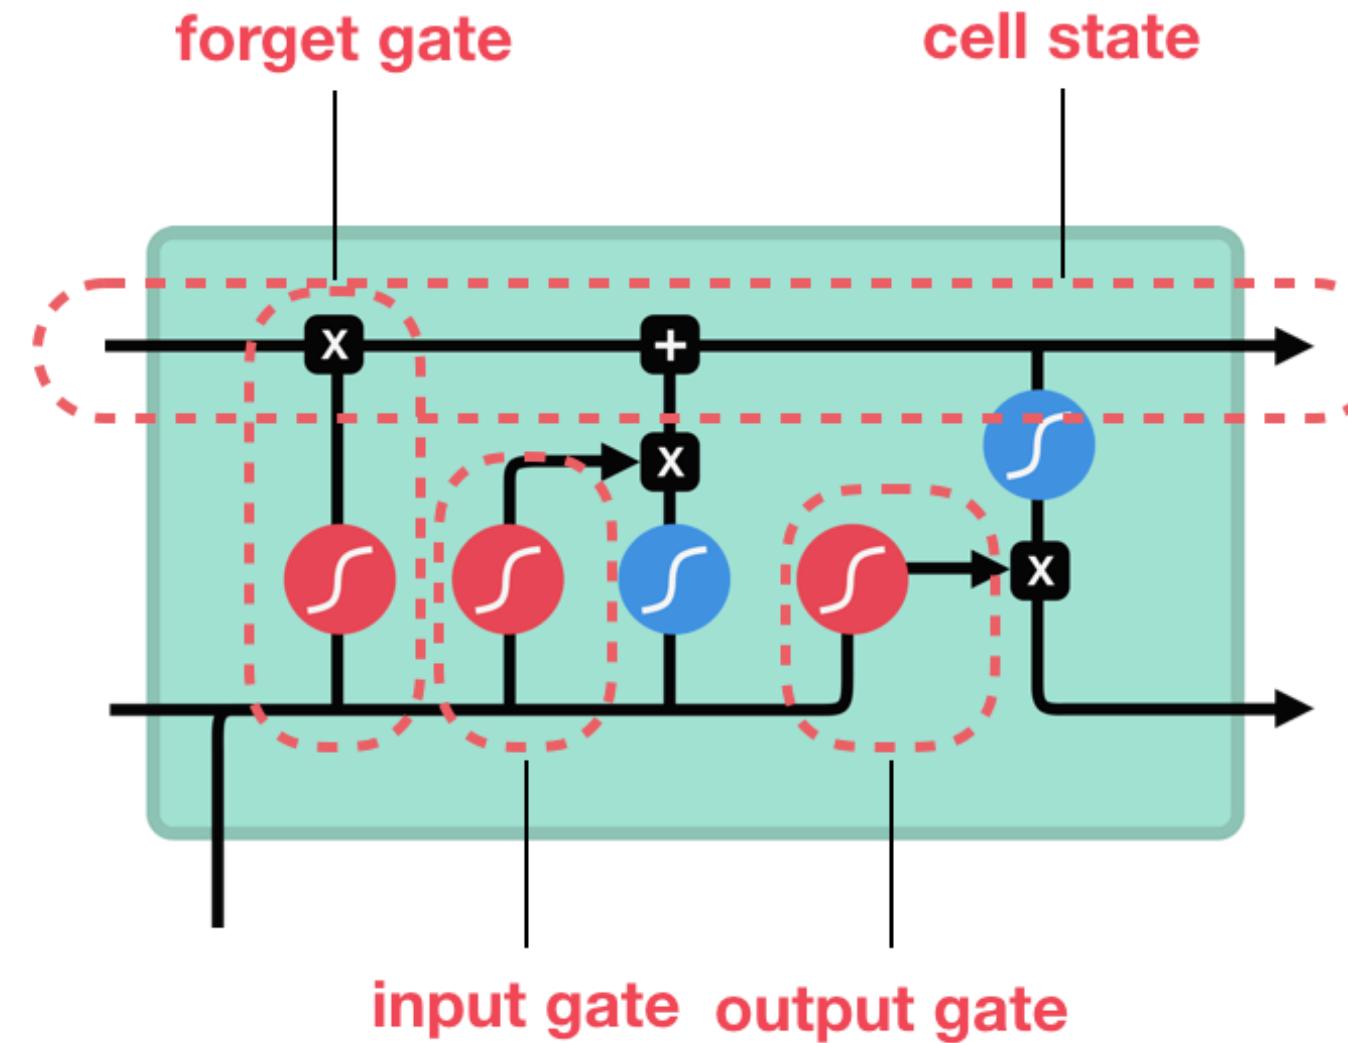
Many values > 1 :
exploding gradients

Gradient clipping to
scale big gradients

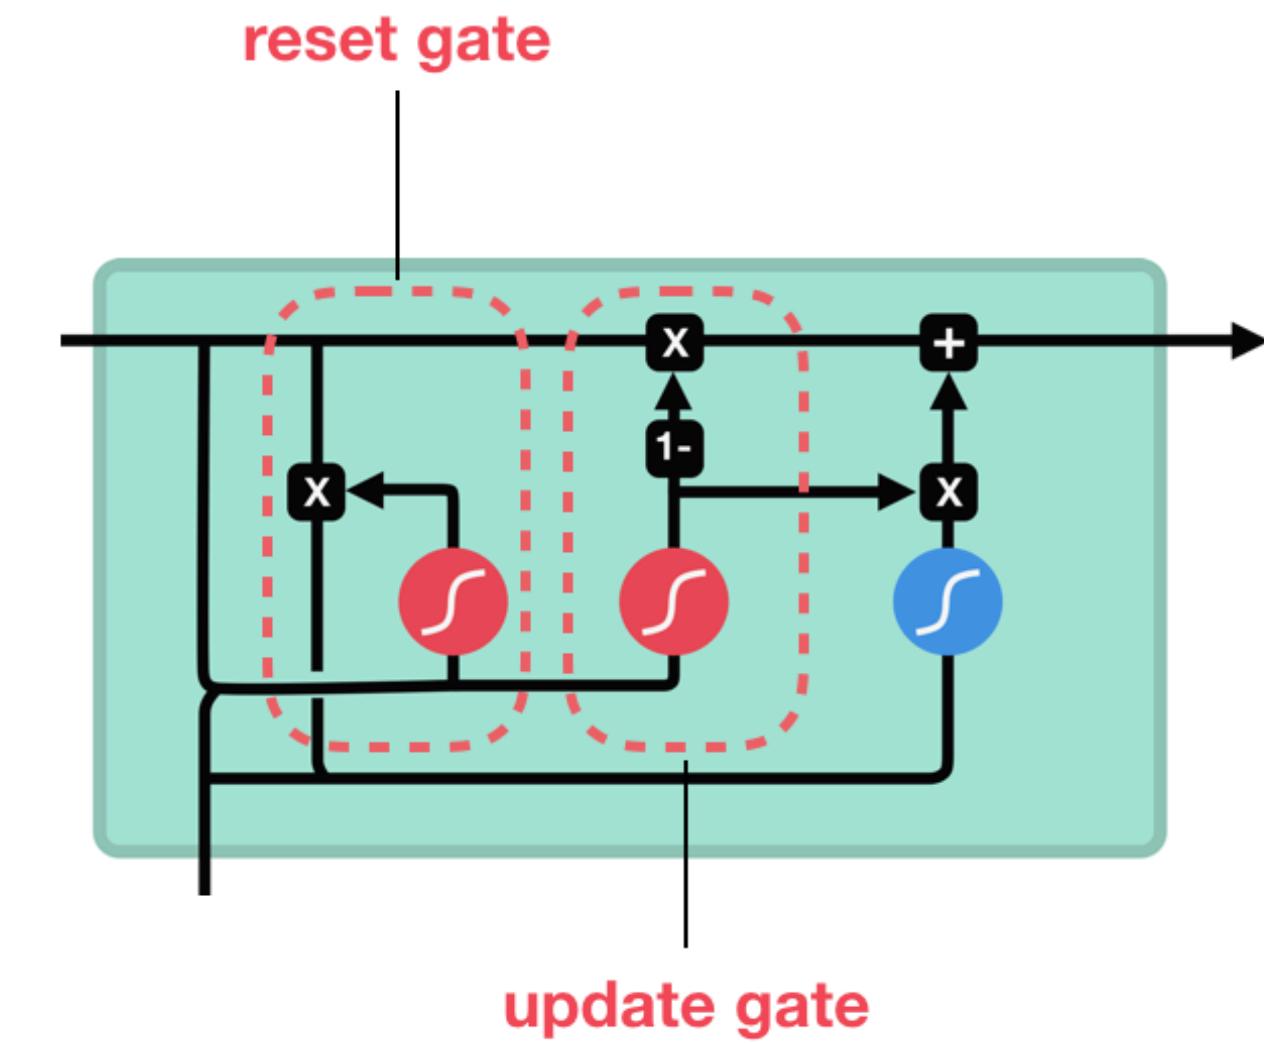
Many values < 1 :
vanishing gradients

1. Activation function
2. Weight initialization
3. Network architecture

LSTM



GRU



sigmoid



tanh



pointwise
multiplication

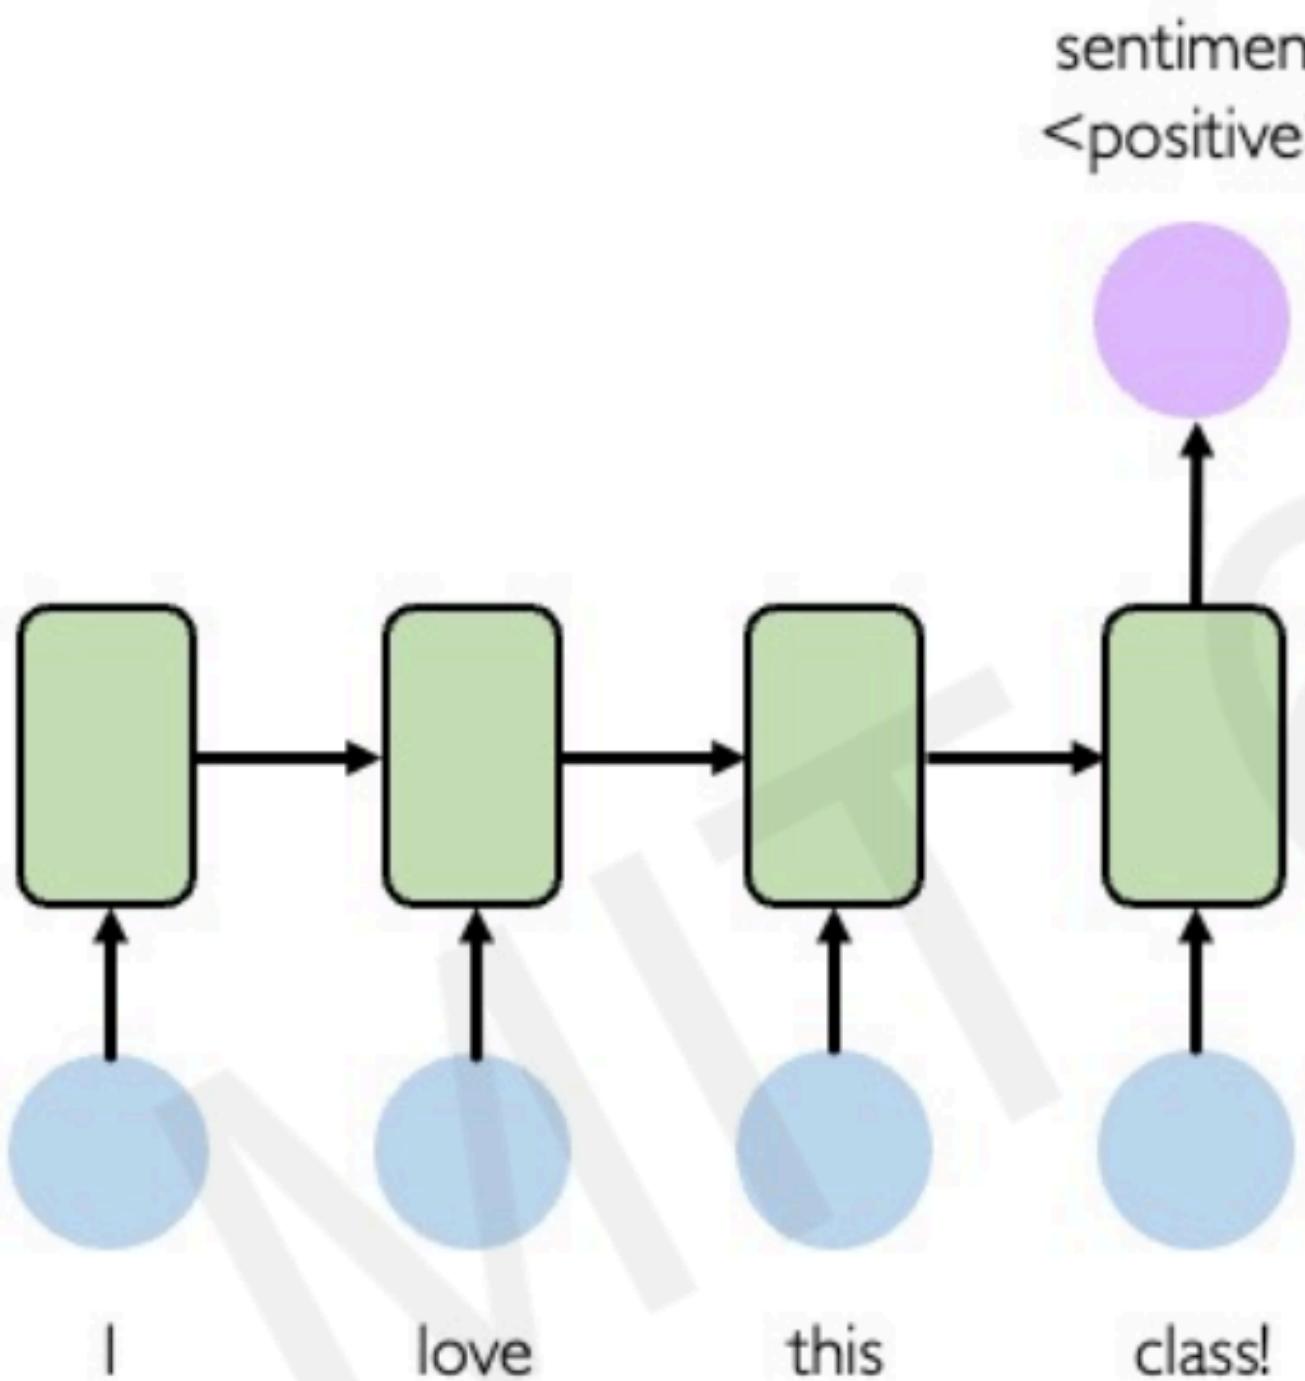


pointwise
addition



vector
concatenation

Example Task: Sentiment Classification



Input: sequence of words

Output: probability of having positive sentiment



`loss = tf.nn.softmax_cross_entropy_with_logits(y, predicted)`

Their customer service





```
EMBEDDING_DIM = 128
LSTM_UNITS = 64

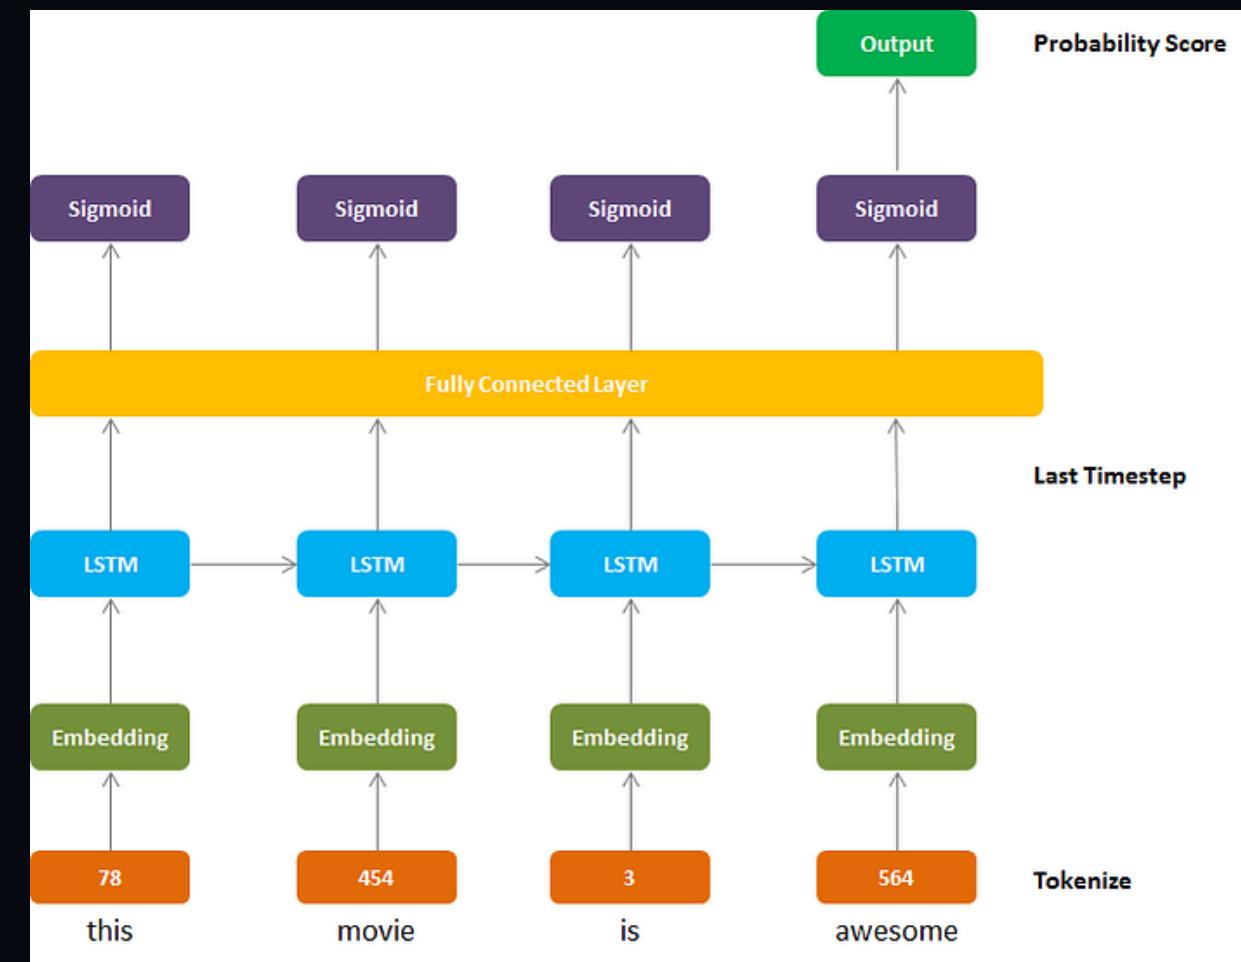
model_lstm = Sequential([
    # Étape 1: Créer un "dictionnaire de sens" pour chaque mot.
    # input_dim est la taille de notre vocabulaire.
    # output_dim est la taille du vecteur de sens que nous voulons pour chaque mot.
    Embedding(input_dim=VOCAB_SIZE, output_dim=EMBEDDING_DIM, input_length=MAX_LEN),

    # Étape 2: La couche LSTM qui lit la séquence et se souvient du contexte.
    # On l'enrobe dans Bidirectional pour lire la phrase dans les deux sens.
    Bidirectional(LSTM(units=LSTM_UNITS, dropout=0.2, recurrent_dropout=0.2)),

    # Étape 3: Une couche dense pour un traitement final.
    Dense(64, activation='relu'),
    Dropout(0.5), # Un autre Dropout pour la régularisation

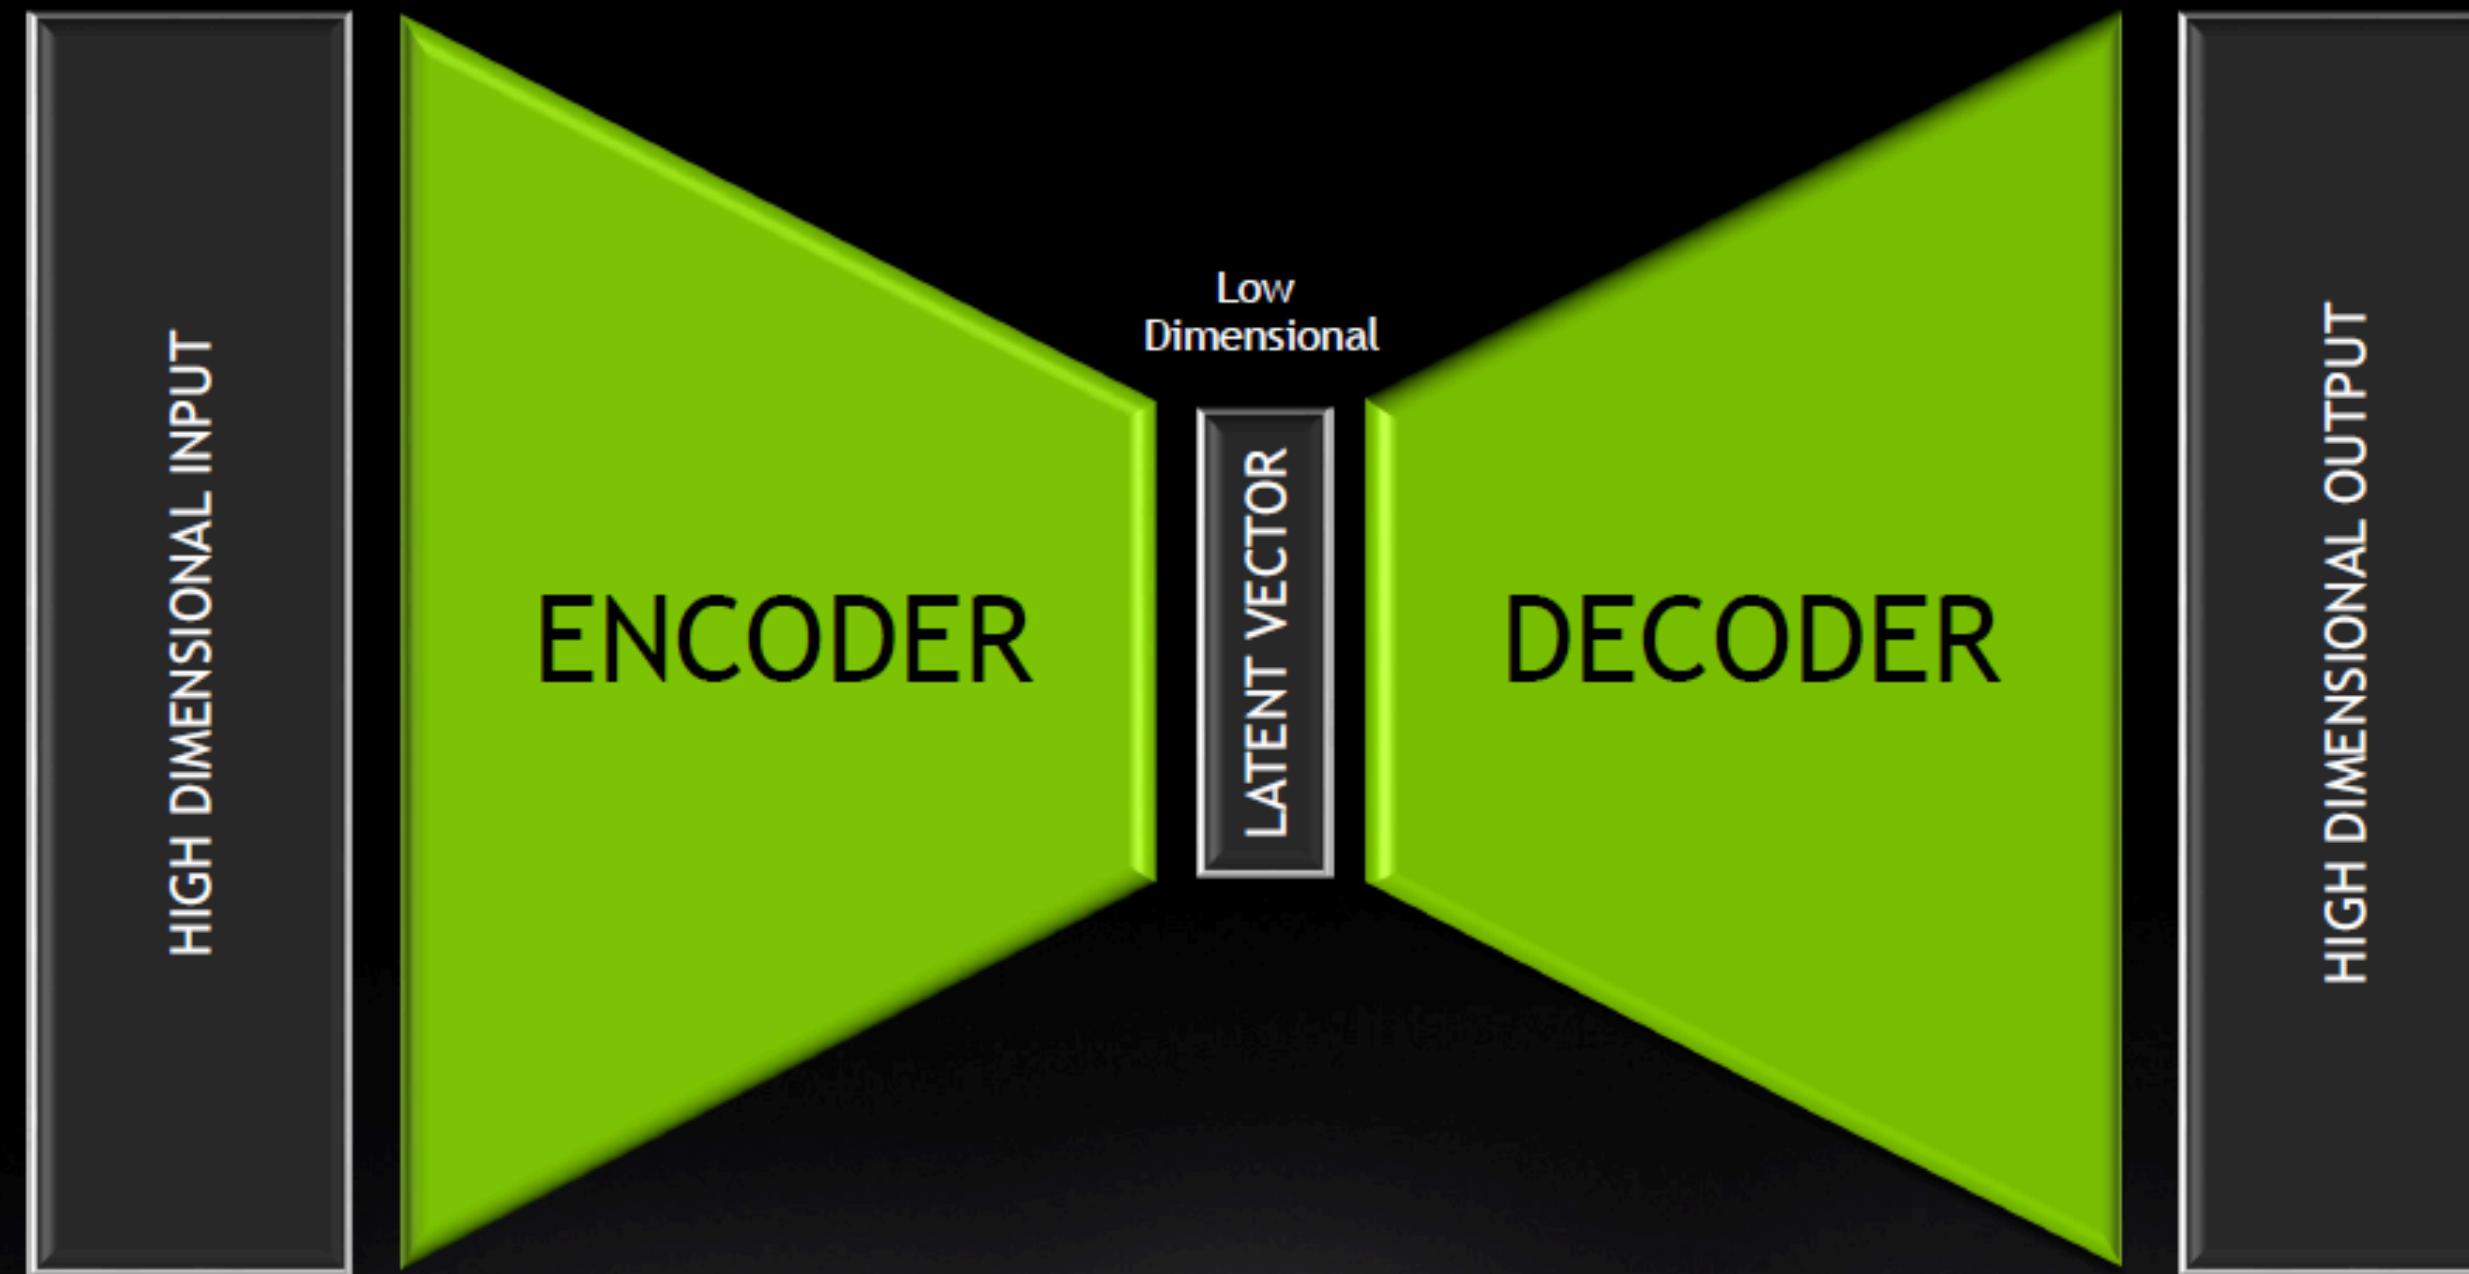
    # Étape 4: Le neurone de sortie qui donne la probabilité d'un sentiment positif.
    Dense(1, activation='sigmoid')
])

# Affichons le plan de notre architecture
model_lstm.summary()
```



ENCODERS AND DECODERS

Networks connecting high and low dimensional spaces



THANK YOU

ANY QUESTION?